

The Silicon Wire

VOL. 1, NO. 005

2026-03-30

COVER STORY

Getting Towards Real Sandbox Containers

Containers are all the rage right now. At the very core of containers are the same Linux primitives that are also used to create application sandboxes.

Jessie Frazelle

Containers are all the rage right now.

At the very core of containers are the same Linux primitives that are also used to create application sandboxes. The most common sandbox you may be familiar with is the Chrome sandbox. You can read in detail about the Chrome sandbox here: chromium.googlesource.com/chromium/src/+master/docs/linux_sandboxing.md. The relevant aspect for this article is the fact it uses user namespaces and seccomp. Other deprecated features include AppArmor and SELinux. Sound familiar? That's because containers, as you've come to know them today, share the same features.

id="why-are-containers-not-currently-being-considered-a-sandbox">Why are containers not currently being considered a "sandbox"?

One of the key differences between how you run Chrome and how you run a container are the privileges used. Chrome runs as your own unprivileged user. Most containers (be it docker, runc, or rkt) run as root.

Yes, we all know that containers run unprivileged processes; but creating and running the containers themselves requires root privileges at some point.

id="how-can-we-run-containers-as-an-unprivileged-user">How can we run containers as an unprivileged user?

Easy! With user namespaces, you might say. But it's not exactly that simple. One of the main differences between the Chrome sandbox and containers is cgroups. Cgroups control what a process can use. Whereas namespaces control what a process can see. Containers have cgroup resource management built in. Creating cgroups from an unprivileged user is a bit difficult, especially device control groups.

If we ignore, for the time being, this huge tire fire that is creating cgroups as an unprivileged user, then unprivileged containers are easy. User namespaces allow us to create all the namespaces without any further privileges. The one key caveat being that the {uid,gid}_map must have the current host user mapped to the container uid that the process will be run as. The size of the {uid,gid}_map can also only be 1. For example if you are running as uid 1000 to spawn the container, your {uid,gid}_map for the process would be 0 1000 1 for uid 0 in the container. The 1 there refers to the size.

id="how-is-this-different-than-the-user-namespace-support-currently-in-docker">How is this different than the user namespace support currently in Docker?

This is quite different, but for very good reason. In Docker, by default, when the remapped user is created, the /etc/subuid and /etc/subgid files are populated with a contiguous 65536 length range of subordinate user and group IDs, starting at an offset based on prior entries in those files. Docker's implementation has a larger range of users that can exist in the container as well as having a more "anonymous" mapped host user. If you want to read more about the user namespace implementation in Docker I would checkout @estesp's blog or the the docker docs .

id="poc-or-gtfo">POC or GTFO

As a proof of concept of unprivileged containers without cgroups I made binctr . Which spawned a mailing list thread for implementing this in runc/libcontainer . Aleksa Sarai has started on a few patches and this might actually be a reality pretty soon!

Update: it took almost a year, but this was added to runc in Mar 2017.

id="where-does-this-put-us-in-the-sandbox-landscape">Where does this put us in the "sandbox" landscape?

With this implementation we get:

namespaces

apparmor

selinux

seccomp

capabilities limiting

all created by an unprivileged user!

Sandboxes should be very application-specific, using custom AppArmor profiles, Seccomp profiles and the like. A generic container will never be equivalent to a sandbox because it's too universal to really lock down the application.

Containers are not going to be the answer to preventing your application from being compromised, but they can limit the damage from a compromise. The world an attacker might see from inside a very strict container with custom AppArmor/Seccomp profiles greatly differs than that without the use of containers. With namespaces we limit the application from seeing various things such as network, mounts, processes, etc. And with cgroups we can further limit what the attacker can use, be it a large amount of memory, cpu, or even a fork bomb.

id="but-what-about-cgroups">But what about cgroups?

We can set up cgroups for memory, blkio, cpu, and pids with an unprivileged user as long as the cgroup subsystem has been chowned to the correct user. Devices are a different story though. Considering the fact you cannot mknod in a user namespace it is not the worst thing in the world.

Let's not completely rule out the devices cgroup. In the future this might be entirely possible. In kernels 4.6+, there is a new cgroup namespace. For now all this does is mask the cgroups path inside the container so it is not entirely useful for unprivileged containers at all. But in the future maybe it could be (if we ask nice enough?).

id="what-is-the-awesome-sauce-we-all-gain-from-this">What is the awesome sauce we all gain from this?

Well judging by the original GitHub issue about unprivileged runc containers, the largest group of commenters is from the scientific community who are restricted to not run certain programs as root.

But there is so much more that this can be used for. One of my most anticipated use cases is the work being done by

Alex Larsson on xdg-app to run applications in sandboxes. Definitely checkout bubblewrap if you are interested in this.

Also subgraph, the container based OS which specializes in security and privacy, have this same idea in mind.

I am a huge fan of running desktop applications in containers as well as solving multi-tenancy for running containers. I definitely hope to help evolve containers into real sandboxes in the future.

{

FEATURES

10 Hacks Every Ring User Should Know

Emily Long · Lifehacker

When setting up any new internet-connected device, don't stick with the default settings. Doing so introduces security risks, and it's also a less-than-optimal way to use the features available to you. If you have a Ring camera or doorbell, there are a whole host of changes you can make to minimize annoyance and maximize privacy.

id="snooze-motion-alerts-when-youre-outside">Snooze Motion Alerts when you're outside

Motion alerts are among the most useful features of any security camera, but you don't need a notification to your phone when you are the one moving around your property. You can snooze alerts in certain situations, such as when you're outside doing yard work or hosting a party. Global Snooze pauses alerts for all cameras and doorbells for a set duration, while Alerts Snooze allows you to pause notifications from a specific device. When Snooze is enabled, you'll still get Doorbell Rings and Priority Alerts.

In the Ring app, tap the motion icon, choose the snooze duration, and tap Start Snooze . For a single device, tap the More icon on the camera you want to snooze and tap the bell icon to turn Alerts Snooze on or off.

If you have monitoring via Virtual Security Guard, you can turn on Motion Snooze for enrolled devices to temporarily pause that service.

id="optimize-motion-zones-to-exclude-certain-areas">Optimize Motion Zones to exclude certain areas

Another way to curate motion alerts is to customize your Ring camera's Motion Zones—for example, to exclude busy streets with lots of cars driving by as well as private, low-traffic areas you don't need to monitor. You can add up to three motion zones per device under your camera's Settings > Motion Settings > Camera Motion Zones . Tap Add Zone , drag the edges of the zone boundaries within your camera's view, and hit Save .

id="set-up-smart-alerts-to-reduce-unnecessary-notifications">Set up Smart Alerts to reduce unnecessary notifications

In addition to narrowing motion zones, you can choose whether you receive motion alerts and/or recordings specifically for people, vehicles, and packages while minimizing notifications from other sources of movement. (Not all Ring devices are compatible with all three Smart Alert categories, and you'll need a Ring subscription to use these features.) To enable Smart Alerts, go to your device's Settings > Motion Settings > Smart Alerts , then tap Enable Feature > Continue and choose your preferred alerts.

id="customize-your-neighborhood-area-for-relevant-alerts">Customize your Neighborhood Area for relevant alerts

Neighbors is an online community through which Ring users can share footage and receive updates for their area. It is the broad umbrella for controversial features like Community Requests and Search Party (which I'll get into below), but it could be useful for keeping abreast of issues in your neighborhood—like fire or other safety alerts—even if you don't make your camera's content public. You can customize your area so you're only getting relevant alerts, especially if your neighborhood is more active on the app. Go to Menu > Neighbors > Settings > Customize Neighborhood to adjust the boundaries of your area.

id="turn-off-community-requests-from-law-enforcement">Turn off Community Requests from law enforcement

Community Requests is a Neighbors feature through which law enforcement can ask users to share video from their Ring devices. While Ring says that footage isn't shared automatically, and law enforcement doesn't have access to live feeds, many users still have privacy concerns related to this type of collaboration. (Note that Ring also had a short-lived partnership with Flock Safety , which would have made it easier for law enforcement agencies to request Ring camera footage using Flock's software.)

You can simply ignore Community Requests in your Neighbors feed, or you can turn these requests off entirely under Neighborhood Settings > Feed Settings . Deselect Community Requests and hit Apply .

id="opt-out-of-search-partys-surveillance">Opt out of Search Party's surveillance

Ring launched its Search Party feature in a Super Bowl ad earlier this year, ostensibly to help users find lost dogs in their neighborhood. Search Party uses AI to identify pets in your Ring's field of vision and pools the footage with other cameras. Obviously, this functionality comes with significant privacy concerns , not least of which is whether and how your footage could be shared with law enforcement to surveil people rather than pets. You can disable Search Party entirely under Control Center > Search Party . Choose the blue Pet icon next to each camera.

id="disconnect-from-amazon-sidewalks-wireless-network">Disconnect from Amazon Sidewalk's wireless network

Amazon Sidewalk uses your Ring device—and others in your neighborhood—to create a mesh network so said devices stay connected to the internet even if your wifi is weak or goes down. Amazon says that this feature provides security (because you'll still get important alerts) and extends the range for other devices like smart lights, smart locks, and pet locators. But you may not want to use your bandwidth for this purpose nor introduce potential privacy concerns to your home network. You can disable Amazon Sidewalk in the Control Center on your Ring app.

id="disable-third-party-provider-sharing">Disable third-

}

{

€9 Ticket

Dennis Felsing

As a relief for rising energy costs and inflation Germany is offering nearly-free public transportation for the months of June, July and August. For just €9 per month you can take any kind of regional public transportation in all of Germany. That's basically everything except the IC/ICE for long distance. This means that a monthly ticket is now as cheap as what I would normally pay for a single route. In this post I want to talk about my experiences with this €9 ticket .

Update 2022-09-01: The €9 ticket ended yesterday without a similar replacement system in place. So many people are again back to paying about €100 per month for much smaller regional tickets instead (see for example this German Reddit thread . Financially it makes more sense for me personally to go back to driving. (I hope the gasoline in the car is still good after three months.) My hope is that these three months will stay in people's minds in Germany and shape the future of public transportation to be cheaper and simpler:

Image via Reddit

id="train-stations">Train Stations

In the above graphic you can see the train stations in our area. Since we are situated in the Rhine-Neckar metropolitan region the railway network is pretty good here.

(map made with umap , routes made with GraphHopper)

The closest train stop from us is St. Ilgen-Sandhausen at 2 km distance. It offers direct connections to the cities of Heidelberg (160k inhabitants), Mannheim (309k), Karlsruhe (313k), Darmstadt (159k) and Frankfurt (753k). For other directions I prefer going to a different train station since switching times are either too long or there is too much risk of missing the connecting train. Only 89% of regional trains are on time in Germany, with some routes and stations having much worse rates. This is a far cry from Switzerland's 95% punctuality with an even stricter definition of 3 minutes versus 5 minutes delay in Germany.

Another closeby train station we commonly use is HD-Weststadt/Südstadt at 6 km distance, which is good to cycle and can also be reached via Tram. Occasionally I used Heidelberg Hauptbahnhof as well as Wiesloch-Walldorf, both at slightly under 10 km distance, which are a bit larger and thus are used as stops for Regional-Express trains. These go faster than the Regionalbahn and S-Bahn trains and thus also have fewer stops.

id="cycling-trips">Cycling Trips

The above map contains some of the cycling trips I took by combining them with the €9 ticket. (map made with umap , routes made with GraphHopper)

Last time I posted about cycling and work was about commuting by bicycle for a year . I kept this up and was super happy with this way of getting to SAP 's headquarter in Walldorf and the included exercise. After a while I moved together with my wife and my bike commute lengthened to 10 km each direction, which was still perfectly fine. Then Covid forced me in March of 2020 to work from home I switched to regular cycling trips in the afternoon since I started working early in the morning, in order to have meetings with colleagues in Korea and China. Two 50 km cycling trips per week got me to the same level as the daily commute used to do, and I have a nice route I regularly take, which covers a lot of forest and lake.

In January I switched to a new job at Yugabyte . As the database we develop is distributed, so are the people I work with. Since most meetings fall into US times, I am now working more during the afternoon and evening. This means my cycling trips moved to the morning instead.

Previously I had a job ticket for three years, but other than the €9 ticket it limits the range I can take public transportation in, so that many of my destinations would be barely outside of the covered range. When looking at a map of Germany with all the different transport associations, each with their own tickets, prices and rules, I feel reminded of historical maps of the Holy Roman Empire :

Good luck figuring out what ticket to get to combine with your job ticket. Luckily the €9 ticket changed this, at least temporarily, making life simpler.

I used the ticket during the working days to discover new cycling routes. By cycling one direction and taking a direct train connection back, I can cover nearly twice the distance as usual with the same starting and stopping point. With this method I took many routes I had never taken before, discovering beautiful cycling paths along the Neckar river for example. Luckily the trains are not too crowded during working days, especially outside of rush hour. The cycling usually takes around three times as long as the train ride back.

The image above shows an almost empty train during a working day.

Even with a single train connection things can go wrong though. One time the train in front of ours started burning and so we got delayed for an unknown time. But since I had my bicycle with me, that meant I could just cycle back the rest of the way and get some extra exercise for free. If I was really out of energy I could have waited in the train, but often busses are used as a replacement, and bicycles on busses don't fit well.

The weather was unusually dry the last months, the last few days had the first few drops of rain I remember since the beginning of June. This is of course bad for nature here, as I saw at the Rhine river, which is much lower than usual. For cycling it's quite good though, only got wet once.

}

{

A Rant on Usable Security

Jessie Frazelle

I recently gave a talk at DevOps Days (slides) and it had a pretty great response. I'm still pretty care-mad about the topics it covered so I figured I would turn some key points from it into a blog post.

The overall outline of the talk covered the past, present, and future of usable security. Let's start with the past.

id="the-past">The Past

A lot of the security tooling of the past (that we still use today) require users to jump through a lot of hoops or learn a hard to grok interface. One of the examples I used was GPG. Contrary to popular opinion, I actually don't find GPG entirely unusable. I obviously agree that it could be easier to use, rotate keys, revoke keys blah blah blah. While I find it not exactly terrible, I can see and completely understand why the majority of criticism I hear about GPG is that it is hard to use.

There is a point at which better security comes at the expense of convenience. This needs to stop happening. Stop compromising convenience for security. Instead find the right balance between the two. Doing this takes collaboration from both security engineers and software engineers.

Dave Cheney recently had a great tweet.

dir="ltr" lang="en">Why is all software shit? Today I discovered the @duosec API returns 200 even if someone denies the 2fa request.

— Dave Cheney (@davecheney) July 25, 2017

I love this tweet because it reeks of the stench that only security engineers built this API. Most software engineers I know would decide to use an HTTP status code... I mean that's what they are for. ;)

When you combine expertise in different areas you build better products. This is not rocket science. However egos tend to get in the way as well as biases towards people who know and like the same things you do. I assure you, though, when security and software engineers work together truly usable security will be the outcome.

id="the-present">The Present

A lot of the content for this portion of the talk focused on how containers make securing your infrastructure easier. I will touch on some of that but if you wish to know more you should checkout the slides or some of my other blog posts on container security.

Two key features in Docker are the default AppArmor and Seccomp profiles. AppArmor and Seccomp are Linux Security Modules that are not exactly usable by someone who is unfamiliar with either.

AppArmor can control and audit various process actions such as file (read, write, execute, etc) and system functions (mount, network tcp, etc). It has its own meta language, so to speak, and I actually have a repo that changes the docs for it to more a readable format via a cron job: github.com/jessfraz/apparmor-docs . The default profile for AppArmor does super sane things like preventing writing to `/proc/{num}` , `/proc/sys` , `/sys` and preventing mount to name a few.

Syscall filters allow an application to define what syscalls it allows or denies. The default in Docker is a whitelist that I initially wrote. Some of the key things it blocks are:

`add_key` , `keyctl` , `request_key` : Prevent containers from using the kernel keyring, which is not namespaced. I wrote a blog post on Two Objects not Namespaced by the Linux Kernel and the keyring was one I mentioned.

`clone` , `unshare` : Deny cloning new namespaces. Also gated by `CAP_SYS_ADMIN` for `CLONE_*` flags, except `CLONE_USERSNS` . I specifically wanted to block cloning new user namespaces inside containers because they are notorious for being points of entry for kernel bugs.

There also is an entire document that I started in the docker repo that outlines what we block and why.

Having written the default seccomp profile for Docker I am pretty familiar with

how hard this would be for other people. It requires a deep knowledge of the

application being contained and the syscalls it requires. This was also a quite

terrifying feature to add to Docker. When I added it, Docker was already very

popular and if anything would break in a big way it would be on the front page

of hacker news and all the maintainers would have a very bad day. So turning

on something that will EPERM by default if we left out any important syscall

is terrifying. I had stress nightmares for weeks. In the end everything went

much smoother than I feared but that was also after HEAVY HEAVY testing. Luckily

I run super obscure things in containers so I even caught that we left out send

and recv right before the release by running Skype (a 32 bit application) in

a container.

}

{

Closing Iterables is a Leaky Abstraction

Reginald Braithwaite · Reginald Braithwaite (raganwald)

id="iterators-and-iterables-a-quick-recapitulation">iterators and iterables, a quick recapitulation

In JavaScript, iterators and iterables provide an abstract interface for sequentially accessing values, such as we might find in collections like arrays or priority queues. 1

An iterator is an object with a `.next()` method. When you call it, you get a Plain Old JavaScript Object (or “POJO”) that has a `done` property. If the value of `done` is `false`, you are also given a `value` property that represents, well, a value. If the value of `done` is `true`, you may or may not be given a `value` property.

Iterators are stateful by design: Repeatedly invoking the `.next()` method usually results in a series of values until `done` (although some iterators continue indefinitely).

Here’s an iterator that counts down:

```
class="highlight">const iCountdown = { value: 10, done: false, next() { this.done = this.done || this.value { done: false, value: 10 }
```

```
iCountdown.next() //=> { done: false, value: 9 }
```

```
iCountdown.next() //=> { done: false, value: 8 }
```

```
// ...
```

```
iCountdown.next() //=> { done: false, value: 1 }
```

```
iCountdown.next() //=> { done: true }
```

An iterable is an object with a `[Symbol.iterator]` method. When invoked, `[Symbol.iterator]()` returns an iterator. Semantically, the iterator returned by `[Symbol.iterator]()` represents an iteration over the values associated with the iterable collection.

For example:

```
class="highlight">const countdown = { [Symbol.iterator]() { const iterator = { value: 10, done: false, next() { this.done = this.done || this.value
```

Or destructure them:

```
ten //=> 10 nine //=> 9 eight //=> 8 rest //=> [7, 6, 5, 4, 3, 2, 1]
```

And now, let’s get started. We’ll begin with a simple problem: How do we iterate over the lines of a text file?

id="reading-lines-from-a-file">reading lines from a file

We wish to create an iterable that successively yields the lines from a text file. Presuming we have some kind of library for opening, reading from, and closing files, we might write something a little like this:

```
class="highlight">function lines ( path ) { return { [Symbol.iterator]() { return { done: false, fileDescriptor: File.open ( path ), next() { if ( this.done ) return { done: true }; const line = this.fileDescriptor.readLine ();
```

```
this.done = line == null ;
```

```
if ( this.done ) { fileDescriptor.close (); return { done: true }; } else { return { done: false, value: line }; } } } }
```

Whenever we want to iterate over all the lines of a file, we call our function, e.g. `lines('./README.md')`, and we get an iterable for the lines in the file.

When we invoke `[Symbol.iterator]()` on our iterable, we get an iterator that opens the file, reads the file line by line when we call `.next()`, and then closes the file when there are no more lines to be read.

So we could output all the lines containing a particular word like this:

The expression `lines('./README.md')` would create a new iterator with an open file, we’d iterate over each line, and eventually we’d run out of lines, close the file, and exit the loop.

What if we only want to find the first line with a particular word in it?

```
class="highlight">for ( const line of lines ( './README.md' ) ) { if ( line.match ( /raganwald/ ) ) { console.log ( line ); break ; }
```

Now we have a problem. How are we going to close the file? The only way it will exhaust the iterations and invoke `this.fileDescriptor.close()` is if the file doesn’t contain `raganwald`. If the file does contain `raganwald`, our program will happily carry on while leaving the file open.

This is not good. And it’s not the only case. We might write iterators that act as coroutines, communicating with other processes over ports. Once again, we’d want to explicitly close the port when we are done with the iterator. We don’t want to just garbage-collect the memory we’re using.

What we need is some way to explicitly “close” iterators, and then each iterator could dispose of any resources it is holding. Then we could exercise a little caution, and explicitly close every iterator when we were done with them. We wouldn’t need to know whether the iterator was holding on to an open file or socket or whatever, the iterator would deal with that.

Fortunately, there is a mechanism for closing iterators, and it was designed for the express purpose of dealing with iterators that must hold onto some kind of resource like a file descriptor, an open port, a tremendous amount of memory, anything at all, really.

Iterables that need to dispose of resources introduce a problem. To solve it, the language introduced a mechanism for

}

standard-article(title: [Reflecting on What Matters After a Terminal Cancer Diagnosis], author: [Harvard Business Review], source-name: [Harvard Business Review], images: (), paragraphs: ([How does someone who's been told he will die much sooner than expected find contentment in the time he has left? As a former therapist, cofounder of the Deeper Coaching Institute, and business book author, Mark Goulston has spent his entire career trying to help others manage their emotions, improve their communication, and find the right balance between the personal and the professional. Faced with his own cancer diagnosis, he's been reflecting on lessons learned in his own life, things he and clients wish they'd done differently, and how to both prepare for a "good" death and leave a meaningful legacy. He shares his newfound perspective and his advice for early, mid- and late-career leaders.],), insert-map: (:), word-count: 115, edited-for-length: false, debug-mode: false,)

standard-article(title: [The Day I Leave the Tech Industry], author: [Jessie Frazelle], source-name: [Jessie Frazelle], images: (), paragraphs: ([I was inspired last night by Cate Huston's post, The Day I Leave the Tech Industry . I decided to write my own, except I'm not as eloquent a writer as Cate so before I go any further please, please, please read her post and not mine.], [Mine is going to be a bit different. Lately I've been thinking more and more about this. It seems imminent. I'm only 27 and let me repeat: it seems imminent.], [I'm going to tell you all the fantasy that plays in my brain for when this happens.], [The day I leave the tech industry will feel like a giant weight has finally been lifted. It will be freeing. There are a few scenarios I've played out for what I will do after.], [teach math in a third world country], [write a book], [play professional poker], [I could do all three. One thing is for sure though, the day I leave the tech industry will be the day I contribute my last piece of code to open source software.], [Today, I am not quite ready to give up this thing I have such a "hate/love" relationship with. Today, I want to get more women contributing so that maybe in the distant future we will feel welcome. Maybe we won't have to fight so hard just to be heard; to have our opinions matter.], [Today is not my last day in the tech industry. But it is comforting to me to plan out this very real future. I am not just "the container girl". I am a human being with feelings, a limit, and a future outside of tech.],), insert-map: (:), word-count: 274, edited-for-length: false, debug-mode: false,)

brief-group(([

Harvard Business Review, Harvard Business Review: Deborah Ancona and Kate Isaacs, researchers at MIT Sloan School of Management, say many companies struggle to be nimble with a command-and-control leadership culture. They studied Xerox's R&D outfit PARC and the materials science company W. L. Gore & Associates and found these highly innovative organizations have three kinds of leaders: entrepreneurial, enabling, and architecting ones. These roles work together to give direction and avoid creative chaos. Ancona and Isaacs are coauthors of the HBR article "Nimble Leadership."

], [

Harvard Business Review, Harvard Business Review: Pankaj Ghemawat, professor at NYU Stern and IESE business schools, debunks common misconceptions about the current state and extent of globalization. (Hint: the world is not nearly as globalized as people think.) He also discusses how popular reactions in Europe and the U. S. against globalization recently could affect the global economy, and how companies will need to adapt to the new reality. Ghemawat is the author of several books on globalization, including "World 3.0" and most recently "The Laws of Globalization and Business Applications."

], [

Supabase Blog, Supabase Blog: This documents our journey from SOC2 Type 1 to SOC2 Type2 and HIPAA compliance. You can start building healthcare apps on Supabase today.

], [

Supabase Blog, Supabase Blog: Open table formats like Apache Iceberg, Delta Lake, and Apache Hudi are transforming how developers manage large-scale data on object storage systems.

], [

Harvard Business Review, Harvard Business Review: Ellen Ernst Kossek, management professor at Purdue University, is researching how the pandemic is putting an enormous strain on working parents and the new challenge that poses for their managers. She shares how supervisors can offer much-needed consistency and predictability for working parents on their teams. She also outlines specific ways to give working parents more flexibility while still holding them accountable. Kossek is the coauthor, with Kelly Schwind Wilson and Lindsay Mechem Rosokha, of

the HBR article “What Working Parents Need from Their Managers.”

], [

Harvard Business Review, Harvard Business Review: Debbie Cohen and Kate Roeske-Zummer, cofounders of HumanityWorks, are sounding an alarm bell for employee retention. Record numbers of people are quitting their jobs due to burnout and better opportunities. Those resignations leave their former colleagues burdened with even more work and a sense of despair. Cohen and Roeske-Zummer argue that employers should re-recruit their existing employees and even think of them as customers. And the two consultants outline steps managers can take to openly appreciate those employees and keep a positive culture. Cohen and Roeske-Zummer wrote the HBR.org article “With So Many People Quitting, Don’t Overlook Those Who Stay.”

], [

Harvard Business Review, Harvard Business Review: Heidi Grant Halvorson, author of “No One Understands You and What to Do About It,” explains the science of perception.

], [

Harvard Business Review, Harvard Business Review: Muriel Wilkins, cofounder of the executive coaching firm Paravis Partners, says that starting a leadership role at a new company or via internal promotion is demanding. Doing so remotely during the Covid-19 pandemic is even more challenging. She says that new senior leaders must focus on two things: connectivity and credibility. And she explains how to build those attributes when much of the job is performed virtually. Wilkins is the host of the new HBR Presents podcast “Coaching Real Leaders.”

], [

Harvard Business Review, Harvard Business Review: John Kerry, former U. S. Secretary of State, shares management and leadership lessons from his long career in public service. He discusses how to win people over to your side, bounce back from defeats, and never give up on your long-term goals. He also calls on private sector CEOs to do more to solve social and political problems. Kerry’s new memoir is “Every Day Is Extra.”

],))

standard-article(title: [An analysis of memory bloat in Active Record 5.2], author: [@sam Sam Saffron], source-name: [Sam Saffron], images: (), paragraphs: ([Not impossible that AR has regressed here further, I have not been tracking, we still use this at Discourse:], [github.com], [discourse/

standard-article(title: [Red Flags You Won’t See on a CEO’s Resume], author: [Harvard Business Review], source-name: [Harvard Business Review], images: (), paragraphs: ([For a long time, we have believed that strong corporate governance is enough to prevent CEO malfeasance. However,

discourse/blob/main/lib/freedom_patches/fast_pluck.rb],
[# frozen_string_literal: true], [# Speeds up #pluck so its
about 2.2x faster, importantly makes pluck avoid creation
of a slew # of AR objects # # class ActiveRecord:: Rela-
tion # Note: In discourse, the following code is included
in lib/sql_builder.rb # # class RailsDateTimeDecoder =
"4.2.0" # @caster ||= ActiveRecord:: Type:: DateTime.new #
@caster.type_cast_from_database(string) # else # ActiveRe-
cord:: ConnectionAdapters:: Column.string_to_time string
end # end # end #], [This file has been truncated. show
original], [style="clear: both;">], [I wish we did not have to,
but I am not sure how to land this in Rails.], [Another change
that has happened since is that we released:], [GitHub],
[GitHub - discourse/mini_sql: a minimal, fast, safe sql ex-
ecutor], [a minimal, fast, safe sql executor. Contribute to
discourse/mini_sql development by creating an account on
GitHub.], [style="clear: both;">], [Which is the piece we use
for any performance sensitive work and totally outperforms
almost anything you can throw at it.], [Maybe run the bench
on some earlier versions 7/6/5 and see what happens?],),
insert-map: (:), word-count: 190, edited-for-length: false,
debug-mode: false,)

new research shows that the lifestyle behaviors of execu-
tives can spell trouble for companies, regardless of the
guardrails in place. Aiysha Dey, an associate professor at
Harvard Business School, has investigated executives’ past
criminal records and the cost of their homes and automom-
biles. Her research has linked an individual’s materialism
and propensity for rule breaking to fraud, insider trading,
and risky business activities. She says that boards and other
hiring bodies should pay more attention to personal behav-
ior when picking organizational leaders. Dey wrote the HBR
article “When Hiring CEOs, Focus on Character.”],), insert-
map: (:), word-count: 111, edited-for-length: false, debug-
mode: false,)

The Silicon Wire

Vol. 1, No. 005 · 2026-03-30

Curated and composed automatically.