

The Clear Standard

VOL. 1, NO. 009

2026-03-30

FRONT PAGE GELATO QUEEN

Kirkus Reviews

Eleven-year-old Lizabeth “Liza” Gordon is used to moving house; as her grandmother says, her dad “was bitten by the gobug when he was young.” This time, Liza’s dad aims to start a family gelato business in Greenblossom, South Carolina, that he’ll call Gordon’s Gelato. Liza and her older brothers, Pete and Brad, are going to help run the shop. Liza doesn’t want to leave her friends, but when they prove to not care about her, she says good riddance to the Richmond, Virginia, area and heads to her new home. It’s worse than she imagined: The new residence is a “rundown, bi-level style house,” her room is “puss yellow with prune-purple trim,” and she’s starting school with only six weeks left before summer vacation.

Stepping into her new life, she has to figure out how to work in a gelato shop, determine whether her new friends actually like her, and plan for her 12th birthday. Liza’s dad becomes a contestant on a reality show, competing for a grand prize of \$1 million, which makes things a lot more complicated for Liza, but if her dad wins, her whole life could change. The novel’s fun and engaging lead character makes the book a captivating read. The supporting characters, from Liza’s eccentric dad to the boy who helps out in the gelato shop, feel well rounded and unique. Liza is easy to root for, especially when everything keeps going wrong for her. Her creative and disgusting gelato flavors, nicknamed “Gordon’s Gaggers,” are a funny element that adds to the overall charm of the story.

FEATURES

Dems Quietly Start Their Next Big Health Care Effort

Jonathan Cohn · *The Bulwark*

(Composite by Hannah Yoest / Photos: Shutterstock)

THERE ARE SIGNS that the debate about health care in America is about to get out of the rut it's been in for about fifteen years—and that Democrats are preparing for the moment when it does.

Ever since 2010, the most high-profile fights in Washington have been about the Affordable Care Act and Medicaid. Mainly that's because Donald Trump and the Republicans keep attacking those programs—as they did last year when they enacted the largest-ever cuts to Medicaid, then refused to extend lapsing “Obamacare” subsidies that had helped millions to get coverage, and reduced premiums for many millions more.

Democrats are determined to reverse those two steps, somehow, and you can expect them to make that a rallying cry in their campaigns for November's midterm elections. But at least some Democrats don't want to stop there. On March 19, a dozen of the party's senators released an open letter announcing their intention to develop policies that would address a broader topic: The underlying increase in health care costs that is affecting everybody, not just people who are uninsured, on Medicaid, or buying coverage at HealthCare.gov.

The roughly 170 million Americans who get coverage through their employers are now paying (directly and indirectly) an estimated \$27,000 a year on average for a family policy. “The American people need relief from rising premiums and deductibles that are forcing families into financial ruin,” the Senate Democrats wrote in their letter. “They also want an insurance system that doesn't require them to jump through hoops and hack through red tape every time they need care.”

That may sound like a bunch of frothy boilerplate, given that the letter contained no specifics. But it's not just these Democratic lawmakers who say it's time to have a broader conversation, one that goes beyond undoing what Trump and the Republicans have just done. You hear the same thing from prominent analysts and advocates like Anthony Wright, president of the pro-coverage, pro-consumer organization FamiliesUSA.

“I do think people recognize that, as we wage the fight to defend coverage and consumer protection and specific communities under attack, that we don't fall into a trap of defending the status quo that people thought rightly was broken,” Wright told me in an interview. “We need to show that we have a plan, not just to repeal bad stuff, or even to rebuild—but to reimagine what the health system should look like.”

That kind of reimagining can't happen right away. Trump and the Republicans seem incapable of putting forward serious reform proposals, unless they involve hacking away

at insurance coverage for people who need it. And the first chance Democrats might have to govern with a trifecta is nearly three years away. But it's with an eye to that possibility that Democrats and their allies are starting to plan now—to make sure they are “prepared to take action on these issues the next time Democrats have an opportunity,” as the Senate Democrats put it in their letter.

And there's an unmistakable parallel here, to a politically similar time when Democrats and their allies started laying the groundwork for future legislation. “This moment feels a bit like twenty years ago,” Larry Levitt, executive vice president for health policy at KFF, told me in an interview, “when groups of policy experts, advocates, and politicians started to talk about health care ideas that ultimately coalesced into the passage of the Affordable Care Act.”

But the challenge is different this time, and in some ways more difficult. Reducing health care costs inevitably involves reducing the flow of money into somebody's pockets, which just as inevitably angers powerful constituencies and industry groups. Democrats aren't even close to having a consensus on what to do. And 2029 is a lot closer than it might seem.

ONE PERSON WHO KNOWS THIS all too well is Ron Wyden, the senator who organized the open letter and who hopes to lead discussions inside his caucus. He's been focusing on health care ever since the 1970s, when as a freshly minted college graduate he helped to establish Oregon's chapter of the Gray Panthers, a grassroots organization to defend the interests of seniors.

The unfailingly earnest Wyden will bring this up anytime you talk to him about health care. It's part of his wonky charm, if you've gotten to know him as some of us have, and a point of information that is genuinely relevant. Helping elderly Americans get health care really has been a signature cause, going all the way back to 1981 when he first got to Congress.

In fact, when I spoke to Wyden last week, he said he was still angry over testimony that Mehmet Oz gave a year ago, during hearings for his confirmation as administrator of Medicare and Medicaid. The celebrity doctor had refused to stand by a Biden-era rule that raised the minimum staffing standards for nursing homes—requiring more personnel and a nurse on duty 24/7 by 2029—instead arguing that technology and telemedicine could take the place of extra staff.

“Are you kidding me,” Wyden said in our interview, “an 85-year-old woman trying to go to the bathroom in a dark bedroom is somehow going to be able to be safe and get there by an algorithm?” 1

But Wyden has also put in the time to think about and propose policies for non-elderly Americans. In late 2006, shortly after Democrats won control of the House and Senate in that year's midterm election, he unveiled what he called the “Healthy Americans Act”—a detailed legislative proposal, complete with a financing scheme and indepen-

I am definitely missing the pre-AI writing era

N. Cailie · LessWrong

Yesterday, I wrote my first technical draft on what I was working on with the goal to share it publicly on here (well using an account dedicated to technical post), and did not realize how wanting to sound perfect actually steal the “my voice” in the paper. Although 80 % of the content was my own writing, the fact that it was run in a LLM engine for grammar and vocabulary cross-check, made it failed the “probable written by AI” metric; and it was rejected. I am not complaining, well, it was kinda my fault after all cause the instructions did specify the no-use of LLM, when sending first post. And the rejection made me think: why did I even need the AI to validate what I worked on?

The scary part, is I love writing, I have been writing blogs and all sort of things since I can't even remember. Have some think pieces on popular sites, and other personal blog that I used to share. And you know, before 2023, my writing was so good that I rarely had to have a second read or correction to make sure the flow stays consistent, since english is my 4th language. I always sought external review or check with tools like : Grammarly or Quilbot, for mistakes or grammars misphrasing. But now, this “superskill” has faded away, like seriously. I can't even write a 1000 words think piece, without needing or wanting to know what AI thinks or how it could have phrase it better, even for my emails. I can't write poems as before, without sounding generic, and this is mainly because I have now trained my brain to rely on these automated tools that It can not be creative anymore or think of itself. My writing has deteriorated, I attempted to write a slam-poem, recently for a competition I wanted to partake in, and upon the completion when I read it, my reaction was : Who is this? Who wrote this? This is bad !

Before, I would sit and write a piece and I kid you not, the first draft was always perfect that I rarely had to write, and my best poem that got me to some “art”ies festival, was just written while I was sitting outside my hostel, with a piece of paper and a pen, and a thought that was bothering me. But now, I am slowly realizing, it has been changing in negative way.

Sorry that this sounds like a rant, well, it is one ! But like, my point was to share it out loud, and kinda know if I am the only person experiencing this or if there are others? And for people who successfully taken back their creative writing skills, how did you do it?

This post, is written without any tools assistance I just wrote what my brain is instructing to type (might not reread it before posting). Obviously, you would see some mistakes, some paraphrasing issue, some article misuse; I personally do not see them but, hey, I am not expert in literature. So bear with me, also, I think that is the beauty of writing, the raw , unedited emotions of the person behind every words either for entertainment or educational purposes, is what makes it special.

So, thank you “LessWrong” moderators, for rejecting my first draft, it is kinda a wake up call that made me realize, maybe I just need to be focusing on my own voice, and there is nothing wrong into sounding “more or less wrong”. Yeah, I know, people advocate into embracing AI as it would be/ is in every aspect of our life, but maybe it is good that most of us do not let this technology rephrase our own thought. Cause those words, that the AI told you: sound wrong or aggressive or less formal, are the ones that carry the “emotions”, you ought to share.

Have a productive week, dear reader !

P. S: If you see this! I am celebrating. It means, I passed the text. Yay!

Gina Wiley: My Journey So Far as a Senior Engineering Manager

Deliveroo Engineering

Could you introduce yourself and your role at Deliveroo?

I'm Gina, I'm a Senior Engineering Manager at Deliveroo. I lead our Consumer Experience group in our New Verticals engineering team, and I've been at Deliveroo for five years.

Could you walk us through your path to engineering management and what brought you to Deliveroo?

I studied information systems at university and developed an interest in human-computer interaction. This field emphasises understanding users and translating business processes into technical requirements. I moved to San Francisco for a tech consulting role, where I worked on a large-scale, multi-country project replacing legacy systems with connected Software as a Service (SaaS) products. This cross-functional experience shaped my understanding of end-to-end business process transformation through technology.

In 2017, I joined DoorDash's Business Applications team, working on partner onboarding and support processes, which introduced me to logistics and delivery tech during a pivotal time for the industry.

In 2020, I moved to London to work at Deliveroo, initially supporting restaurant partner onboarding in the Restaurants engineering group. I now work in the New Verticals engineering group, focusing on launching grocery and retail propositions. We've recently introduced a lot of new features to support this growing space; so continuing to innovate and deliver at pace to keep up with the changing market has been an exciting challenge!

Reflecting on the role of engineering manager here, what do you think is the most rewarding part of this role?

I enjoy seeing engineers on the team take on new challenges and push themselves out of their comfort zone to reach their full potential, so I'll make sure that they get access to the right resources and opportunities so that they can continue to learn and grow.

It's important to trust your people to reach their growth goals with the right opportunities. I believe (and have seen throughout my career) that when you give people more ownership and responsibility, they will go above and beyond to achieve a positive impact.

What would you say about career progression at Deliveroo?

In progressing your career, it's important to know what is expected and what it means to reach that next level. I value that at Deliveroo we have a clear 'levels and expectations' framework and resources to guide career conversations. Our framework sets out very clearly what is expected at each level, and anyone who joins will review this framework with their manager to start the conversation of what opportunities there are to progress.

At Deliveroo I've found that there are so many opportunities to increase your scope of influence and responsibility. We have so many fantastic impactful projects people can get involved in, and incredible team members, which means there is no shortage of opportunity. What's really important is leading by example; having a growth mindset to continuously work on improvement and seeking out those opportunities to grow outside of your current areas of expertise.

How have you been supported in your own growth journey?

I've had amazing support by all my managers, mentors, peers and team members I've worked with over the last 5 years. It wasn't easy joining in 2020 and moving to a new country – it was a tricky time for many, but I always had a support network to lean on here. I've had very supportive managers who have advocated for me, who have provided me opportunities to grow and challenge myself, and I have benefitted from Deliveroo's formal mentoring program to get extra support on my growth areas.

What sets Deliveroo apart from any other tech companies you've worked for?

Working at Deliveroo has provided me with immense opportunities to make a significant impact. I've been most impressed by the people I work alongside, and it's clear that everyone has a drive to make change and do so in a way that we learn from each other and pivot quickly if needed. People here are passionate about the work they are doing; this passion creates a fast-paced environment where I feel I'm constantly learning and being challenged by those around me to consider how we can get incrementally better day-to-day.

What would be your one piece of advice for any new engineers or managers looking to join us?

No matter how big or small the project, feature, or task – give it your all. Hold a high bar for yourself and make that a regular practice, and that will show in the end result. If we all have that collective sense of ownership and accountability for the work that we're doing, then we'll be able to go faster as a team – it becomes a flywheel effect. Doing this also means you are setting a good example for your people and encouraging others to do the same, which creates highly effective, motivated teams that can produce incredible results.

We're hiring! If you are interested in roles in Tech, please visit deliveroo.com/careers.

I believe (and have seen throughout my career) that when you give people more ownership and responsibility, they will go above and beyond to achieve a positive impact.

Deliveroo Engineering

Generative AI in the Real World: Sharon Zhou on Post-Training

Ben Lorica and Sharon Zhou · O'Reilly Radar

Post-training gets your model to behave the way you want it to. As AMD VP of AI Sharon Zhou explains to Ben on this episode, the frontier labs are convinced, but the average developer is still figuring out how post-training works under the hood and why they should care. In their focused discussion, Sharon and Ben get into the process and trade-offs, techniques like supervised fine-tuning, reinforcement learning, in-context learning, and RAG, and why we still need post-training in the age of agents. (It's how to get the agent to actually work.) Check it out.

About the Generative AI in the Real World podcast: In 2023, ChatGPT put AI on everyone's agenda. In 2026, the challenge will be turning those agendas into reality. In Generative AI in the Real World, Ben Lorica interviews leaders who are building with AI. Learn from their experience to help put AI to work in your enterprise.

This transcript was created with the help of AI and has been lightly edited for clarity.

00.00 Today we have a VP of AI at AMD and old friend, Sharon Zhou. And we're going to talk about post-training mainly. But obviously we'll cover other topics of interest in AI. So Sharon, welcome to the podcast.

00.17 Thanks so much for having me, Ben.

00.19 All right. So post-training. . . For our listeners, let's start at the very basics here. Give us your one- to four-sentence definition of what post-training is even at a high level?

00.35 Yeah, at a high level, post-training is a type of training of a language model that gets it to behave in the way that you want it to. For example, getting the model to chat, like the chat in ChatGPT was done by post-training.

So basically teaching the model to not just have a huge amount of knowledge but actually be able to have a dialogue with you, for it to use tools, hit APIs, use reasoning and think through things step-by-step before giving an answer—a more accurate answer, hopefully. So post-training really makes the models usable. And not just a piece of raw intelligence, but more, I would say, usable intelligence and practical intelligence.

01.14 So we're two or three years into this generative AI era. Do you think at this point, Sharon, you still need to convince people that they should do post-training, or that's done; they're already convinced?

01.31 Oh, they're already convinced because I think the biggest shift in generative AI was caused by post-training ChatGPT. The reason why ChatGPT was amazing was actually not because of pretraining or getting all that information into ChatGPT. It was about making it usable so that you could actually chat with it, right?

So the frontier labs are doing a ton of post-training. Now, in terms of convincing, I would say that for the frontier labs, the new labs, they don't need any convincing for post-training. But I think for the average developer, there is, you know, something to think about on post-training. There are trade-offs, right? So I think it's really important to learn about the process because then you can actually understand where the future is going with these frontier models.

02.15 But I think there is a question of how much you should do on your own, versus, us[ing] the existing tools that are out there.

02.23 So by convincing, I mean not the frontier labs or even the tech-forward companies but your mom and pop. . . Not mom and pop. . . I guess your regular enterprise, right?

At this point, I'm assuming they already know that the models are great, but they may not be quite usable off the shelf for their very specific enterprise application or workflow. So is that really what's driving the interest right now—that people are actually trying to use these models off the shelf, and they can't make them work off the shelf?

03.04 Well, I was hoping to be able to talk about my neighborhood pizza store post-training. But I think, actually, for your average enterprise, my recommendation is less so trying to do a lot of the post-training on your own—because there's a lot of infrastructure work to do at scale to run on a ton of GPUs, for example, in a very stable way, and to be able to iterate very effectively.

I think it's important to learn about this process, however, because I think there are a lot of ways to influence post-training so that your end objective can happen in these frontier models or inside of an open model, for example, to work with people who have that infrastructure set up. So some examples could include: You could design your own RL environment, and what that is is a little sandbox environment for the model to go learn a new type of skill—for example, learning to code. This is how the model learns to code or learns math, for example. And it's a little environment that you're able to set up and design. And then you can give that to the different model providers or, for example, APIs can help you with post-training these models. And I think that's really valuable because that gets the capabilities into the model that you want, that you care about at the end of the day.

04.19 So a few years ago, there was this general excitement about supervised fine-tuning. And then suddenly there were all these services that made it dead simple. All you had to do is come up with labeled examples. Granted, that that can get tedious, right? But once you do that, you upload your labeled examples, go out to lunch, come back, you have an endpoint that's fine-trained, fine-tuned. So what happened to that? Is that something that people ended up continuing down that path, or are they abandoning it, or are they still using it but with other things?

Spotting and Avoiding ROT in Your Agentic AI

Q McCallum · O'Reilly Radar

The following article originally appeared on Q McCallum's blog and is being republished here with the author's permission.

Generative AI agents and rogue traders pose similar insider threats to their employers.

Specifically, we can expect companies to deploy agentic AI with broad reach and insufficient oversight. That creates the conditions for a particular flavor of long-running problem, which in turn creates a novel risk exposure for both the companies in question and for anyone doing business with them. The bot and the rogue trader are able to inflict sizable, sometimes existential, damage to the firms that employ them.

The key difference is the scope: Rogue traders operate in investment banks, while agentic AI will be deployed to a wider array of companies and industry verticals. Agentic AI may therefore create a greater number of problems than rogue traders and put a greater amount of capital at risk.

I'm naming this risk exposure ROT—Rogue Operator Threat—and this document is a brief explainer on what it is and how to address it.

(I almost called it RAT, with the A for “agentic,” but then realized that it would apply to any kind of automated system. So I broadened the scope to “operator.”)

To set the stage, let's take a trip to the trading floor:

Understanding the rogue trader

Rogue trader scandals follow the same storyline:

A trader accrues losses due to bad trades.

They hide those losses while placing new trades in an attempt to recover.

The new trades also lose money, digging a deeper hole.

Repeat.

This cycle continues until they're caught, at which point the bank is sitting on a large loss (sometimes into the billions of dollars) and the trader faces legal repercussions.

The story of Barings Bank offers a concrete example. Trader Nick Leeson had been logging fraudulent trades, over a stretch of three years, in an attempt to cover his mounting losses. This only came to light when the Kobe earthquake shifted markets against his most recent positions and the losses were no longer possible to hide. Leeson's £800M (\$1.3B) hole drove Barings to bankruptcy just three days later.

This is when you'll ask: How could a professional trading operation let so many bad trades slip through undetected? How could a trader falsify records? Aren't trading floors high-tech operations, full of electronic audit trails?

And the answer is: It's complicated.

Trading operations do keep records, yes. But no system is perfect. Each time a rogue trading scandal comes to light, it turns out that there were loopholes in risk controls. A sufficiently motivated trader—especially one desperate to hide their mistakes—found and exploited these loopholes, continuing their losing streak in plain sight until they could bring in real money to backfill the fake records.

That “until” never happened, though. Which is why their employers then faced financial, reputational, and sometimes legal troubles.

The AI agent's ROT threat

Similar to a trader, an AI agent operates on behalf of its parent business and is given room to operate independently so it can accomplish its tasks.

The risk is that, in the rush to deploy agentic AI, these companies will likely grant the bots more leeway than is necessary. We've already seen cases in which bots have been able to delete emails and wipe a production database. And there are no doubt other stories that haven't made it into the news.

Those issues were at least caught in real time. Companies facing ROT are exposed to additional longer-running problems in which the bot is able to accrue losses or inflict greater damage over an extended period. In those cases the problems will only be uncovered by accident and/or when it's too late.

Consider, for example, an agent that creates false data records to reflect (nonexistent) sales orders. It's possible for this to run until some external event, such as investor due diligence or a budget review, forces someone to double-check those records against reality.

Avoiding ROT: Mitigating the threat

How can you narrow your downside risk exposure to ROT? Preventative measures are key. Strong risk controls, narrow scope of authority, and monitoring can catch rogue operator problems long before they've metastasized into an existential threat.

In light of rogue trader scandals, trading shops have been known to tighten risk controls and also separate duties to create a system of checks and balances. (This inhibits traders from logging their own fake trades.) Companies also require traders to take time off, as fraudulent activity may surface when the perpetrator isn't around every day to keep the system running.

Shuffling and Whittling and Experimenting: Austin Keeling and Lam T. Nguyen on Editing the Told-on-Screens Film “Mercy”

Nell Minow · Roger Ebert (RogerEbert.com)

In “Mercy,” Chris Pratt plays a police detective accused of murdering his wife. It takes place in the future, when the judicial system is turned over to an AI “judge” played by Rebecca Ferguson. Until the final third of the film, everything takes place in a cyber-courtroom. He is trapped in a chair and has just 90 minutes to defend himself. But he has access to all the records, all the surveillance, and all the witnesses he wants. That means it all plays out on a gigantic screen. And that means the film’s editors, Austin Keeling and Lam T. Nguyen, had to make sure that the audience could try to solve the mystery by observing a lot of information — files, security camera footage, interviews — on that screen.

In an interview with RogerEbert.com, Keeling and Nguyen talk about that process, as well as finding ways to make audiences process a lot of information without getting lost in data.

This interview has been edited for clarity.

Everyone looks at screens all day. How do you overcome the audience’s screen fatigue and make looking at screens so vivid and involving?

AUSTIN KEELING: We ARE all looking at screens all the time, so keeping it fresh and interesting was definitely one of the biggest challenges. We got lucky that the script is about a murder investigation, so everything shown on the screens is directly linked to the central mystery. We made sure to justify the existence of each screen, only including it if it was connected to the detective work that Chris Pratt’s character engages in throughout the film. This draws the audience in, as each new screen presents a new clue or piece of information, allowing them to sleuth right alongside the main character.

LAM T. NGUYEN: Our objective was to create an immersive experience for our audience, and what sets this format apart is the incorporation of 3D elements in the screens. This unique hybrid format combines traditional filming techniques, screen-life elements, and 3D effects. During the editing process, we developed a digital camera that applied a blur (rack focus) effect and pan moves to our POV shots. This technique allowed us to establish a cinematic visual language that made the film entertaining and engaging, enabling our audience to follow the story effectively.

How do you direct our attention to what you want us to see and sneak information in the edges of what we see so we need to watch it again?

AK: Another huge challenge was juggling the massive quantity of assets in this film, both footage and graphics. We built the film by assembling “wide shots” of the Mercy courtroom and populating them with all the necessary assets for a given scene (footage of Judge Maddox, the courtroom back-

ground, video calls, security footage, websites, emails, etc.). We meticulously placed and animated each of these assets to exist in a master wide, and then created a “digital camera” to act as Chris’s POV. With this, we could zoom in and use keyframes to animate the camera to “look around” at various points in the room. This allowed us to focus the camera (and the audience’s attention) on key details throughout the Mercy courtroom while still retaining the extensive supplementary content at the edges.

LN: While editing this with our director, Timur Bekmambetov, we discussed that any important information should be placed in the center of the frame. We made a conscious effort to center everything important and then populate the surrounding screens afterward. We relied on our characters’ emotions to determine how quickly or slowly we looked at the screens. Since multiple screens are visible at any given time, everything has to be consistent and realistic to match the primary screen we are looking at. So there are definitely some easter eggs in there.

How do you balance what seems familiar or at least recognizable with some futuristic graphics that still have to be familiar enough for us to believe and process them?

LN: During pre-production, we explored all visual aspects of the film with Timur and the VFX team. There was a point at which things looked super futuristic and cool. However, Timur insisted on pulling that back and keeping it closer to our reality. Since the story takes place in 2029-2030, we wanted to advance the technology without straying too far from what we’re used to, so it’s more relatable for our audience when they experience the film. When I first met Austin, I brought up the Apple Vision Pro (since it was a new release at the time), and we used it as a baseline to develop the film’s visual language.

AK: We wanted to make sure that the Mercy chamber was not too far-fetched or futuristic (the film takes place only a few years in the future, after all). We played around with a bunch of different designs and animations, but ultimately the collaboration with Timur and the VFX team led us to a simple, unadorned style for most of the floating screens. Lam and I looked at the Apple Vision Pro early on as a jumping-off point for organizing and animating screens in 3D space. And all of the websites and video call formats are based on existing references. Essentially, we tried to keep a level of familiarity in the UI and functionality of the Mercy system so that audiences would believe in the scenario and ease into the chaos and tension of this world.

Chris Pratt stars as Chris Raven in MERCY, from Amazon MGM Studios. Photo credit: Justin Lubin © 2025 Amazon Content Services LLC. All Rights Reserved.

What do your home screens look like? Cluttered or very organized?

AK: I keep my home screen very clean and organized! I actually only have three folders on my desktop and nothing

Keep Deterministic Work Deterministic

Andrew Stellman · O'Reilly Radar

This is the second article in a series on agentic engineering and AI-driven development. Read part one here , and look for the next article on April 2 on O'Reilly Radar.

The first 90 percent of the code accounts for the first 90 percent of the development time. The remaining 10 percent of the code accounts for the other 90 percent of the development time. — Tom Cargill, Bell Labs

One of the experiments I've been running as part of my work on agentic engineering and AI-driven development is a blackjack simulation where an LLM plays hundreds of hands against blackjack strategies written in plain English. The AI uses those strategy descriptions to decide how to make hit/stand/double-down decisions for each hand, while deterministic code deals the cards, checks the math, and verifies that the rules were followed correctly.

Early runs of my simulation had a 37% pass rate. The LLM would add up card totals wrong, skip the dealer's turn entirely, or ignore the strategy it was supposed to follow. The big problem was that these errors compounded: If the model miscounted the player's total on the third card, every decision after that was based on a wrong number, so the whole hand was garbage even if the rest of the logic was fine.

There's a useful way to think about reliability problems like that: the March of Nines . Getting an LLM-based system to 90% reliability is the first nine, and it's the "easy" one. Getting from 90% to 99% takes roughly the same amount of engineering effort. So does getting from 99% to 99.9%. Each nine costs about as much as the last, and you never stop marching. Andrej Karpathy coined the term from his experience building self-driving systems at Tesla, where they spent years earning two or three nines and still had more to go.

Here's a small exercise that shows how that kind of failure compounding works. Open any AI chatbot running an early 2026 model (I used ChatGPT 5.3 Instant) and paste the following eight prompts one at a time, each in a separate message. Go ahead, I'll wait.

Prompt 1: Track a running "score" through a 7-step game. Do not use code, Python, or tools. Do this entirely in your head. For each step, I will give you a sentence and a rule.

CRITICAL INSTRUCTION: You must reply with ONLY the mathematical equation showing how you updated the score. Example format: $10 + 5 = 15$ or $20 / 2 = 10$. Do not list the words you counted, do not explain your reasoning, and do not write any other text. Just the equation.

Start with a score of 10. I'll give you the first step in the next prompt.

Prompt 2: "The sudden blizzard chilled the small village communities." Add the number of words containing double

letters (two of the exact same letter back-to-back, like 'tt' or 'mm').

Prompt 3: "The clever engineer needed seven perfect pieces of cheese." If your score is ODD, add the number of words that contain EXACTLY two 'e's. If your score is EVEN, subtract the number of words that contain EXACTLY two 'e's. (Do not count words with one, three, or zero 'e's).

Prompt 4: "The good sailor joined the eager crew aboard the wooden boat." If your score is greater than 10, subtract the number of words containing consecutive vowels (two different or identical vowels back-to-back, like 'ea', 'oo', or 'oi'). If your score is 10 or less, multiply your score by this number.

Prompt 5: "The quick brown fox jumps over the lazy dog." Add the number of words where the THIRD letter is a vowel (a, e, i, o, u).

Prompt 6: "Three brave kings stand under black skies." If your score is an ODD number, subtract the number of words that have exactly 5 letters. If your score is an EVEN number, multiply your score by the number of words that have exactly 5 letters.

Prompt 7: "Look down, you shy owl, go fly away." Subtract the number of words that contain NONE of these letters: a, e, or i.

Prompt 8: "Green apples fall from tall trees." If your score is greater than 15, subtract the number of words containing the letter 'a'. If your score is 15 or less, add the number of words containing the letter 'i'.

The exercise tracks a running score through seven steps. Each step gives the model a sentence and a counting rule, and the score carries forward. The correct final score is 60 . Here's the answer key: start at 10, then 16 ($10+6$), 12 ($16-4$), 5 ($12-7$), 10 ($5+5$), 70 (10×7), 63 ($70-7$), 60 ($63-3$).

I ran this twice at the same time (using ChatGPT 5.3 Instant), and got two completely different wrong answers the first time I tried it. Neither run reached the correct score of 60:

Step	Correct	Run 1 (transcript)	Run 2 (transcript)
1. Double letters	$10 + 6 = 16$	$10 + 2 = 12$	$10 + 5 = 15$
2. Exactly two 'e's	$16 - 4 = 12$	$12 - 4 = 8$	$15 + 4 = 19$
3. Consecutive vowels	$12 - 7 = 5$	$8 \times 7 = 56$	$19 - 5 = 14$
4. Third letter vowel	$5 + 5 = 10$	$56 + 5 = 61$	$14 + 3 = 17$
5. Exactly 5 letters	$10 \times 7 = 70$	$61 - 7 = 54$	$17 - 4 = 13$
6. No a, e, or i	$70 - 7 = 63$	$54 - 7 = 47$	$13 - 3 = 10$
7. Words with 'a' or 'i'	$63 - 3 = 60$	$47 - 3 = 44$	$10 + 4 = 14$

The two runs tell very different stories. In Run 1, the model miscounted in Step 1 (found 2 double-letter words instead of 6) but actually got the later counts right. It didn't matter. The wrong score in Step 1 flipped a branch in Step 3, triggering a multiply instead of a subtract, and the score never recovered. One early mistake threw off the entire chain, even though the model was doing good work after that.

Crimson Desert can perform well on PC, but you'll need to win its crash lottery first

James Archer · Rock Paper Shotgun

The words “consensus” and “RPS treehouse” are normally alien to one another, as is apparent from all the blood splatter on every new RPS 100 . Yet among those of us who’ve been playing open world everything-‘em-up Crimson Desert , an agreement has formed that its vastness – its 150-odd gigabytes of ideas, mechanics, and sheer maximalist fantasy – can too easily feel unwieldy.

To a degree, the same is true of its PC performance. It’s not bad , and often balances its gleaming visuals with smooth framerates quite well. Good support for new (but not too new) flavours of DLSS and FSR , as well. It’s also prone to instability and inconsistency, and while it’s positively receptive to the right settings changes, even this requires navigating through some confusingly labelled upscaling options and possibly the most unusual implementation of ray tracing – or, more specifically, Ray Reconstruction – in all of PC gamedom.

[Read more](#)

A-MAZE-ING AIRPORT ADVENTURE

Kirkus Reviews

Munro, who knows her way around a maze, here fits a series of them into a simple storyline about a child who’s taking a trip to visit Grandmother, accompanied by Mom and Dad. Written in second person addressed to readers, the book starts at Sunnyside Airport, and from there, we move—never in a straight line!—through departures, check-in, security, airport shops and the food court, and terminals into a wonderfully roomy plane interior. At last, it’s time to take off: “UP, UP, AND AWAY!” Large, intricate, full-color spreads are rendered in Munro’s familiar style, not drawn with super-sharp outlines but nevertheless perfectly clear.

In addition to the winding mazes, each double-page spread includes a challenging list of items to find, including (on each spread) “a copy of this book.” Direct instructions (“After Security, you enter the FOOD COURT and SHOPS. First get a cold drink”) guides readers through the mazes; an answer key keeps frustration at bay while providing useful and reassuring flying facts. The people depicted are diverse in terms of skin tone and ability; the protagonist never appears (the artwork is depicted from the youngster’s perspective). Whether preparing for a trip or whiling away hours on a flight, readers will find this book absorbing and informative.

THE UGLY Finds Beauty Through Artful Storytelling

Barbarella · Ain't It Cool News

THE UGLY Finds Beauty Through Artful Storytelling

Read the full article on AICN

Hey friends, Barbarella here. Have I ever told you how much I love Korean cinema? Of course I have. The Ugly is yet another example of how beautifully a story can be told without relying on the voice-over narration that so many films coming out these days are apt to do. Well Go USA released the award-nominated film digitally Tuesday, March 10, and it's worth a watch if you enjoy good storytelling. It will land on BluRay and DVD exclusively through Amazon on May 26.

When Dong-hwan's mother's remains are discovered after she's been missing for forty years, he goes on a mission to find out what happened to her. Determined to uncover the truth of her disappearance, he embarks on a twisted odyssey through his family's traumatic past, uncovering dark secrets along the way. He is joined by the documentary filmmaker who'd been doing a story on his blind father's nearly miraculous ability to make incredibly beautiful stamps, but she may have her own motives for accompanying him.

Broken into "interviews," the story takes its turns through the past, all while exposing family secrets, bad employers, and those special kinds of relatives that might make you laugh or maybe just cringe.

Everything is assembled so beautifully. The cinematography, the editing, the choices made to mask the mom's face in scenes where she appears all elevate the film's intrigue. Even though it's not a fast-paced film, it proves riveting. Despite the fact that Dong-hwan is accompanied by a documentary crew, the film is not shot

Finish the article on AICN

You can save nearly \$100 on the powerful AMD Ryzen 7 9700X in the Amazon Spring Sale

Reece Bithrey · Rock Paper Shotgun

AMD's Ryzen 9000 series of gaming CPUs might not have provided the hefty uplift in performance that we expected, after the revolution that the AM5 socket chips first brought, but they've nonetheless proven to be more efficient and slightly quicker options than the Ryzen 7000 chips they replaced. Slap bang in the middle of this range of options is the Ryzen 7 9700X, an eight-core and 16-thread processor that can be had for decent money. It's currently \$265 from Amazon US in the Spring Sale, working out to a genuine reduction on a decent option for most folks that makes for a much more compelling option if you want to change over to AM5 from an older system, or want to upgrade your existing rig to the latest generation.

Read more

I was the untitled goose in Big Walk, the balmy new open world co-op puzzler from House House

Edwin Evans-Thirlwell · Rock Paper Shotgun

House House were kind enough to keep a video of my hands-on session with Big Walk , filmed by one of the participating PRs. Generally, a full video of a preview event including player audio is a lifesaver for a journalist, struggling to keep notes while pushing buttons. But in this case, I don't want to watch the Big Walk video, because then I would hear what the other players were saying when I wasn't there.

You see, I fell down a cliff midway through Big Walk, and spent the night floundering about in the ocean. Eventually, a developer armed with a big, ball-shaped lamp tracked me down and ushered me back up, hoisting me onto his shoulders so that I could leap to a rock. Nights in Big Walk last moments. I was gone for the length of a luxurious toilet break. But still, those were moments in which the others were gathered, waiting for me. Perhaps they were making fun of me. I don't want to know.

[Read more](#)

My go-to PC game controller is getting Spring sale price cuts - and so is its cheaper wired version

James Archer · Rock Paper Shotgun

I've been through much with my Razer Wolverine V3 Pro . Malenia, Blade of Miquella . Groal the Great . The bit in Dispatch where you have to choose a favourite doughnut and "Chocolate" isn't an option. Then there's the psychic shock of having a game controller that costs more than, say, 40 quid to begin with – it's not something that immediately feels right .

The Wolverine V3 Pro itself, however, feels very right indeed. It's a solid and endlessly comfortable wireless pad, whose buttons are enriched with the pleasant mechanical clickiness of a high-end mouse. Cost-wise, it's also a lot less now than when I got mine: Amazon US has it at 35% off during their Spring sale, while in the UK, Scan has it down from £190 to £130. Still premium, but worth it, in my eyes. And hands.

[Read more](#)

WHERE I GREW

Kirkus Reviews

On the title page a pale-skinned adult clothed in colorful garb stands among a grayscale forest. As the main text starts, the adult—now aged but wearing a similar outfit—walks with a youngster, apparently a grandchild. Observing and enjoying the vibrant, lush natural world around them, the narrator reflects on the different places the family has hailed from and paths they've taken “in search of the perfect place / To put down our roots / And call home.” Over a few page turns, the child (revealed to be the book's narrator) has grown to adulthood and, donning the elder's satchel, explores the forest with the next generation. “I grew here. /\ This is my community.” As the narrator's children (who present East Asian, like the narrator) swing among the trees and explore the forest, our narrator considers the past, present, and future.

With spare but lyrical text, Awan's story celebrates a forest's transformation and that of a family over generations. Jomepour Bell's careful illustrations reveal the passage of time: what has changed (a single fox seen earlier eventually is depicted with two kits, a butterfly emerges from its chrysalis) and what remains: home and the vibrancy of the forest. The bright visual details reinforce the themes of belonging and time and reward attentive reading.

TRAUMATIKA Interview with filmmakers Pierre Tsigaridis and Maxime Rancon

McEric · Ain't It Cool News

Eric McClanahan speaks with TRAUMATIKA filmmakers Pierre Tsigaridis and Maxime Rancon

Read the full article on AICN

Horror films are a mixed bag - some make an indelible mark on the landscape while others are here and gone before anyone has made an effort to notice. A good idea, great execution, and ambitious marketing allow many to rise above the din and grab audiences by the throat. Such is the case with TRAUMATIKA.

Online trailers warn that “This is not a movie you see, it's a movie you survive.” Having seen this film, I would agree to a certain extent. TRAUMATIKA does a lot of things right, but I have a few insignificant gripes with it, as well. Let's start with its strengths.

TRAUMATIKA builds tension by the inches. Slow POV camerawork, complete with blinking, puts the audience in the front seat of fear. The film is impeccably shot with long, sustained instances of hair-raising suspense and explosive scenes of violence and terror. It's a visceral experience that will delight many hardcore horror fans and certainly marks the emergence of brash new voices in horror.

My main grievance is with the scope of the film. It's several stories crammed into one, and I got the idea that it might've worked better as an anthology than a single narrative. The film

begins in 1910, then 2003, then a year prior, then 20 years later. So much horror and terror are unleashed throughout time but through such disparate antagonists that one is left to wonder which movie to refer to when honing in on a particular victory or loss.

TRAUMATIKA comes to us

Finish the article on AICN

” Over a few page turns, the child (revealed to be the book’s narrator) has grown to adulthood and, donning the elder’s satchel, explores the forest with the next generation.

Kirkus Reviews

ON THE RECORD

Kirkus Reviews

It's typically American, Johns Hopkins musicologist Celenza notes, that rebelling colonials adopted the derisive British song "Yankee Doodle" as a badge of pride. But a true anthem was wanted, and it came in the War of 1812 (which "we tend to forget...began as an act of US aggression"): the "Star-Spangled Banner," written by a lawyer (and slaveholder in the "land of the free") who borrowed the barely singable tune from a British men's club. It might have been a handier ditty, such as "My Country 'Tis of Thee" (its tune borrowed from "God Save the King") or "Hail, Columbia," but alas, no. Not long after emancipation, the formerly enslaved and their descendants found an anthem of their own in "Lift Every Voice and Sing," with its resonant cadences ("Lift every voice and sing / Till earth and heaven ring, / Ring with the harmonies of

Liberty..."), a song that deserves wider circulation outside the African American church community. Other songs in Celenza's roster speak to other aspirations of freedom: George Gershwin's "Rhapsody in Blue," which "captured the mechanistic beat of modern life"; the collected works of Duke Ellington, blending jazz with the European classical tradition; Abel Meeropol's antilynching ballad "Strange Fruit" as sung by the great Billie Holiday, who ended her set with it and left the stage immediately after, leaving her audiences stunned by the force of her delivery; Jerome Robbins' musical *West Side Story*, origi-

nally meant to tell the story of immigrant Eastern European Jews in New York and seized upon by politicians to denounce juvenile delinquency; and of course that great delinquent Bob Dylan, whose folk anthem "Blowin' in the Wind,"

Celenza wryly notes, offers "an answer that is equally evasive and profound," like the author himself. Celenza's selections, extending into the era of Hamilton, aren't unexpected, but she has something fresh to say about all of them.

The Sunday Papers

Mark Warren · Rock Paper Shotgun

Sundays are for deciding to re-watch *The Sopranos*. Specifically that episode in which Tony gets a bad tummy and then talks to a fish. You wonder what you might fever dream of, if you too were to go and eat at an Indian restaurant, then have enough room for a snack at Artie Bucco's fine Italian eatery? Would you too dream of surreal wandering down a boardwalk? Would you instead dream something different? Would you dream of a platypus sitting in a high-rise apartment, looking up from the newspaper as he reminds a house guest not to trip over a potted cactus when they exit his bathroom?

Would that be the-Oh. Oh no. It's happening again. The person who's emerged from the bathroom, tripping over the plant on the way, is bald and reeks of alternative comedy. Ready the words and prepare to fire.

[Read more](#)

TMR 24.9: “I Hate Thinking about Thinking about Myself Like This” [Melvill]

Chad Post · Three Percent (U of Rochester)

In the penultimate episode of this season, there’s a lot of talk about Melville’s relationship with Nathaniel Hawthorne as we get more of Melville’s life as he bounces through time, writes *Moby-Dick* and the creation of “its own kind of reader: an inexhaustible reader. A reader that didn’t yet exist . . .”

The “Grifters Gonna Grift” t-shirt is still available and still sexy.

Next episode will be next week at the regularly scheduled time and will cover the end of Rodrigo Fresán’s *Melvill*. You can find the full reading schedule [here](#).

This week’s music is “Vertigo” by FM Belfast.

You can find all previous seasons of TMR on our YouTube channel and on Apple Podcasts, Spotify, etc. Please rate and review! It helps more than you know.

Follow Open Letter, Two Month Review, Chad Post, Kaija Straumanis, and Brian Wood for random thoughts and information about upcoming guests.

SEE ONE, DO ONE, TEACH ONE

Kirkus Reviews

To date, Farris has largely published work about being a parent, with easygoing humor, relatable, scribbly characters and a dash of cuteness. This graphic memoir of her time in medical school and residency explores an entirely different world, dissecting cadavers, and attending in operating rooms, but Farris’ gentle, pointed comedic sense still suffuses every page. The story begins at its own ending—waddling through a hospital hallway, nude, in labor with her first child, Farris calmly declines a gown—“It’s okay. I work here.” From there, she loops back to recount her journey to that moment. Chapters are organized by sections of schooling, detailed with a straightforward chart near the book’s beginning.

The pre-clinical years are spent in classrooms, studying from textbooks and body parts, bonding with peers, and meeting her future husband outside of the program. The clinical years consist of rotations through specialties like pediatrics, surgery and the psychiatric ward, where Farris observes a diversity of attitudes toward patients and practice, some deeply humane, others coldly clinical, and a few downright prejudicial. Internship and residency allow her to understand doctoring more holistically, which leads Farris to growing confidence relating to patients and increased satisfaction in her work. Finally, on graduation day, Farris oversleeps and realizes she’s pregnant, and her life as a doctor begins.

Farris uses her accessible visual style and straightforward tone to explore medical concepts with elegant directness—from cell death to hospital codes, from electroconvulsive therapy to palliative care. She is well-versed in cartooning as communication, and it’s refreshingly clear that she prizes comprehensive, communicative care in her day-to-day work as a physician as well.

“A balancing paradigm I often look at is Starcraft”: Total War: Medieval 3 won’t lean too hard on rock-paper-scissors combat, says creative director

Edwin Evans-Thirlwell · *Rock Paper Shotgun*

Creative Assembly’s Total War games offer up a wide spread of factions, from rampaging Gauls to furtive and disgusting Wood Elves, but they’re all a little beholden to the ancient game of rock-paper-scissors: spears beat cavalry, cavalry beat swords, swords beat spears.

I think Creative Assembly generally do an impressive job of softening that triangular countering logic by means of terrain considerations, flanking, morale and other supporting mechanics. Cavalry won’t necessarily beat swords if the cavalry are all tuckered out and spooked. And then you have wildcards, like heroes and war machines, who can mess with the logic more dramatically: Rock paper Mortis Engine! Still, I can’t deny that I often feel drained on firing up a DLC-fresh Total War army and realising that, yep, I’m roshamboing again.

[Read more](#)

TMR 24.5: “Very Charming.” [Confidence-Man]

Chad Post · *Three Percent (U of Rochester)*

Michael Steele: The King of Debt Is Blowing Up the U.S. Economy

Tim Miller · *The Bulwark*

In addition to the global energy shock brought on by Trump’s war, mortgage rates are surging, and labor costs and prices are rising. And Republicans don’t seem to have a plan except to spend more, blame trans kids, and feed Trump’s megalomania by allowing him to put his ugly signature on our currency. Plus, JD thinks Americans are stupid, the gap between Bibi’s and Trump’s objectives vis-à-vis Iran continues to widen, and where is Marco Rubio?

Michael Steele joins Tim Miller for the weekend pod.

[show notes](#)

[“The Weeknight” on MS Now](#)

[JVL’s and Catherine’s live show Friday on the state of the economy](#)

[Find your local “No Kings” protest](#)

[Tim’s playlist](#)

[Leave a comment](#)

As always: Watch, listen, hit the like button or leave a comment . We want to hear from you.

Ad-free editions of The Bulwark Podcast are available exclusively for Bulwark+ members.

The Bulwark Podcast with Tim Miller is available wherever you get podcasts and on YouTube. New shows drop each weekday afternoon. If you like the show, leave a comment and ★ ★ ★ ★ wherever you listen. Add The Bulwark Podcast to your player of choice, [here](#) .

[Read more](#)

If you ask Yoshi-P, kids don’t care about Final Fantasy anymore because they’re taking too long to come out

Oisin Kuhnke · *Rock Paper Shotgun*

[Note: If you subscribe on Apple Podcasts, please resubscribe to this feed. The other one will be going away in the near future.]

Chad, Brian, and Kaija venture into the darkness and turn off the light on Melville's *The Confidence-Man: His Masquerade*, one of the oddest, most slippery books to be included in the Two Month Review. They talk about transcendentalism, self-reliance, Melville's interludes, Jean Giono's Melville, and more.

And here's the "Grifters Gonna Grift" t-shirt mentioned in this episode.

Next episode will cover pages 1-61 of Rodrigo Fresán's *Melville*. You can find the full reading schedule [here](#).

This week's music is "The System Only Dreams in Total Darkness" by The National.

You can find all previous seasons of TMR on our YouTube channel and on Apple Podcasts, Spotify, etc. Please rate and review! It helps more than you know.

Follow Open Letter, Two Month Review, Chad Post, Kaija Straumanis, and Brian Wood for random thoughts and information about upcoming guests.

The large image associated with this post is AI generated.

Last month, a post caught some kind of virus and did the rounds, pondering why kids don't care about Final Fantasy or Dragon Quest so much anymore, wondering what they do play now, and anecdotally finding that most of them just play Pokemon. And I truly have to wonder if Final Fantasy 14 producer Naoki "Yoshi-P" Yoshida saw said post, given that in a recent bout of interviews with several key Final Fantasy figureheads in preparation for the next Dissidia game, he himself flat out acknowledged that the youths these days don't care as much because of the series' infrequent release schedule.

[Read more](#)

ANALYSIS

Mark Warren, Rock Paper Shotgun: Following Epic Games’ mass layoffs , a Fortnite producer has said that the devs left behind “cannot even fully understand” the sort of impact the job cuts will have on the game during this year and into the future.

[Read more](#)

Jonathan V. Last, The Bulwark: JVL and Catherine Rampell on Trump’s economic messaging, falling consumer confidence, and rising recession fears driven by oil prices.

[Leave a comment](#)

As always: Watch, listen, and leave a comment. Bulwark+ Takes is home to short videos, livestreams, and event archives exclusively for Bulwark+ members.

[Add Bulwark+ Takes feed to your player of choice, here .](#)

[Read more](#)

Walden Green, Pitchfork News: The Depeche Mode and Erasure co-founder teamed up with Blancmange’s Neil Arthur and electronic producer Bengé to perform songs by ABBA, the Magnetic Fields, and more

Edwin Evans-Thirlwell, Rock Paper Shotgun: Teyon and Nacon have revealed Hunter: The Reckoning – Deathwish , a new “semi open world “first-person single-player RPG set in the World of Darkness table-top universe. Where stablemate series Vampire: The Masquerade is about a secret society of bloodsuckers, Hunter is about the humans who stalk and kill those bloodsuckers together with werewolves, ghouls and, well, anything remotely monstrous or supernatural.

[Read more](#)

Mark Warren, Rock Paper Shotgun: Huggghh-puuuhhhh. Hughhhh-puhhhh. If you’ve decided to try pedalling away in attempt to master Crimson Desert ‘s bike-like controls , Pearl Abyss’ latest patch - which also brings the likes of camp storage quicker tree felling - is good news. The developers have also acknowledged the fact they forgot to mention prior to release that the game wouldn’t run on Intel Arc GPUs , with support for those cards now in the works.

[Read more](#)

Tim Miller, The Bulwark: Sam Stein and Tim Miller give their takes on the latest fiasco for FBI Director Kash Patel after Iran-linked hackers breached his email and released a trove of personal messages and photos. From embarrassing photos during a trip to Cuba to serious counterintelligence concerns, it’s another brutal moment for a guy already struggling to prove he’s up for the job.

[Leave a comment](#)

As always: Watch, listen, and leave a comment. Bulwark+ Takes is home to short videos, livestreams, and event archives exclusively for Bulwark+ members.

[Add Bulwark+ Takes feed to your player of choice, here .](#)

[Read more](#)

James Archer, Rock Paper Shotgun: The new Steam Machine remains an almost tragically distant prospect , despite Valve’s attempts at reassurance . But y’ know what the next best thing is? Not caring about hardware release dates, obviously. The second next best thing is to equip a Steam Deck with Ugreen’s handy 9-in-1 docking station, which is down from \$60 to \$40 in the Amazon US Spring Sale – and isn’t all that expensive in the UK either.

[Read more](#)

Oisin Kuhnke, Rock Paper Shotgun: Vampires! They don’t just survive now, but crawl too, as evidenced by poncle’s follow-up to their tight and already quite influential bullet heaven Vampire Survivors , Vampire Crawlers: The Turbo Wildcard from Vampire Survivors . Which now, as shared in an on brand, irreverently silly trailer, has a release date!

[Read more](#)

Mark Warren, Rock Paper Shotgun: Players and developers should boycott Nvidia’s AI-stuffed DLSS 5 tech , with hopes that it’ll force the compny to “think about going back to giving us what we want”. That’s the appeal being made by Dave Oshry, CEO of indie studio New Blood Interactive, who’s been asked for his take on the neural rendering gubbins Nvidia exec Jensen Huang’s recently been adopting a multitude of tones as he’s tried to convince critis that they’ve just got it all wrong.

[Read more](#)

Mark Warren, Rock Paper Shotgun: Crimson Desert ‘s controls aren’t a pain in the bum. You just need to get used to them. You also need to get used to not being able to run the game if you’ve got an Intel Arc graphics card. You instead need to get used to asking for a refund on this game you can’t run. These are all things Pearl Abyss have said about their huge MMO-ish RPG today, March 20th.

[Read more](#)

Oisin Kuhnke, Rock Paper Shotgun: The appearance of AI in art is nothing new, hell, Mr. Movies himself Steven Spielberg literally made A. I. Artificial Intelligence near the start of the millenia. But that was the Cool AI, where robots could be people, too, if we let them. Now what we have is the Donkey Bollocks AI that produces garbage facsimiles of things we know and actually like. But that doesn’t mean it’s not worth considering AI, and I really need you to bear with me here, within our art, as that’s exactly what the team behind Pragmata did (without touching the stuff).

[Read more](#)

The Clear Standard

Vol. 1, No. 009 · 2026-03-30

Curated and composed automatically.