

The Edition

VOL. 1, NO. 032

2026-03-30

FEATURES

How American Kids Have Been Collateral Damage in Trump's Immigration Crackdown

Nicole Foy · ProPublica

For much of last year, Trump administration officials insisted that no Americans were caught up in the government's immigration dragnet.

ProPublica and many others repeatedly documented that is not true: Americans have even been kicked, dragged and detained for days by immigration agents.

On Tuesday, House and Senate Democrats are spotlighting a particularly troubling part of the crackdown: the American children who have been collateral damage in the deportation campaign.

The forum the lawmakers are holding is part of an ongoing congressional investigation prompted by ProPublica's report last fall that more than 170 U. S. citizens have been detained by immigration agents for some amount of time. That included Americans who have been handcuffed, held at gunpoint or simply prevented from leaving their location.

As of last October, more than 20 of those citizens were children, ranging from toddlers to teens. A toddler, a preschooler and a 7-year-old — all citizens — were deported despite their documented parents claiming they wanted to keep the children in the U. S.

In response to questions, Department of Homeland Security spokesperson Lauren Bis said in a statement that Immigration and Customs Enforcement “does NOT deport United States citizens or separate families,”

American children held along with their families will be sharing their stories at Tuesday's forum. That includes two families whose accounts were featured in ProPublica investigations.

Eighteen-year-old Fernando Hernández García, who is using a pseudonym to protect the safety of his family in Mexico, is speaking on behalf of his 11-year-old sister. Both siblings are citizens .

Last year, the family was driving to Houston to get emergency treatment for the girl, who was recovering from brain cancer. Border Patrol agents ignored a hospital letter that the family had used previously to go through checkpoints. This time, agents held the family until they were deported the next day to Mexico. With few other options, the American children went with their parents — except for Hernández García, who had not been detained and stayed to earn money and send medicine home.

The family's lawyers say they have not been able to access the care they need for their daughter in Mexico, and they have applied for humanitarian parole to return. Customs and Border Protection previously told ProPublica the family's account was inaccurate but declined to provide specifics.

We Found That More Than 170 U. S. Citizens Have Been Held by Immigration Agents. They've Been Kicked, Dragged and Detained for Days.

Some Americans Have Already Been Caught in Trump's Immigration Dragnet. More Will Be.

We Found More Than 40 Cases of Immigration Agents Using Banned Chokeholds and Other Moves That Can Cut Off Breathing

Also speaking is 16-year-old Arnolando Bazan. As ProPublica detailed earlier this year, Bazan was tackled and choked by immigration agents who were chasing his undocumented father in Houston.

Bystanders filmed the teen screaming that he was a minor and a U. S. citizen. After agents knelt on his neck and put him in a choke hold, then they handcuffed him.

Bazan told ProPublica that when he was in a choke hold, “I felt like I was seeing the light.” He said he's now speaking up — including on Capitol Hill — to help keep others from going through the same. “I don't think nobody's safe anymore.”

DHS said in its statement that Bazan elbowed an officer in the face as he was detained, which the teen denies. The agency's spokesperson added that any allegations that agents assaulted Bazan “are FALSE.”

It's unclear exactly how many American kids have been held. The government doesn't disclose how many Americans are detained , even briefly, during immigration enforcement.

Former immigration officials told ProPublica that it used to be rare to encounter, let alone hold, American children for any amount of time. While the officials couldn't recall a specific policy prohibiting it, they said past administrations just didn't prioritize arresting families during immigration enforcement in the interior of the country. (A ProPublica investigation published Monday found that in his second term, President Donald Trump has deported mothers of U. S. children at four times the rate Biden did.)

In a report shared with ProPublica , the minority staff from the Senate Permanent Subcommittee on Investigations and House Committee on Oversight and Government Reform compiled 128 cases of children — a mix of citizens and noncitizens — who were injured, left unattended or otherwise put at risk by enforcement operations conducted by Department of Homeland Security agents.

The review found that citizen children caught up in immigration operations were also exposed to chemical agents , were placed in restraints or required medical attention, and some were held at gunpoint, were left unattended when agents detained their parents, or were present when agents smashed car windows or rammed their vehicles .

How Do You Know That You Love Somebody? Philosopher Martha Nussbaum's Incompleteness Theorem of the Heart's Truth, from Plato to Proust

Maria Popova · *The Marginalian* (Brain Pickings)

"The alternations between love and its denial, suffering and denial of suffering ... constitute the most essential and ubiquitous structural feature of the human heart."

"The state of enchantment is one of certainty," W. H. Auden wrote in his commonplace book. "When enchanted, we neither believe nor doubt nor deny: we know, even if, as in the case of a false enchantment, our knowledge is self-deception." Nowhere is our capacity for enchantment, nor our capacity for self-deception, greater than in love — the region of human experience where the path to truth is most obstructed by the bramble of rationalization and where we are most likely to be kidnapped by our own delicious delusions. There, it is perennially difficult to know what we really want; difficult to distinguish between love and lust; difficult not to succumb to our perilous tendency to idealize; difficult to reconcile the closeness needed for intimacy with the psychological distance needed for desire.

How, then, do we really know that we love another person?

That's what Martha Nussbaum, whom I continue to consider the most compelling philosopher of our time, examines in her 1990 book *Love's Knowledge: Essays on Philosophy and Literature* (public library) — the sandbox in which Nussbaum worked out the ideas that would become, a decade later, her incisive treatise on the intelligence of emotions.

Devising a sort of incompleteness theorem of the heart's truth, Nussbaum writes:

We deceive ourselves about love — about who; and how; and when; and whether. We also discover and correct our self-deceptions. The forces making for both deception and unmasking here are various and powerful: the unsurpassed danger, the urgent need for protection and self-sufficiency, the opposite and equal need for joy and communication and connection. Any of these can serve either truth or falsity, as the occasion demands. The difficulty then becomes: how in the midst of this confusion (and delight and pain) do we know what view of ourselves, what parts of ourselves, to trust? Which stories about the condition of the heart are the reliable ones and which the self-deceiving fictions? We find ourselves asking where, in this plurality of discordant voices with which we address ourselves on this topic of perennial self-interest, is the criterion of truth? (And what does it mean to look for a criterion here? Could that demand itself be a tool of self-deception?)

With an eye to Proust's *In Search of Lost Time* and its central theme of how our intellect blinds us to the wisdom of the heart, Nussbaum contemplates the nature of those experiences "in which the self-protective tissue of rationalization is in a moment cut through, as if by a surgeon's knife": Proust's protagonist, Marcel, has rationally convinced him-

self that he no longer loves his beloved, Albertine, but is jolted into confronting the falsity of that rationalization upon receiving news of her death; in the shock of his intense sorrow, he instantly gains the knowledge, far deeper and more sinewy than the intellect's, that he did, in fact, love Albertine.

In a testament to Proust's assertion that "the end of a book's wisdom appears to us as merely the start of our own," Nussbaum writes:

Proust tells us that the sort of knowledge of the heart we need in this case cannot be given us by the sciences of psychology, or, indeed, by any sort of scientific use of intellect. Knowledge of the heart must come from the heart — from and in its pains and longings, its emotional responses.

Such a conception of love's knowledge, to be sure, stands radically against the long intellectual tradition of rationalism stretching from Plato to Locke like an enormous string of reason that plays only one note, deaf to the symphonic complexity of the emotional universe. The Proustian view calls for a restoration of lost nuance. Pointing to "the pseudotruths of the intellect," Nussbaum revisits Marcel's predicament, wherein the intellect has imposed an illusory sense of order and structure upon the entropy of the emotions:

The shock of loss and the attendant welling up of pain show him that his theories were forms of self-deceptive rationalization — not only false about his condition but also manifestations and accomplices of a reflex to deny and close off one's vulnerabilities that Proust finds to be very deep in all of human life. The primary and most ubiquitous form of this reflex is seen in the operations of habit, which makes the pain of our vulnerability tolerable to us by concealing need, concealing particularity (hence vulnerability to loss), concealing all the pain-inflicting features of the world — simply making us used to them, dead to their assaults. When we are used to them we do not feel them or long for them in the same way; we are no longer so painfully afflicted by our failure to control and possess them. Marcel has been able to conclude that he is not in love with Albertine, in part because he is used to her. His calm, methodical intellectual scrutiny is powerless to dislodge this "dream deity, so riveted to one's being, its insignificant face so incrusting in one's heart." Indeed, it fails altogether to discern the all-important distinction between the face of habit and the true face of the heart.

Nussbaum considers how our over-reliance on the intellect for clarity about love produces instead a kind of myopia:

Intellect's account of psychology lacks all sense of proportion and depth and importance... [Such a] cost-benefit analysis of the heart — the only comparative assessment of which intellect, by itself, is capable — is bound, Proust suggests, to miss differences of depth. Not only to miss them, but to impede their recognition. Cost-benefit analysis is a way of comforting oneself, of putting oneself in control

This JBL Boombox Is 40% Off During Amazon's Big Spring Sale

Naima Karp · Lifehacker

We may earn a commission from links on this page. Deal pricing and availability subject to change after time of publication.

Old school boomboxes might be a relic of the past, but they've been reinterpreted in recent years with sleek, modern, and more portable silhouettes. While not as compact as some Bluetooth speakers, they pack a bigger punch in sound, especially when it comes to playing music outdoors. One of the best-rated ones on the market is the JBL Boombox 3, which earned a PCMag Editor's Choice Award—and it's currently 40% off for Amazon's Spring Sale, bringing it down to its lowest price ever of \$299.95 (originally \$499.95).

JBL Boombox 3

\$299.95 at Amazon

\$499.95 Save \$200.00

\$299.95 at Amazon

\$499.95 Save \$200.00

At 14.7 pounds, it might not be the most lightweight speaker, but it's still technically portable given the built-in handle that lets you carry it from place to place. The control bar lets you link to other JBL speakers, control volume, power, play/pause and Bluetooth pairing, but doesn't allow you to go back to the previous track. It has a covered rear panel with a compartment housing a 3.5mm aux input, a USB-A port that allows you to charge other devices, and a connector for the included power cable. Aesthetically, it's similar to the Boombox 2, but adds an extra subwoofer.

It's compatible with Bluetooth 5.3 and supports the SBC codec, but not AAC or AptX. Still, the bass and volume have impressive output in both AC and battery power modes. Battery life lasts up to 24 hours on a charge and takes around 6.5 hours to fully charge from empty. An IP67 rating makes it dust-proof and waterproof (up to one meter of depth for 30 minutes), so poolside hangs or unexpected rain shouldn't pose an issue. Just note that the snap-shut cover on the back needs to be closed for the rating to apply. The companion app includes a customizable three-band EQ, firmware updates, and battery life monitoring.

If you're looking for a loud, durable outdoor speaker with deep bass and adjustable EQ, and you don't mind slightly limited Bluetooth codec support or a heavier build, the size-to-power ratio makes this award-winning speaker a reliable choice. With summer right around the corner, and the JBL Boombox 3 at nearly half off, now is a great time to invest.

Our Best Editor-Vetted Amazon Big Spring Sale Deals Right Now

Apple AirPods Pro 3 Noise Cancelling Heart Rate Wireless Earbuds

— \$199.00

(List Price \$249.00)

Apple iPad 11" 128GB A16 WiFi Tablet (Blue, 2025)

— \$299.00

(List Price \$349.00)

Samsung Galaxy Tab A11+ 128GB Wi-Fi 11" Tablet (Gray)

— \$209.99

(List Price \$249.99)

Sony WH1000XM6- Best Wireless Noise Canceling Headphones

— \$398.00

(List Price \$459.99)

Apple Watch Series 11 (GPS, 42mm, S/M Black Sport Band)

— \$299.00

(List Price \$399.00)

Blink Video Doorbell Wireless (Newest Model) + Sync Module Core

— \$35.99

(List Price \$69.99)

Fire TV Stick 4K Max Streaming Player With Remote

— \$34.99

(List Price \$59.99)

Amazon Kindle Colorsoft 16GB 7" eReader (Black)

— \$169.99

(List Price \$249.99)

Deals are selected by our commerce team

10 Shows Like ‘The Comeback’ You Should Watch Next

Ross Johnson · Lifehacker

We may earn a commission from links on this page.

Valerie Cherish is back, and not a moment too soon. Twenty years after a modestly rated but absolutely brilliant first season, mega-*cringe* Hollywood satire *The Comeback* has returned to HBO for one more painfully embarrassing season (complimentary) Co-created by and starring Lisa Kudrow, it's the hilarious, dark, and occasionally moving story of a C-list celebrity determined to stay in the limelight at whatever cost.

Valerie Cherish—clueless, egotistical, and all-too-human—is one of TV's indelible characters, simultaneously a victim of ageism, misogyny, and celebrity culture, and also a woman who doesn't care about any of that provided she can find a way to scrabble to the top. Likewise, these 10 shows spotlight women (mostly) of a certain age navigating career challenges with...mixed results.

Hacks (2021 –)

On the surface, Valerie Cherish and Deborah Vance don't have all that much in common, but they often feel like two sides of a coin in the ways in which they battle, and concede to, the pressures of being a woman in modern show business. *Hacks* finds 25-year-old writer Ava (Hannah Einbinder) struggling to get her career back on track after getting cancelled over a tweet. She reluctantly takes a job with Deborah Vance (Jean Smart), a comedy trailblazer who remains popular with an older Vegas crown, but whose career is largely on autopilot. They're an entirely mismatched pair, but their chemistry is ultimately explosive, with Jean Smart doing some of the best work of her incredible career as the often deeply unlikeable Vance, and Einbinder more than holding her own in return. It's funny, bitchy, and surprisingly moving when it wants to be. Stream *Hacks* on HBO Max .

Hacks (2021 –)

Girls5Eva (2021 – 2024)

The premise is very much about a comeback: A very '90s group (Sara Bareilles, Busy Philipps, Paula Pell, and Renée Elise Goldsberry) who made it big very briefly with exactly one hit song gets its own shot at a career resurgence (they're Girls5Eva because they expected to be in the game “longer than 4 ever”—turned out, not so much; also there used to be a fifth member, but she died in a tragic infinity pool accident). After an up-and-coming rapper samples one of their old hits, the band reunites to try to find new relevancy in the 2020s. It doesn't go great! Tina Fey is one of the executive producers, and the (very fun) '90s-inspired original music was composed and written for the show by the creators. Stream Girls5Eva on Peacock and Netflix .

Girls5Eva (2021 – 2024)

at Peacock

at Peacock

Veep (2012 – 2019)

There's a sweet vulnerability to Valerie Cherish absent from pretty much anyone in *Veep* , but the air of cringeworthy desperation hangs heavy over both shows. *Veep* asks: What if the very worst people imaginable had control of all the levers of power in America? What would that be like? In the highly satirical and in no way politically relevant series, Julia Louis-Dreyfus brilliantly plays Vice President Selina Meyer: part icon, part worst-case-scenario for the feminist dream. The supporting cast consists almost entirely of entitled rich people climbing over each other for infinitesimal extra bits of power while trying to appear as folk heroes to the rubes (us voters). It's one of the funniest shows ever on TV, with the nastiest put-downs...if you can handle the fact that history has overtaken it in favor of a reality that's even more absurd. Stream *Veep* on HBO Max.

Veep (2012 – 2019)

at HBO Max

at HBO Max

The Studio (2025 –)

Industry exec Matt Remick (*Comeback* vet Seth Rogen) loves movies, and when he signs on for a high-profile role at the fictional Continental Studios, he feels like his time has come. He aims to make a real difference in returning an increasingly IP-driven movie industry to its creative roots. —until about a minute into his new job, when the CEO (Bryan Cranston) reveals his first order of business will be marketing the *Kool-Aid Man Movie* . Hollywood satires of yore have focused on the industry as one that eats people up and spits them out, but the spin here is that Remick yearns for those good old days, much as Valerie can never quite seem to move forward. The late Catherine O'Hara, Ike Barinholtz, Chase Sui Wonders (*Bodies Bodies Bodies*), and Kathryn Hahn also star, and like *The Comeback* , the show boasts a long list of celebrity cameos. Stream *The Studio* on Apple TV+ .

The Studio (2025 –)

at Apple TV+

at Apple TV+

Enlightened (2011 – 2013)

Before *The White Lotus* , Mike White co-created (with star Laura Dern) this beloved, if short-lived, comedy-drama about a middle-aged woman who has a complete mental breakdown following a demotion at the job to which she's devoted her life. Following a two-month stay in an holistic treatment facility, Dern's Amy Jellicoe becomes determined to approach life with a new perspective, focusing on medi-

Overconfidence is how wars are lost – lessons from Vietnam, Afghanistan and Ukraine for the war in Iran were ignored

Monica Duffy Toft, Professor of International Politics and Director of the Center for Strategic Studies, The Fletcher School, Tufts University · The Conversation (Politics)

Plumes of smoke and fire rise after debris from an intercepted Iranian drone struck an oil facility, according to authorities, in Fujairah, United Arab Emirates, on March 14, 2026. AP Photo/Altaf Qadri

Wars are rarely lost first on the battlefield. They are lost in leaders' minds – when leaders misread what they and their adversaries can do, when their confidence substitutes for comprehension, and when the last war is mistaken for the next one.

The Trump administration's miscalculation of Iran is not an anomaly. It is the latest entry in one of the oldest and most lethal traditions in international politics: the catastrophic gap between what leaders believe going in and what war actually delivers.

I'm a scholar of international security, civil wars and U. S. foreign policy, and author of the book "Dying by the Sword," which examines why the United States repeatedly reaches for military solutions and why such interventions rarely produce durable peace. The deeper problem with the U. S. war in Iran, as I see it, was overconfidence bred by recent success.

Dismissed concerns

Before the conflict involving Iran, Israel and the U. S. escalated, Energy Secretary Chris Wright dismissed concerns about oil market disruption, noting that prices had barely moved during the 12-day war in June 2025 between Israel and Iran. Other senior officials agreed.

What followed was significant: Iranian-aimed missile and drone barrages against U. S. bases, Arab capitals and Israeli population centers. Then Iran effectively closed the Strait of Hormuz, through which roughly 20% of the world's oil supply passes daily – not with a naval blockade, not with mines or massed anti-ship missiles, but with cheap drones.

A few strikes in the vicinity of the strait were enough. Insurers and shipping companies decided the transit was unsafe. Tanker traffic dropped to zero, although the occasional ship has made it through recently. Analysts are calling it the biggest energy crisis since the 1970s oil embargo.

President Donald Trump expressed anger on March 17, 2026, at allies who did not agree to help the U. S. force the reopening of the Strait of Hormuz to tanker traffic.

Iran's new supreme leader, Mojtaba Khamenei, has since vowed to keep the strait closed. U. S. Sen. Chris Murphy, a Connecticut Democrat, reported after a closed-door briefing that the administration had no plan for the strait and did not know how to get it safely back open.

With no embassy in Tehran since 1979, the U. S. relies heavily for intelligence on CIA networks of questionable quality and Israeli assets who have their own country's interests in mind. So the U. S. did not anticipate that Iran had rebuilt and dispersed significant military capacity since June 2025, nor that it would strike neighbors across the region, including Azerbaijan, widening the conflict well beyond the Persian Gulf.

The war has since reached the Indian Ocean, where a U. S. submarine sank an Iranian frigate 2,000 miles from the theater of war, off the coast of Sri Lanka – just days after the ship had participated in Indian navy exercises alongside 74 nations, including the U. S.

The diplomatic damage to Washington's relationships with India and Sri Lanka, two countries whose cooperation is increasingly important as the United States seeks partners to manage and mitigate Iran's blockade, was entirely foreseeable. Washington has put them in a difficult position, with India choosing diplomacy with Iran to secure passage for its vessels and Sri Lanka opting to retain its neutrality, underscoring its vulnerable position.

But U. S. planners didn't foresee any of this.

The wrong lesson from Venezuela

The swift military intervention by the U. S. in Venezuela in January 2026 produced rapid results with minimal blowback – appearing to validate the administration's faith in coercive action.

But clean victories are dangerous teachers.

They inflate what I call in my teaching the "hubris/humility index" – the more a leadership overestimates its own abilities, underestimates the adversary's and dismisses uncertainty, the higher the score and the more likely disaster will ensue. Clean victories inflate the index precisely when skepticism is most needed, because they suggest the next adversary will be as manageable as the last.

Political scientist Robert Jervis demonstrated decades ago that misperceptions in international relations are not random but follow patterns. Leaders tend to project their own cost-benefit logic onto opponents who do not share it. They also fall into "availability bias," allowing the most recent operation to stand in for the next.

The higher the hubris/humility index, the less likely there is to be the kind of strategic empathy that might ask: How does Tehran see this? What does a regime that believes its survival is at stake actually do? History shows that such a regime escalates, improvises and takes risks that appear irrational from an outside perspective but are entirely rational from within.

Recent cases reveal this unmistakable pattern.

How we use Abstract Syntax Trees (ASTs) to turn Workflows code into visual diagrams

André Venceslau · Cloudflare Blog

Cloudflare Workflows is a durable execution engine that lets you chain steps, retry on failure, and persist state across long-running processes. Developers use Workflows to power background agents, manage data pipelines, build human-in-the-loop approval systems, and more.

Last month, we announced that every workflow deployed to Cloudflare now has a complete visual diagram in the dashboard.

We built this because being able to visualize your applications is more important now than ever before. Coding agents are writing code that you may or may not be reading. However, the shape of what gets built still matters: how the steps connect, where they branch, and what's actually happening.

If you've seen diagrams from visual workflow builders before, those are usually working from something declarative: JSON configs, YAML, drag-and-drop. However, Cloudflare Workflows are just code. They can include Promises, Promise.all, loops, conditionals, and/or be nested in functions or classes. This dynamic execution model makes rendering a diagram a bit more complicated.

We use Abstract Syntax Trees (ASTs) to statically derive the graph, tracking Promise and await relationships to understand what runs in parallel, what blocks, and how the pieces connect.

Keep reading to learn how we built these diagrams, or deploy your first workflow and see the diagram for yourself.

Here's an example of a diagram generated from Cloudflare Workflows code:

Dynamic workflow execution

Generally, workflow engines can execute according to either dynamic or sequential (static) execution order. Sequential execution might seem like the more intuitive solution: trigger workflow → step A → step B → step C, where step B starts executing immediately after the engine completes Step A, and so forth.

Cloudflare Workflows follow the dynamic execution model. Since workflows are just code, the steps execute as the runtime encounters them. When the runtime discovers a step, that step gets passed over to the workflow engine, which manages its execution. The steps are not inherently sequential unless awaited — the engine executes all unawaited steps in parallel. This way, you can write your workflow code as flow control without additional wrappers or directives. Here's how the handoff works:

An engine, which is a “supervisor” Durable Object for that instance, spins up. The engine is responsible for the logic of the actual workflow execution.

The engine triggers a user worker via dynamic dispatch, passing control over to Workers runtime.

When Runtime encounters a step.do, it passes the execution back to the engine.

The engine executes the step, persists the result (or throws an error, if applicable) and triggers the user Worker again.

With this architecture, the engine does not inherently “know” the order of the steps that it is executing — but for a diagram, the order of steps becomes crucial information. The challenge here lies in getting the vast majority of workflows translated accurately into a diagnostically helpful graph; with the diagrams in beta, we will continue to iterate and improve on these representations.

Parsing the code

Fetching the script at deploy time, instead of run time, allows us to parse the workflow in its entirety to statically generate the diagram.

Taking a step back, here is the life of a workflow deployment:

To create the diagram, we fetch the script after it has been bundled by the internal configuration service which deploys Workers (step 2 under Workflow deployment). Then, we use a parser to create an abstract syntax tree (AST) representing the workflow, and our internal service generates and traverses an intermediate graph with all WorkflowEntrypoints and calls to workflows steps. We render the diagram based on the final result on our API.

When a Worker is deployed, the configuration service bundles (using esbuild by default) and minifies the code unless specified otherwise. This presents another challenge — while Workflows in TypeScript follow an intuitive pattern, their minified Javascript (JS) can be dense and indigestible. There are also different ways that code can be minified, depending on the bundler.

Here's an example of Workflow code that shows agents executing in parallel:

```
const summaryPromise = step.do( `summary agent (loop ${loop})`, async () => { return runAgentPrompt( this.env, SUMMARY_SYSTEM, buildReviewPrompt( 'Summarize this text in 5 bullet points.', draft, input.context ) ); } );
const correctnessPromise = step.do( `correctness agent (loop ${loop})`, async () => { return runAgentPrompt( this.env, CORRECTNESS_SYSTEM, buildReviewPrompt( 'List correctness issues and suggested fixes.', draft, input.context ) ); } );
const clarityPromise = step.do( `clarity agent (loop ${loop})`, async () => { return runAgentPrompt( this.env, CLARITY_SYSTEM, buildReviewPrompt( 'List clarity issues and suggested fixes.', draft, input.context ) ); } );
```

Some notes on upgrading Hugo

Julia Evans

Warning: this is a post about very boring yakshaving, probably only of interest to people who are trying to upgrade Hugo from a very old version to a new version. But what are blogs for if not documenting one's very boring yakshaves from time to time?

So yesterday I decided to try to upgrade Hugo. There's no real reason to do this – I've been using Hugo version 0.40 to generate this blog since 2018, it works fine, and I don't have any problems with it. But I thought – maybe it won't be as hard as I think, and I kind of like a tedious computer task sometimes!

I thought I'd document what I learned along the way in case it's useful to anyone else doing this very specific migration. I upgraded from Hugo v0.40 (from 2018) to v0.135 (from 2024).

Here are most of the changes I had to make:

change 1: template “theme/partials/thing.html is now partial thing.html

I had to replace a bunch of instances of `{{ template “theme/partials/header.html” . }}` with `{{ partial “header.html” . }}`.

This happened in v0.42 :

We have now virtualized the filesystems for project and theme files. This makes everything simpler, faster and more powerful. But it also means that template lookups on the form `{{ template “theme/partials/pagination.html” . }}` will not work anymore. That syntax has never been documented, so it's not expected to be in wide use.

change 2: `. Data. Pages` is now `site.RegularPages`

This seems to be discussed in the release notes for 0.57.2

I just needed to replace `. Data. Pages` with `site.RegularPages` in the template on the homepage as well as in my RSS feed template.

change 3: `. Next` and `. Prev` got flipped

I had this comment in the part of my theme where I link to the next/previous blog post:

“next” and “previous” in hugo apparently mean the opposite of what I'd think they'd mean intuitively. I'd expect “next” to mean “in the future” and “previous” to mean “in the past” but it's the opposite

It looks they changed this in ad705aac064 so that “next” actually is in the future and “prev” actually is in the past. I definitely find the new behaviour more intuitive.

downloading the Hugo changelogs with a script

Figuring out why/when all of these changes happened was a little difficult. I ended up hacking together a bash script to download all of the changelogs from github as text files , which I could then grep to try to figure out what happened. It turns out it's pretty easy to get all of the changelogs from the GitHub API.

So far everything was not so bad – there was also a change around taxonomies that's I can't quite explain, but it was all pretty manageable, but then we got to the really tough one: the markdown renderer.

The blackfriday markdown renderer (which was previously the default) was removed in v0.100.0 . This seems pretty reasonable:

It has been deprecated for a long time, its v1 version is not maintained anymore, and there are many known issues. Goldmark should be a mature replacement by now.

Fixing all my Markdown changes was a huge pain – I ended up having to update 80 different Markdown files (out of 700) so that they would render properly, and I'm not totally sure

why bother switching renderers?

The obvious question here is – why bother even trying to upgrade Hugo at all if I have to switch Markdown renderers? My old site was running totally fine and I think it wasn't necessarily a good use of time, but the one reason I think it might be useful in the future is that the new renderer (goldmark) uses the CommonMark markdown standard , which I'm hoping will be somewhat more futureproof. So maybe I won't have to go through this again? We'll see.

Also it turned out that the new Goldmark renderer does fix some problems I had (but didn't know that I had) with smart quotes and how lists/blockquotes interact.

finding all the Markdown problems: the process

The hard part of this Markdown change was even figuring out what changed. Almost all of the problems (including #2 and #3 above) just silently broke the site, they didn't cause any errors or anything. So I had to diff the HTML to hunt them down.

Here's what I ended up doing:

Generate the site with the old version, put it in `public_old`

Generate the new version, put it in `public`

Diff every single HTML file in `public/` and `public_old` with this `diff.sh` script and put the results in a `diffs/` folder

Run variations on `find diffs -type f | xargs cat | grep -C 5 ‘(31m|32m)’ | less -r` over and over again to look at every single change until I found something that seemed wrong

And the only thing I needed depends on the size of the list markers. Here's a reference in Com-

Investigating multi-vector attacks in Log Explorer

Jen Sells · Cloudflare Blog

In the world of cybersecurity, a single data point is rarely the whole story. Modern attackers don't just knock on the front door; they probe your APIs, flood your network with "noise" to distract your team, and attempt to slide through applications and servers using stolen credentials.

To stop these multi-vector attacks, you need the full picture. By using Cloudflare Log Explorer to conduct security forensics, you get 360-degree visibility through the integration of 14 new datasets, covering the full surface of Cloudflare's Application Services and Cloudflare One product portfolios. By correlating telemetry from application-layer HTTP requests, network-layer DDoS and Firewall logs, and Zero Trust Access events, security analysts can significantly reduce Mean Time to Detect (MTTD) and effectively unmask sophisticated, multi-layered attacks.

Read on to learn more about how Log Explorer gives security teams the ultimate landscape for rapid, deep-dive forensics.

The flight recorder for your entire stack

The contemporary digital landscape requires deep, correlated telemetry to defend against adversaries using multiple attack vectors. Raw logs serve as the "flight recorder" for an application, capturing every single interaction, attack attempt, and performance bottleneck. And because Cloudflare sits at the edge, between your users and your servers, all of these events are logged before the requests even reach your infrastructure.

Cloudflare Log Explorer centralizes these logs into a unified interface for rapid investigation.

Zone-Scoped Logs

Focus: Website traffic, security events, and edge performance.

HTTP Requests

As the most comprehensive dataset, it serves as the "primary record" of all application-layer traffic, enabling the reconstruction of session activity, exploit attempts, and bot patterns.

Provides critical evidence of blocked or challenged threats, allowing analysts to identify the specific WAF rules, IP reputations, or custom filters that intercepted an attack.

DNS Logs

Identify cache poisoning attempts, domain hijacking, and infrastructure-level reconnaissance by tracking every query resolved at the authoritative edge.

NEL (Network Error Logging) Reports

Distinguish between a coordinated Layer 7 DDoS attack and legitimate network connectivity issues by tracking client-side browser errors.

For non-web applications, these logs provide visibility into L4 traffic (TCP/UDP), helping to identify anomalies or brute-force attacks against protocols like SSH, RDP, or custom gaming traffic.

Track and audit unauthorized changes to your site's client-side environment such as JavaScript, outbound connections.

Examine how third-party tools and trackers are interacting with user data, which is vital for auditing privacy compliance and detecting unauthorized script behaviors.

Account-Scoped Logs

Focus: Internal security, Zero Trust, administrative changes, and network activity.

Tracks identity-based authentication events to determine which users accessed specific internal applications and whether those attempts were authorized.

Provides a trail of configuration changes within the Cloudflare dashboard to identify unauthorized administrative actions or modifications.

CASB Findings

Identifies security misconfigurations and data risks within SaaS applications (like Google Drive or Microsoft 365) to prevent unauthorized data exposure.

Magic Transit / IPSec Logs

Helps network engineers perform network-level (L3) monitoring such as reviewing tunnel health and view BGP routing changes.

Tracks user actions inside an isolated browser session (e.g., copy-paste, print, or file uploads) to prevent data leaks on untrusted sites

Details the security health and compliance status of devices connecting to your network, helping to identify compromised or non-compliant endpoints.

DEX Application Tests

Monitors application performance from the user's perspective, which can help distinguish between a security-related outage and a standard performance degradation.

DEX Device State Events

Provides telemetry on the physical state of user devices, useful for correlating hardware or OS-level anomalies with potential security incidents.

Being the co-founder of a VC backed startup

Matt Aimonetti

I've now been the co-founder and CTO of Splice for almost 6 years. It's officially the longest place I ever worked at. It's funny because it does feel like we just started yesterday. I decided to write down some unedited thoughts and share them here. I remembered thinking that I'd love to know what it was really like to be a founder, but that's not something I could find on the internet. Hopefully a future entrepreneur will find this post and will get a glimpse at what it is like for me. (sorry for the raw/unedited version, but I felt the post would otherwise stay in my draft folder forever)

I remember so clearly building the first prototype and hearing the excitement in the voice of my co-founder. The first pitches, the first hires, the first people we had let go ☹️ The first arguments, the first fights, the first wins, the first 1:1s, the first dollars we made, the firsts disappointments. The first bugs, the first crashes, the first breakthroughs, the first townhome meetings. The first term sheets, the first board meetings, the first no, the first yes! The first employee departures, the first M&A, the first team off-site. The first exec meeting, the first OKR meetings, the first advisors, the first \$10M paid to creators! Wow, so many memories, I certainly spent too much time thinking about tomorrow and not enough celebrating.

Kruskal-Szekeres spacetime diagram

It's been and still is a wild ride. It feels like we started yesterday, yet it feels like it's been forever. Often people want to know what it is like to be the founder of a startup. They have this glamorous silicon valley image clashing with the image reflected by Silicon valley, the HBO TV show. The reality is hard to describe, it's full of pain, excitement, highs and lows. It's a processes that forces you to reevaluate and reinvent yourself constantly. It's lonely yet you are surrounded by others going through the same thing.

The experience seems to be different for each founder depending on their personality, cultural, social and economical background (and unfortunately also their gender and skin color). Running a VC backed company is like playing a board game where only a few select group of people know the rules. It's not like the rules are secret or anything, it's just that they aren't the rules most of us grow up learning. Those rules seem very foreign at first. Worse, those rules are implicit, they are enforced guidelines that are believed to guide you to a predefined notion of success. Mind you, those rules change with the market. The way you were raising money in 2013, would probably not work in 2019. So you have to stay up to date with those implicit and changing rules. Let's be honest, this part of the system is very unfair, it favors a small group of privileged people who have the social capital to be taught those rules. Talent is everywhere, opportunities are not. I have to admit that this is something that really bothers me. I had the chance to discuss this topic with some VC partners I trust. It's clear they are very

aware of the situation and in some cases try to induce local change. It's clearly not enough but I honestly appreciate their effort when they try.

A thing I also learned the hard way is that you rarely have an opportunity to correct your mistakes. Running a VC backed startup is like strapping a rocket to your chair, it might get you to the finish line faster, but a few small mistakes and you are out (and a crappy rocket might blow up). It's an exciting process, especially if you have little patience, but the g-force you have to endure will take its toll on most of us. Another funny thing about this analogy is that if you don't pass you, you get used to the initial speed and when the rocket is slowing down because you are getting close to orbit, you miss the initial speed and thrill.

Robert Courter in the Bell Rocket Chair.

I'm not a nostalgic person and I don't live in the past regretting things, but I often think of the many mistakes I made so I don't make them again. I wish I could fix those mistakes now that I know the consequences, but the time window is often closed. The best I can do is assume the consequences and not making those mistakes again as well as helping others who are going through the same challenges. As much as we love saying that we embrace failure, ask a founder how they are doing, the most common answer is probably a version of: crushing it! Unfortunately, the reality is probably very different. Self-doubt is probably the best feeling describing being a founder. But we don't want people to doubt us, otherwise investors won't support us, we won't be able to hire top talent and we will fail to make the "X under X" list. So many founders at all stages struggle with depression, substance abuse, deep feeling of loneliness. It's quite sad to see peers going through that while the media love glorifying or at least "glamourising" our self destructive behaviors.

VCs are not your friends, but they aren't your enemy either. Like you, they are doing their job and their job is to make sure you exponentially increase the financial value of their investment. You can have a close relationship with some investors, they can provide you with great insights and be supportive. But their primary duty is to their firm and they LPs. They are being kept accountable and have to show results. This is not a secret, it's their duty. It's too easy to blame VCs when things go bad. Don't get me wrong, there is a fair share of bad VCs out there, but remember that you are the one who agreed to tape a rocket to your chair. They provided you with the rocket, they didn't force you, being VC backed is certainly not for everyone or every project. You told them you'd manage the pressure, the growth etc.. But did you make sure you were aligned on the conditions attached to using the rocket? If they don't think you will make it, don't be mad if they don't want to add more gas in the rocket or suggest you stop the race. If they feel that you can make it but you are slowing down too quickly on purpose, don't be mad if they pressure you to add more gas and take bigger risks. That's part of the game.

AI Security for Apps is now generally available

Liam Reese · Cloudflare Blog

Cloudflare's AI Security for Apps detects and mitigates threats to AI-powered applications. Today, we're announcing that it is generally available.

We're shipping with new capabilities like detection for custom topics, and we're making AI endpoint discovery free for every Cloudflare customer—including those on Free, Pro, and Business plans—to give everyone visibility into where AI is deployed across their Internet-facing apps.

We're also announcing an expanded collaboration with IBM, which has chosen Cloudflare to deliver AI security to its cloud customers. And we're partnering with Wiz to give mutual customers a unified view of their AI security posture.

A new kind of attack surface

Traditional web applications have defined operations: check a bank balance, make a transfer. You can write deterministic rules to secure those interactions.

AI-powered applications and agents are different. They accept natural language and generate unpredictable responses. There's no fixed set of operations to allow or deny, because the inputs and outputs are probabilistic. Attackers can manipulate large language models to take unauthorized actions or leak sensitive data. Prompt injection, sensitive information disclosure, and unbounded consumption are just a few of the risks cataloged in the OWASP Top 10 for LLM Applications .

These risks escalate as AI applications become agents. When an AI gains access to tool calls—processing refunds, modifying accounts, providing discounts, or accessing customer data—a single malicious prompt becomes an immediate security incident.

Customers tell us what they're up against. "Most of Newfold Digital's teams are putting in their own Generative AI safeguards, but everybody is innovating so quickly that there are inevitably going to be some gaps eventually," says Rick Radinger, Principal Systems Architect at Newfold Digital, which operates Bluehost, HostGator, and Domain.com.

What AI Security for Apps does

We built AI Security for Apps to address this. It sits in front of your AI-powered applications, whether you're using a third-party model or hosting your own, as part of Cloudflare's reverse proxy . It helps you (1) discover AI-powered apps across your web property, (2) detect malicious or off-policy behavior to those endpoints, and (3) mitigate threats via the familiar WAF rule builder.

Discovery — now free for everyone

Before you can protect your LLM-powered applications, you need to know where they're being used. We often hear from security teams who don't have a complete picture of AI deployments across their apps, especially as the LLM market evolves and developers swap out models and providers.

AI Security for Apps automatically identifies LLM-powered endpoints across your web properties, regardless of where they're hosted or what the model is. Starting today, this capability is free for every Cloudflare customer, including Free, Pro, and Business plans.

Cloudflare's dashboard page of web assets, showing 2 example endpoints labelled as cf-llm

Each prompt is run through multiple detection modules for prompt injection, PII exposure, and sensitive or toxic topics.

Liam Reese

Discovering these endpoints automatically requires more than matching common path patterns like /chat/completions . Many AI-powered applications don't have a chat interface: think product search, property valuation tools, or recommendation engines. We built a detection system that looks at how endpoints behave , not what they're called. To confidently identify AI-powered endpoints, sufficient valid traffic is required.

AI-powered endpoints that have been discovered will be visible under Security → Web Assets , labeled as cf-llm . For customers on a Free plan, endpoint discovery is initiated when you first navigate to the Discovery page . For customers on a paid plan, discovery occurs automatically in the background on a recurring basis. If your AI-powered endpoints have been discovered, you can review them immediately.

AI Security for Apps detections follow the always-on approach for traffic to your AI-powered endpoints. Each prompt is run through multiple detection modules for prompt injection, PII exposure, and sensitive or toxic topics. The results—whether the prompt was malicious or not—are attached as metadata you can use in custom WAF rules to enforce your policies. We are continuously exploring ways to leverage our global network, which sees traffic from roughly 20% of the web , to identify new attack patterns across millions of sites before they reach yours.

New in GA: Custom topics detection

The product ships with built-in detection for common threats: prompt injections, PII extraction , and toxic topics . But every business has its own definition of what's off-limits. A financial services company might need to detect discussions of specific securities. A healthcare company might need to flag conversations that touch on patient data. A retailer might want to know when customers are asking

new rails plugin mimetype_fu

Matt Aimonetti

mimetype_fu/ is a new plugin I just wrote. It's simple and it can be really useful if you need to get the mime type of a file already on your server.

During one of my project, I add to migrate old assets from a legacy system to a new Rails app. The new app uses attachment_fu and even though techno weenie did an amazing job, attachment_fu validation is based on the content type. A_fu gets the content type coming from the CGI query.

Unit test has a helper faking this process but in real life, if you use a Flash uploader (Flash doesn't give you the proper mime type/content type) or if you want to migrate files, the attachment_fu validation won't work for you.

The solution is simple: mimetype_fu/

mimetype_fu/ extends the File class and is really easy to use:

```
File.mime_type?(@file)
```

Check it out <https://code.google.com/p/mimetype-fu/>

Expect a post showing how a ninja would use the mimetype_fu / attachment_fu combo :)

fun with macruby

Matt Aimonetti

To be ready for 2010, I'm taking some time off relaxing and spending time with my family in Florida.

During my free time, I've been reading, catching up on movies and TV shows and worked on the MacRuby book that I am writing for O'Reilly.

I wrote a bunch of small apps, played with various APIs and every single time I was amazed by all the goodies Apple makes available to developers. My most recent discovery is very simple but I wanted to share it with you.

I often type text in English, French and Spanish and I even mix the languages from time to time. SnowLeopard comes with a great spellchecker that auto detects the language I'm typing in and is most of the time correct. It's a very impressive feature and I was wondering if, as a MacRuby developer, I could use one of Apple's lib to detect what language is being used. I dug through the documentation but didn't find anything. I started looking at some header files and found the API to use :)

```
framework 'Foundation'
class String def language
  CFStringTokenizerCopyBestStringLanguage(self,
    CFRangeMake(0, self.size))
end end
```

```
puts "Bonne année!".language # => "fr"
puts "Happy new year!".language # => "en"
puts "¡Feliz año nuevo!".language # => "es"
puts "Felice anno nuovo!".language # => "it"
puts "أعياد سعيدة".language
```

```
# => "ar" puts "العيد سعيد".language # => "ja"
```

The documentation says that the result is not guaranteed to be accurate and that typically 200-400 characters are required to reliably guess the language of a string. (CFStringTokenizer Doc)

Probably not the most useful piece of code, but really cool none the less :)

Happy new year!

become better at merb

Matt Aimonetti

I hear a lot of people asking for better Merb documentation as they are learning Merb. I remember struggling with the very same problem when I started using Rails a few years ago.

The good news is that the core team is working hard on improving the state of documentation.

You have the Merb Book on which I am actively work on with a team of 20 people (writers, editors and translators).

You also have the documentation browsing application we are developing to let you browse documentation based on where you are in your app.

And of course, you have the 2 beta books (Merb in Action and the Merb way) plus two other ones in the process of being written (Beginning Merb and Merb: What You Need to Know).

Finally, you also have other resources like the Merb peepcode .

The point being that we are actively working on improving the overall state of documentation. However, if you wish to learn more about the internals, how to some more advanced stuff like writing plugins, building your own stack, knowing when to use what component, I would suggest you attend the training class that Yehuda Katz (Merb lead developer) and myself are giving in Phoenix, AZ, January 19-21 .

This would give you a rare opportunity to spend time

with some of people directly involved with the day to day development of the framework as well as other like minded developers.

Website: <https://merbclass.com> Curriculum: <https://merbclass.com/curriculum.html> Date: Jan 19-21, 2009 Location: Phoenix, AZ

attachment_fu updated

Matt Aimonetti

I recently bugged Rick Olson so much about attachment_fu that he gave me SVN access to fix few bugs.

Rick being really busy with ActiveRecord he didn't spend too much time maintaining attachment_fu .

On my side of things, I've been using attachment_fu on a lot of projects and I've been fixing bugs and adding new features.

My first contribution to attachment_fu is a fix for the ImageScience processor.

Attachment_fu is very flexible and let you use your favorite image processor:

RMagick based on ImageMagick and GraphicsMagick (known to leak memory and being a pain to setup)

minimagick based on ImageMagick

ImageScience based on FreeImage .

Like many rubyists, I like ImageScience for its simplicity and efficiency. However, attachment_fu had few problems when being used with ImageScience .

File sizes for thumbnails were not saved correctly in the database. Fixed

Thumbnails based on a gif files were not processed properly. So, this was the big problem. FreeImage has issues dealing with resizing gif files because of the gif palette limitation (256 colors). to avoid this problem thumbnails of

gif files are converted to png. However the thumbnail content type info in the database was not saved properly. That's now fixed.

Because of the gif bug reported above, any thumbnail link was broken since it was trying to link to the thumbnail version with a gif extension instead of a png one. Fixed

I also fixed a small bug related to S3 storage and the fact that a_fu had issues loading the config file. (Fixed)

I'll also try be able to add some of the S3 features I've been working on. As well as maybe enhancing the validation process.

In the mean time, you might want to read this blog post about better validation with attachment_fu.

If you are heavily using attachment_fu or starting using it and think that a google group would be great idea, please let me know in the comment and I'll try to convince Technoweenie that we need to set that up :)

Ooohh I almost forgot, if you fixed some bugs you found while using a_fu, please contact me so we can get a_fu bug free.

Ultrahuman's New Ring Pro Is Finally Available in the US

Beth Skwarecki · Lifehacker

Ultrahuman finally has a smart ring back on the U. S. market, and it's the Ring Pro, the new model with the feature-packed charging case. Pre-orders are now open, with early bird prices starting at \$349. The regular price will be \$479 for the ring plus case.

The ring will begin shipping May 15. The company is using a tiered model for early bird pricing, where the price bumps up after a certain number of rings are sold. Here's the schedule:

\$349 — first 1,000 orders (\$130 off)

\$389 — next 2,000 orders (\$90 off)

\$429 — next 3,000 orders (\$50 off)

\$449 — until pre-orders close.

Ultrahuman's smart rings have been off the U. S. market since October 2025, due to legal action by Oura. As I detailed [here](#), Oura sued several other smart ring makers over patent infringement and won a judgment by the U. S. International Trade Commission. As a result, most other smart ring companies had to leave the U. S. market or reach a deal with Oura. Ultrahuman found a third option—continuing to develop a ring they already had in the works, betting that the ITC would rule its design is different enough that it does not infringe Oura's patent.

That bet seems to have paid off. U. S. Customs and Border Protection has officially signed off on the new design, allowing it to be sold here. Ultrahuman's older designs, like the Ring Air I reviewed last year, are still banned.

What's new with the Ultrahuman Ring Pro

I wrote more about the Ring Pro [here](#). It's got a faster processor with "on-chip machine learning," a 15-day battery life, and the ability to hold 250 days of data if you're unable to sync your ring for extended periods. The physical shape is similar to the Ring Air (no sensor bumps, but there is a flattened area at the bottom of the interior), and it has break points for easier removal of the ring in an emergency.

The cleverest feature of the Ring Pro is its charging case, which can store data, flash the firmware, and deliver haptic alerts. It also has a speaker for alarms and for sounding when you use the "find my case" feature. The charger can charge wirelessly.

Ultrahuman has since announced that you'll have the option to buy the ring packaged with the charging case for \$479, or with just a "mini charger" (presumably similar to the Ring Air's charger) for \$399. The case will be available separately for \$100.

ANALYSIS

The New Echo Dot Max Is Its Lowest Price Ever for the Amazon Spring Sale

Khamosh Pathak · Lifehacker

We may earn a commission from links on this page. Deal pricing and availability subject to change after time of publication.

The Echo Dot Max is Amazon's "Pro" tier smart speaker, designed for people who want better sound quality in a relatively small package and at an affordable price. This upgraded smart speaker is powerful, with a spherical body and a concave cutout for volume and mute controls, surrounded by an LED ring (yes, it looks like the Death Star).

Released last October, the Echo Dot Max is currently at its lowest price ever during Amazon's Big Spring Sale . You can get one for \$74.99 instead of the usual \$99.99.

\$74.99 at Amazon

\$99.99 Save \$25.00

\$74.99 at Amazon

\$99.99 Save \$25.00

The new Echo Dot Max is a step up from the Echo Dot, and compared to the 2022 model, it features three times louder bass (thanks to a new 0.8-inch tweeter and a 2.5-inch woofer). It's powered by Amazon's custom Za3 chip with a built-in AI Accelerator, meaning it's built for the Alexa+ AI service .

The Max also features a built-in smart home hub with support for Matter, Zigbee, and Thread Border Router; it can also work as an eero extender for your mesh Wi-Fi router. There's also a temperature sensor and ultrasonic presence detector, so you can use it to trigger Alexa routines as soon as someone walks into the room.

Tech wise, the Echo Dot Max supports Wi-Fi 6E for faster speeds, Bluetooth 5.3, lossless high-definition audio, and automatic room adaption. It's also easy to pair two Echo Dot Max speakers for a stereo home theater setup with any compatible Fire TV devices. Echo Dot Max comes in three colors: Graphite, Glacier White, and Amethyst. All three are available on discount right now.

PCMag gave the Echo Dot Max an "Outstanding" 4.5 star rating , along with an Editor's Choice award, noting, "the Amazon Echo Dot Max stands out as a strong successor to the fourth-generation Echo, comfortably occupying the \$100 sweet spot for smart speakers." And what's better than a \$100 sweet spot? A \$75 sweet spot.

Our Best Editor-Vetted Amazon Big Spring Sale Deals Right Now

Apple AirPods Pro 3 Noise Cancelling Heart Rate Wireless Earbuds

— \$199.00

(List Price \$249.00)

Apple iPad 11" 128GB A16 WiFi Tablet (Blue, 2025)

— \$299.00

(List Price \$349.00)

Samsung Galaxy Tab A11+ 128GB Wi-Fi 11" Tablet (Gray)

— \$209.99

(List Price \$249.99)

Sony WH1000XM6- Best Wireless Noise Canceling Headphones

— \$398.00

(List Price \$459.99)

Apple Watch Series 11 (GPS, 42mm, S/M Black Sport Band)

— \$299.00

(List Price \$399.00)

Blink Video Doorbell Wireless (Newest Model) + Sync Module Core

— \$35.99

(List Price \$69.99)

Fire TV Stick 4K Max Streaming Player With Remote

— \$34.99

(List Price \$59.99)

Amazon Kindle Colorsoft 16GB 7" eReader (Black)

— \$169.99

(List Price \$249.99)

Deals are selected by our commerce team

All three are available on discount right now.

Khamosh Pathak

The names of birds, part 1

ericlippert · Fabulous Adventures in Coding (Eric Lippert)

For the next part in my Bean Machine retrospective to make sense I'll need to make a short digression. In looking back on the almost 20 years I've been blogging, it is surprising to me that I've only briefly alluded to my appreciation of combinatory logic. In the next couple of episodes, I'll do a quick introduction based on the delightful book that introduced it to me: *To Mock A Mockingbird*, by the late Raymond Smullyan.

Imagine a forest containing some birds, possibly finitely or infinitely many. These are unusual birds. When you call the species name of a bird in the forest to a bird in the forest, it calls one back to you. Maybe the same, maybe different, but when you tell a bird the name of a bird, it names a bird back to you. There might be a Red Capped Cardinal in the forest and when you call out Great Blue Heron to it, it calls back Belted Kingfisher.

(Photos by me; click for higher resolution.)

We will notate "I called Q to P and got response R" as $PQ = R$. If I then called out S to R and R responded with T, we'll notate that as $PQS = RS = T$. We'll use parentheses in the obvious way: $PQS = (PQ)S$ and this might be different from $P(QS)$. The latter is "I called S to Q, and then called Q's response to P". I'll use capital letters to represent specific bird names, and small letters to represent variables.

The question we're considering here is: under what circumstances will a bird call back the same name you called to it? That is, for a given bird y , under what circumstances does $yx = x$?

Smullyan calls birds with this relationship "fondness" — that is, "y is fond of x" means that $yx = x$. If y is fond of x then x is said to be a "fixpoint" of y.

A forest is said to be "compositional" if for every pair of birds (a, b) — a and b can be the same or different — there is a bird c in the forest such that $cx = b(ax)$ for all x. That is, if we call any name x to a, and then call its response to b, we get the same result as simply calling x to c. c is the composition of "call x to a, and then call that response to b".

A mockingbird is the bird M with the property that $Mx = xx$. That is, for any bird name x, M tells you what x's response is to its own name. (The attentive reader will note that we have not said what MM is, but at least we know from the definition that $MM = MM$, so give me some credit for consistency at least.)

Theorem : if a compositional forest contains a mockingbird then every bird in the forest is fond of at least one bird.

Proof: Try to prove it yourself; scroll down for the solution.

.
.
.
.
.
.
.

Let p be any bird. Let c be the bird that composes p with M. Therefore $cx = p(Mx)$ for all x.

In particular, that's true for c, so $cc = p(Mc)$. But $Mc = cc$, thus $cc = p(cc)$, and we've found a bird that p is fond of. Since p was any bird, every bird in the forest is fond of at least one bird. Or, put another way, every bird in a compositional forest with a mockingbird has at least one fixpoint.

Next time on FAIC: we'll take a look at a few more interesting birds, and then discuss why this whimsy is relevant to compilation before getting back to Bean Machine.

How Robert F. Kennedy Jr.'s Vaccine Agenda Risks a Resurgence of Deadly Childhood Plagues

Patricia Callahan · ProPublica

Dr. Adam Ratner hovered over a gravely ill infant in a New York City intensive care unit on a grim day in 2022. The 3-month-old girl spiked a fever two days earlier and had become lethargic. Soon she was having seizures and struggling to breathe.

She didn't register Ratner's towering frame or the bright hospital lights. Her eyes stared up and to the right, eerily frozen.

He ran his hand over the soft spot on her head, which should have been flat. Instead, it bulged, a sign that too much fluid was building up inside her skull.

The baby's life was in danger, and Ratner needed to figure out why. He worried the culprit was bacterial meningitis, an infection of the membranes that protect the brain.

What came back on her lab tests was something out of the history books.

The infant's meningitis was caused by invasive *Haemophilus influenzae* type b, or Hib, a type of bacteria that used to kill nearly 1,000 children a year in the U. S. A shot introduced in the late 1980s was so effective that Ratner, a veteran pediatric infectious disease doctor, was among the generations of physicians who had never seen a case. But the baby's parents, Ratner learned, had chosen not to vaccinate her.

Disheartened, he told his colleagues, "This should be a never event."

It wasn't. The following year, Ratner treated another infant with Hib, then another, each of them unvaccinated. Two went home, but one had to be discharged to a rehabilitation facility. That 5-month-old boy had huge black pupils that didn't respond to light, and he needed a ventilator to breathe. Ratner and his colleagues noted an "absence of brain stem reflexes," indicating severe damage.

The U. S. government took a half century to build a vaccination system that shielded children from such a fate. Its success depended on two fundamental pillars: parents trusting in vaccines and children having access to them. Both are now in peril, thanks in no small part to the man steering America's health policy.

Health and Human Services Secretary Robert F. Kennedy Jr., who founded an antivaccine group and once likened the immunization of children to a holocaust, is transforming a government that long championed the lifesaving benefits of shots into one that spreads doubts about their safety here and abroad.

Kennedy is also considering changes that could prompt the few companies that make vaccines for American kids

to abandon the U. S. market, leaving parents who want the shots unable to get them.

The threat to vaccine access reaches across the globe after Kennedy yanked the government's \$1.6 billion pledge to the aid group that provides shots for the world's poorest children. For decades, the U. S. had funded such work not just as a humanitarian mission but as a way to keep Americans safe from unchecked contagions.

Kennedy's efforts to reshape vaccine policies have been well chronicled, but ProPublica wanted to take a broader look at how the changes might affect Americans' health in the years to come.

We found that long-forgotten plagues have roared back, killing and maiming children in parts of the world where access to vaccines or trust in them faltered. What seemed like subtle changes to a country's vaccine policies had disastrous consequences years later.

Even in places that offer highly advanced health care, doctors have felt impotent trying to undo the damage when these horrors return. Modern medicine can't reverse paralysis from polio. Surgeons can intervene when a baby is born blind, deaf and with heart defects after being exposed to rubella in the womb, but the child is still likely to face a life shaped by disability.

ProPublica reviewed hundreds of studies on vaccines and outbreaks of the diseases they prevent and interviewed more than three dozen people who have worked on U. S. immunization programs here and abroad, dating back to the days of smallpox. Some had never spoken publicly about their experiences.

They shared a pit-of-the-stomach dread that American children will end up fighting for their lives against infections that have long been preventable.

"I think there always was a worst-case scenario," said Dr. Melinda Wharton, who retired last September after more than three decades leading immunization programs at the Centers for Disease Control and Prevention. "I don't think I imagined it could or would be this bad."

This week, Kennedy's agency indicated it planned to appeal a federal court ruling that halted, at least temporarily, some of his changes. Among those was the decision to drop six diseases from the routine childhood immunization schedule.

HHS declined to make Kennedy available for an interview. In an emailed response to detailed questions, HHS spokesperson Andrew Nixon said that the agency has not limited access to or insurance coverage for vaccines. During the prior administration, federal health agencies "told the public that questioning vaccine policy was off limits," Nixon said. "That posture contributed to a collapse in trust in U. S. health care."

Sunday at La Bombonera

Juan Villoro · *The Paris Review Blog*

Clásicos are like Christmas for football. In these high-tension matches between fierce rivals, expectation almost always outstrips results. For months, fans visualize goals with the unrealistic yearning of a child who hopes for a new PlayStation from Santa Claus in exchange for a few cookies left out for his tired reindeer.

For me, the Superclásico between Buenos Aires's Boca Juniors and River Plate on May 4, 2008, was preceded by thirty-four years of anticipation. In 1974 I went to the Estadio Monumental to see River-Boca, but I had never been to the reverse fixture in La Bombonera, that exceptional stadium that should have been examined by Elias Canetti in *Crowds and Power*.

The wait had charged the occasion with so much emotion that it was almost a shame it actually had to take place. Friends from Mexico, Colombia, and Spain had all similarly circled the date of May 4—the Argentine derby appeals not only to those who sleep in shirts emblazoned with the Quilmes beer logo but to an entire global tribe.

Like Everest or the Mona Lisa, the fame of Boca's stadium is impossible to deny—look no further than the crowds of tourists who come to snap pictures. But does it really represent the pinnacle of footballing passion?

I spoke about the Superclásico with the taxi driver who picked me up at the Ezeiza airport and he replied with indignation: "But we hate each other more!" He came from Rosario and was referring to the bad blood between his hometown clubs Newell's Old Boys ("the Lepers") and Rosario Central ("the Swine"). On the drive he told me about his family's marvelous wrath and the betrayal of his aunt Teresita, the heretic who refused to support the Swine. At the core of his story was the issue of rancor: on its biggest days, football comes down to contempt, and nobody hates each other more than Swine hate Lepers. In his opinion, the lesser rivalry between Boca-River was inflated by the press. The driver summed up his argument with theological flair: "God is everywhere, but performs his tricks in Buenos Aires."

Don't Worry: It's Just the Earth Shaking

On April 16, my friend Daniel Samper Pizano, a renowned Colombian journalist, organized a dinner in Madrid to prepare for the Superclásico. The final of the Copa del Rey had just concluded between Valencia and Getafe, but we weren't interested. We preferred to talk about the football of the future, which was to say the upcoming game on May 4. The other guest justified our interest; after the meal, Jorge Valdano told us about the first time he faced Boca as a visiting player. Lacing up his boots in the locker room, he said, everything around him seemed to be moving. One of the veterans came over to reassure him: "It's not you, pibe, it's the stadium." Playing at La Bombonera means overcom-

ing an arena about to collapse under the weight of its own passion. No other ground imposes itself so forcefully upon its visitors.

In his stupendous book *Boquita*, Martín Caparrós reminds us that it was Argentina where fans were first baptized as "the twelfth man." Accustomed to adversity, we Mexicans consider the scoreline a suggestion we can ignore. Argentine fans, on the other hand, seek to alter the result using three basic tactics: holding their breath, insulting their rivals, and singing love songs. It's no accident that one of the most prominent barras bravas is called *La Doce*, or "the twelfth." These ultras aren't there to merely watch a game, but to participate in it through their shouts.

For years, magical realism has gone missing from Latin American literature, finding refuge in aviation instead. To fly across the continent is to endure a saga of detours, delays, and strange schedules that lead to a parallel reality. Maybe air traffic control is cheaper in the early morning and this determines the itineraries of the hemisphere. At any rate, I found myself on a four-hour predawn flight from Bogotá to Buenos Aires. Anyone without the meditative powers of a yogi arrived at the destination a zombie. Among the many surprises to come that day, I was made to experience May 4 in an altered state of time.

The relationship between football and aviation is no trifle—the Copa Libertadores will never be truly competitive until the continent modifies its fixture calendar and flight paths. When I was a kid, I used to take medications with labels that read "Shake well before using." Thanks to a miserable night on the plane, I arrived at the Superclásico in a state of total agitation.

Entering the stadium was likewise an extreme sport. I was lucky to be accompanied by my friend Leo Tarifeño, a diehard River fan who had sworn to never set foot in La Bombonera.

Leo is convinced that Argentines live for confrontation, eagerly disregard established rules, object impulsively, and justify themselves only through negativity, taking issue with anything they don't agree with. After asserting this theory, he put it into practice. When I praised the chanting of the Boca fans, he replied, "Deep down, their joy is bitter."

Being with Leo was the opposite of having a human shield. The route to the stadium was blocked, so we made our way through a wasteland thick with the invigorating smoke from street stalls selling choripán. Our barren surroundings gradually became a funnel. Police barricades flanked us on either side. We continued on until someone—an invisible leader far ahead—committed a grave mistake. The crowd was pushed back by the sound of riot guns and retreated to a checkpoint, where Leo and I asked for directions to the press box. An officer waved his hand as if trying to hypnotize us. We "understood" to make our way to the other side of a roundabout, but instead we ducked into the first alley we came across. Another crowd pulled us in before being repelled again by the riot guns. Everyone ran en masse to

These Are the Best Deals on Home Gym Equipment I've Seen During Amazon's Big Spring Sale

Meredith Dietz · Lifehacker

We may earn a commission from links on this page. Deal pricing and availability subject to change after time of publication.

I dream of building out a garage gym one day, but for now, I'm content with my "workout corner" in my apartment bedroom. Hey, a home gym can look like anything. Whether you're eyeing a smart rowing machine, a compact massage gun for post-workout recovery, or just a clean set of hex dumbbells with a rack to keep your space organized, there's something worth buying during Amazon's Big Spring Sale right now.

The best deals on at-home cardio equipment

I know how quickly cardio equipment can turn into major investment of your time, money, and space. That's why I recommend equipment that can fold up to store vertically against the wall, so your treadmill doesn't just become a fancy coat rack.

Original Peloton Bike , \$1,295, down from \$1,695

Hydrow Origin Rowing Machine , \$1,895, down from \$2,195

NordicTrack T 5 Treadmill , \$479, down from \$599

Peloton Exercise Cross Training Bike

\$1,295.00 at Amazon

\$1,695.00 Save \$400.00

\$1,295.00 at Amazon

\$1,695.00 Save \$400.00

The best deals on at-home strength training essentials

In my experience, adjustable dumbbells are the single highest-impact upgrade you can make to a home gym. But sometimes you just want the real thing, without dealing with pins or spin locks or any sort of fuss. In that case, you definitely want a quality storage rack keeps your gym organized, your equipment accessible, and your floor clear.

PowerBlock Elite EXP Adjustable Dumbbells , \$339.15, down from \$449

CAP Adjustable Dumbbells , \$249.99, down from \$299.99

CAP Rubber Coated Dumbbell Weight Set , \$159.99, down from \$189.99

TRX Resistance Band Set , \$148.71, down from \$174.95

PLKOW storage and weight rack , \$79.98, down from \$129.99

PowerBlock Elite EXP Adjustable Dumbbells (Pair, 5-50lb)

\$339.15 at Amazon

\$449.00 Save \$109.85

\$339.15 at Amazon

\$449.00 Save \$109.85

The best deals on muscle recovery and stretching tools

I know I've personally got my eye on massage guns and foam rollers to help me survive my upcoming marathon training.

Bob and Brad Q2 Ultra Mini Massage Gun , \$78.82, down from \$99.99

1-Inch Thick Yoga Mat , \$29.99, down from \$36.99

TriggerPoint Grid 1.0 Foam Roller , \$27.99, was \$39.99

Bob and Brad Q2 Ultra Mini Massage Gun

\$78.82 at Amazon

\$99.99 Save \$21.17

\$78.82 at Amazon

\$99.99 Save \$21.17

While we're at it, plenty of other gadgets could technically fall into the realm of home gym expenses. Check out more of the best deals we've found on fitness wearables , on headphones and earbuds , and all other tech during Amazon's Big Spring Sale here .

Our Best Editor-Vetted Amazon Big Spring Sale Deals Right Now

Apple AirPods Pro 3 Noise Cancelling Heart Rate Wireless Earbuds

— \$199.00

(List Price \$249.00)

Apple iPad 11" 128GB A16 WiFi Tablet (Blue, 2025)

— \$299.00

(List Price \$349.00)

Samsung Galaxy Tab A11+ 128GB Wi-Fi 11" Tablet (Gray)

— \$209.99

(List Price \$249.99)

Sony WH1000XM6- Best Wireless Noise Canceling Headphones

Akari Brings Authenticity to New York's Washed Sauna Scene

Adrian Madlener · Design Milk Architecture

There are strong links between Japan and the Nordic region (a grouping of countries that includes much more than just Scandinavia). That's especially true when it comes to sauna culture. Though the word is Finnish, translating as bathhouse, the concept of pumping steam through enclosures firmly roots in Neolithic Greenland and Newfoundland. It's been an integral part of daily life across these, especially frigid, Northern reaches for eons.

In Japan, sentō and onsen facilities have been around since the 1200s. While the former is generally defined as a communal bathhouse, the latter is specified as geothermal hot springs that are encircled by traditional inns.

Gaining popularity as an alternative, far less harmful, type of third place, saunas and other types of bathhouses have cropped up across bustling metropolises like London and New York. While many tout unique amenities—some more gimmicky than others—few of these highly publicized venues hold true to the Japanese traditions mentioned before.

Cue Akari , two immersive saunas in Williamsburg and Greenpoint, Brooklyn that are as comprehensive as they are unassuming. There's no need for flashy light shows or poolside DJ-sets at these quaint, smartly outfitted wellness spaces. Comprising dry saunas, cold plunges, and quiet relaxation areas, both membership—non reservation—based venues champion the idea that New Yorkers should slow down more often and go to the bathhouse in order to reset rather than party. While the entirely renovated Williamsburg location features an incredibly rare outdoor cold plunge pool, the Greenpoint one was imagined in partnership with celebrated thermal consultant Kiel Moe .

International architecture firm Stroop Design staged the latter with intentional detailing. The aim was not to overpower the restorative act of saunaing with bold finishes or sculptural elements but to succinctly facilitate the experience with a consistent design vocabulary. Many of these accents—2 by 2 inch light green tiles—harken back to traditional sentōs. Tabletops laminated in washi paper by master craftsperson Hanato Watar were repurposed as wall art throughout.

Sparingly placed furnishings and fixtures are modestly modern, solidly made but light, reflecting the brand's name Akari, which translates as light in Japanese. The idea here is to make saunaing a daily ritual—one that has both physical and mental benefits—and remove any of the social expectation or architectural fanfare that might make it feel like a chore.

To learn more about the creative synergy between Akari Sauna and Stroop Design, visit akarisauna.com and stroop.design .

The Titanium Apple Watch Series 10 Is Just \$449 During the Amazon Big Spring Sale

Jake Peterson · Lifehacker

We may earn a commission from links on this page. Deal pricing and availability subject to change after time of publication.

Flagship Apple Watches tend to start at \$399, but the price only goes up from there. If you want a larger size, that's a slight bump in cost; if you want cellular connectivity, that's an even larger hit. But the titanium Apple Watch—the “premium” material, if you will—is the most expensive upgrade outside of the Ultra line. These models typically start at \$699, but cost \$749 for the larger size, and a whopping \$799 if you want the stainless steel band. Right now, you can get a titanium Apple Watch for as low as \$449 during Amazon's Big Spring Sale—so long as you're okay with using last generation's watch.

Apple Watch Series 10 (Titanium)

\$449.00 at Amazon

\$749.00 Save \$300.00

\$449.00 at Amazon

\$749.00 Save \$300.00

Here's the deal: Amazon currently has the titanium versions of the Apple Watch Series 10 on discount for the Big Spring Sale. You can get the Natural or Slate models with rubber sport bands for \$449, or if you want the stainless steel band, you can get the watch in gold for \$499. That's a pretty great price, considering the watch once retailed for \$799.

Now, this is the Series 10. The Series 11 is Apple's newest flagship watch. If you ask me, however, that isn't a big deal. The Series 10 is still the best Apple Watch for most people, since it has most of the Series 11's features, while often being available at a discount. They both have large displays that can reach 2,000 nits of brightness; they both have heart rate and sleep tracking with alerts for each; they both support Emergency SOS, and are water resistant up to 50m; they both have the same S10 chip, and support the same gestures; and both have Precision Finding, so you can use your iPhone if you misplace it.

Really, the key advantages of the Series 11 are few, but noteworthy: The battery life is better, up to 24 hours instead of 18 hours; you get 38 hours in Low Power Mode versus 36 hours; and while both watches support fast charging (0-80% in 30 minutes), the Series 11 can get eight hours of battery life in 15 minutes of charge time. The Series 11 also supports 5G, if you opt for the cellular plan. That's largely it, however. Other than that, the Series 10 is virtually the same watch. Amazon does have the Series 11 discounted at this time, starting at \$649. You just need to ask yourself whether those extra features are worth \$200.

Our Best Editor-Vetted Amazon Big Spring Sale Deals Right Now

Apple AirPods Pro 3 Noise Cancelling Heart Rate Wireless Earbuds

— \$199.00

(List Price \$249.00)

Apple iPad 11" 128GB A16 WiFi Tablet (Blue, 2025)

— \$299.00

(List Price \$349.00)

Samsung Galaxy Tab A11+ 128GB Wi-Fi 11" Tablet (Gray)

— \$209.99

(List Price \$249.99)

Sony WH1000XM6- Best Wireless Noise Canceling Headphones

— \$398.00

(List Price \$459.99)

Apple Watch Series 11 (GPS, 42mm, S/M Black Sport Band)

— \$299.00

(List Price \$399.00)

Blink Video Doorbell Wireless (Newest Model) + Sync Module Core

— \$35.99

(List Price \$69.99)

Fire TV Stick 4K Max Streaming Player With Remote

— \$34.99

(List Price \$59.99)

Amazon Kindle Colorsoft 16GB 7" eReader (Black)

— \$169.99

(List Price \$249.99)

Deals are selected by our commerce team

few changes in merb edge

Matt Aimonetti

I thought I should share few changes that may affect your apps if you want to upgrade to Edge.

`merb_helpers` which was previously in `merb-plugins` now moved to `merb-more` and got renamed `merb-helpers`. What that means for you is that you probably want to install `merb-helpers` and change the reference in your `init.rb` from dependencies “`merb_helpers`” to dependencies “`merb-helpers`”

People started reporting problems with templates not being reloaded in dev mode etc.. The reason is that we made some changes to the `config/environments/development.rb` file few weeks ago and people did not notice. Here is the new generated `development.rb` file. Note the following interesting change:

```
c[:reload_templates] = true
```

You will need to update this setting if you want Merb to auto reload your templates.

Also, if you didn't read it yet, go check on Yehuda's explanation of Merb new master process.

Finally, I have a bit of bad news. We were hoping to release Merb 1.0 final during MerbCamp next week end. Unfortunately it looks like we will only release 1.0RC.

The reason behind this choice is simple, we have been adding a lot of features, fixed a lot of bugs and stabilized the API. However we need more feedback from users to be confident

enough to release a final 1.0 release.

rspec bundle for textmate

Matt Aimonetti

I recently realized that a lot of people using RSpec didn't use the great TextMate bundle designed for them.

You might not think you need an extra bundle, or maybe you never used a bundle ever and don't know why you would try one. I'll try to change your mind, let's look at some screenshots first:

This is what you will want to see before checking in your code:

This is what you will see before you write your code:

This is how you will run your specs:

This is all you get for free to save you some precious minutes every day:

How to install?

Easy, checkout the file from RSpec repository:

```
svn co svn://rubyforge.org/var/svn/rspec/trunk/RSpec.tmbundle
```

Double click on the file and that's it! You are ready to go. (you would obviously need RSpec plugin or gem).

Thank you to all the RSpec team for creating such a great tool.

Note the following interesting change: `c[:reload_templates] = true` You will need to update this setting if you want Merb to auto reload your templates.

Matt Aimonetti

BRIEFS

IN BRIEF

The Hacker News, The Hacker News (cybersecurity): A critical security flaw impacting Langflow has come under active exploitation within 20 hours of public disclosure, highlighting the speed at which threat actors weaponize newly published vulnerabilities. The security defect, tracked as CVE-2026-33017 (CVSS score: 9.3), is a case of missing authentication combined with code injection that could result in remote code execution. “The POST /api/v1

Bonface Landi, Inhabitat: Soil erosion is an age-long predicament that always shows up whenever humans embark on any land use activity. Whether it is agriculture, construction or gardening, the need to tame erosion is apparent, especially in regions that experience significant precipitation. The U. N. Food and Agriculture Organization (FAO) warns that 90% of the world’s topsoil risks erosion by 2050. It is, therefore, the responsibility of those who participate in land use activities to prevent such a mess.

Jupiter Hadley, Indie Games Plus: Dark Tides is an adventure game where you play as a boardwalk psychic unravelling the demonic mysteries of a seaside Victorian town. Playing as Theodora Green, you are tasked with figuring out...

The post ‘Dark Tides’ Reads Minds to Discover an Old Town’s True Fate appeared first on Indie Games Plus .

Jupiter Hadley, Indie Games Plus: Ghost Haunting is an adventure game about a young girl whose grandmother has been trapped in the realm of the dead. You are a young 8-year old child currently visiting your grandfather....

The post ‘Ghost Haunting’ Follows Unruly Children Attempting Resurrections appeared first on Indie Games Plus .

Stripe Blog, Stripe Blog: Adaptive Pricing is now available for subscriptions, allowing businesses to automatically localize prices in 150+ countries while Stripe handles currency conversion. In an A/B test across 1.5 million subscription checkouts, businesses saw 4.7% higher conversion and 5.4% higher LTV per session, on average.

The Edition

Vol. 1, No. 032 · 2026-03-30

Curated and composed automatically.