

The Civic Digest

VOL. 1, NO. 056

2026-03-30

COVER STORY

Apple intelligence and AI maximalism

No-one outside Apple has really used any Apple Intelligence features yet.

Benedict Evans

It won't launch until the autumn, in stages, and even then it won't work on 80 or 90% of the iPhone installed base, since it needs on-device processing that's only on the iPhone 15 Pro. Nothing is ever as good as the demos (except the original iPhone), Apple might mess up the developer incentives again, and it all might just not be as good as it looks. We'll find out.

But meanwhile, if you step back from the demos and screenshots and look at the strategy, Apple is pointing to most of the key questions and points of leverage in generative AI, and proposing a thesis for how this is going to work that looks very different to all the hype and evangelism.

The 'AI Maximalist' view is that general-purpose chatbots, with multimodal and 'multi-agentic' capabilities, will be able to take over broad classes of complex multi-stage tasks and problems that today need lots of different specially-written software, and will also be able to automate whole new class of complex tasks that could never be done in software before, again in one single general-purpose interface. The chatbot might replace all software with a prompt - 'software is dead'. I'm skeptical about this, as I've written here , but Apple is proposing the opposite: that generative AI is a technology, not a product.

Unbundled into individual features (ditto), which are inherently easier to run as small power-efficient models on small power-efficient devices on the edge (paid for by users, not your capex budget) - which is just as well, because...

This stuff will never work for the mass-market if we have marginal cost every time the user presses 'OK' and we need a fleet of new nuclear power-stations to run it all.

To begin, then: Apple has built an LLM with no chatbot. Apple has built its own foundation models , which (on the benchmarks it published) are comparable to anything else on the market, but there's nowhere that you can plug a raw prompt directly into the model and get a raw

output back - there are always sets of buttons and options shaping what you ask, and that's presented to the user in different ways for different features. In most of these features, there's no visible bot at all. You don't ask a question and get a response: instead, your emails are prioritised, or you press 'summarise' and a summary appears. You can type a request into Siri (and Siri itself is only one of the many features using Apple's models), but even then you don't get raw model output back: you get GUI. The LLM is abstracted away as an API call.

[View fullsize](#)

[View fullsize](#)

[View fullsize](#)

This means that Apple's foundation model won't suggest putting glue on your pizza, as Gemini famously did , simply because you can't get it to answer those kinds of open-ended questions at all. Apple is treating this as a technology to enable new classes of features and capabilities, where there is design and product management shaping what the technology does and what the user sees, not as an oracle that you ask for things.

Instead, the 'oracle' is just one feature, and Apple is drawing a split between a 'context model' and a 'world model'. Apple's models have access to all the context that your phone has about you, powering those features, and this is all private, both on device and in Apple's 'Private Cloud'. But if you ask for ideas for what to make with a photo of your grocery shopping, then this is no longer about

your context, and Apple will offer to send that to a third-party world model - today, ChatGPT. A world model does have an open-ended prompt and does give you raw output, and it might tell you to put glue on your pizza, but that's clearly separated into a different experience where you should have different expectations, and it's also, of course, OpenAI's brand risk, not Apple's. Meanwhile, that world model gets none of your context, only your one-off prompt.

class="">We have yet to see how well Apple's context model really works, but in principle it does look pretty defensible. Neither OpenAI nor any of the other cloud models from new companies (Anthropic, Mistral etc) have your emails, messages, locations, photos, files and so on. Google does have both a world model, and access to your context if you use Android, but that's a distinct minority in the USA (while even less of the Android base than the iPhone base would be able to do any of this locally). Microsoft's AI PCs have some of this context, especially for a work context, but the smartphone is the primary device with all the real context for most people now, not the PC. Does Meta have that context? Some of it, maybe. There will be an interesting anti-trust conversation here, at some point. But the point of leverage is that you need to have your own billion-scale platform already before you can build this: you can't get to it from zero, with a website.

class="">Last May a leaked Google memo claimed there are no moats in LLMs, because everyone has essentially the same access to training data and there would be good open source models. That's pretty much what has happened: the only moat is capital, and access to Nvidia chips (for now), and depending on how you count it there are anything from half-a-dozen to a dozen top-tier models, with OpenAI ahead, but not by enough. Apple doesn't claim its new foundation model is the best for everything, but it does appear to be good enough for the features it wants to provide. This isn't going to play out like search, or operating systems - there is no sign, yet, of clear winner-takes-all effects. Apple could build its own foundation models - it's just money.

class="">Hence, OpenAI is being given (apparently) 'free' distribution to a few hundred million Apple users, for which it bears all the inference costs, in exchange for a chance to upsell to a premium subscription (though looking at all of Apple's WWDC demos, it's not clear how it would do that). But it's also being treated as an interchangeable plug-in. There's a very obvious parallel here with Google paying Apple \$20bn a year to be the default search engine. Apple's AI chief John Giannandrea made this comparison explicitly after the event this week - "I think of it a bit like the way Safari deals with search engines" - and Craig Federighi said that he thought there might be different 'world models' used for different questions. By implication, Apple might send a flights question to one world model and a cooking query to another.

class="">But is web search the right comparison, or should we look at Maps? Apple decided that it makes no sense to try to build a search engine as good as Google, and no-one else

has really managed it either. On the other hand, Apple did build Maps, and though it messed up badly at the beginning, Apple Maps are now at least 'good enough,' because again, there was no real moat except capital. It's already clear that OpenAI is not the new Google: there will not be one winner.

class="">And, surely, the foundation model that Apple built itself to run in its private cloud IS a 'world model', and you could ask it for pizza recipes- it's just that so far, Apple has decided not to offer that UI. Apple is letting OpenAI take the brand risk of creating pizza glue recipes, and making error rates and abuse someone else's problem, while Apple watches from a safe distance. The next step, probably, is to take bids from Bing and Google for the default slot, but meanwhile, more and more use-cases will be quietly shifted from the third party to Apple's own models. It's Apple's own software that decides where the queries go, after all, and which ones need the third party at all.

class="">Of course, none of this is new - it was obvious when Llama 3 came out (if not far earlier) that LLMs would be commodities sold at marginal cost, and the question would be what product you built on top - hence OpenAI hired Kevin Weil as head of product. But Apple is also arguing that a whole category of LLM product will be built in places that the cloud LLMs can't go, or in places where they're just an API call.

class="">This takes me to a broader point. There's an old line that everyone in tech is trying to commoditise someone else, give their product away for free, or both. Meta is giving Llama away for free (both the model and weights and, for now, even embedding free queries in its apps) where the hyper-scalers want to charge for models, because Meta wants this to be cheap commodity infrastructure, and it will differentiate with services and features on top. Apple is doing something very similar. A lot of the compute to run Apple Intelligence is in end-user devices paid for by the users, not Apple's capex budget, and Apple Intelligence is free. (We have no indication on how much the Apple Private Cloud will cost, nor the likely mix of local versus cloud queries.) Nvidia sold \$25bn of AI chips last quarter, and the hyper-scalers will probably spend \$150bn or so on data centres this year, but the global smartphone market is over \$400bn, and the PC market over \$200bn, and your users pay for that. These numbers are not directly comparable (obviously!), but it's a relevant comparison. No-one can know for sure what this will look like in a few years - the models will get bigger but also more efficient, and the edge will get faster - but there are some very powerful incentives to get as much as possible onto the device.

class="">Commoditisation is often also integration. There was a time when 'spell check' was a separate product that you had to buy, for hundreds of dollars, and there were dozens of competing products on the market, but over time it was integrated first into the word processor and then the OS. The same thing happened with the last wave of machine learning - style transfer or image recognition were products for five minutes and then became features. Today

‘summarise this document’ is AI, and you need a cloud LLM that costs \$20/month, but tomorrow the OS will do that for free. ‘AI is whatever doesn’t work yet.’

I’m also curious what this means for Nvidia. As I often and carefully note, I am not a chips analyst, and actual chips and data centre analysts talk a lot about Nvidia’s moats, both in the silicon itself and in the software for developers that it’s built on top. Apple is big enough to take its own path, just as it did moving the Mac to its own silicon: it controls the software and APIs on top of the silicon that are the basis of those developer network effects, and it has a world class chip team and privileged access to TSMC. Who else compares to that today? Google TPUs? It seems unlikely that many other tech companies, even huge ones, will build their own completely custom silicon-to-GUI AI stack. The place for change is in where the models get run: the full foundation model doesn’t fit on a phone, yet, but the more that the real use cases come from unbuilding that model and the ‘oracle’ into features, the more that the inference could shift to the edge a lot faster. Again, Apple is pointing to the question.

Finally, all of these are still just theses. None of this worked just 18 months ago, and we don’t have product-market fit yet. Anything up to half of the population of many developed world countries has already tried generative AI, but half of those never tried it twice. That’s great awareness for something so new, but it’s not traction. We don’t know what the product will be, nor the market, nor the science, and everything is still changing really fast. There may be new breakthroughs in agents or error rates that change the use cases entirely.

Meanwhile, incumbents always try to make the new thing a feature. Google and Microsoft have spent the last 18 months spraying LLMs all over their products, and so has every enterprise SaaS company: my old colleague Steven Sinofsky says that ‘every text box on the internet is going to get an LLM’. Apple is doing something slightly different - it’s proposing a single context model for everything you do on your phone, and powering features from that, rather than adding disconnected LLM-powered features at disconnected points across the company. But it’s still trying to make the ‘disruptive’ new tech into a feature, and it’s trying to put ChatGPT into a box, limited to a pretty narrow set of use cases, and interchangeable with Anthropic or Gemini.

As I said at the beginning, that’s really the single core question about the future of Generative AI - is this a new general purpose tool, where one product from one company does the work of hundreds of pieces of software from hundreds of companies, or is this a generic technology that will enable features inside products from hundreds or thousands of companies?

{

FRONT PAGE

Here We Go Again

Penny Arcade

Ronia is, in fact, on stage jumping around going fucking crazy in Mamma Mia, which is essentially like a jazzercise class with a lyric book. Like, on some Tae Bo shit. I would last part of one song before my lungs tried to crawl out of my mouth. Gabriel was under some misconceptions regarding the production that I was happy to correct by striking him repeatedly with one of those little ice fishing bats. Did you know that those little clubs are called “priests”? Gabriel didn’t know either. I made sure he understood as it was coming down .

}

{

FEATURES

Bowflex SelectTech Adjustable Dumbbells Are Over 50% Off in the Amazon Spring Sale

Eric Song · IGN

Why pay \$479 for a brand new pair of Bowflex adjustable dumbbells? As part of Amazon's Mega Spring Sale, Amazon Resale is offering a pair of used "like new" Bowflex Results Series SelectTech 552 Adjustable Dumbbells for just \$227.99 with free shipping after an extra 40% off discount is applied during checkout. I haven't seen Bowflex dumbbells sell for this low of a price for at least a few years. You better hurry because Amazon Resale inventory is usually limited.

Amazon Resale (formerly Amazon Warehouse) is Amazon's official account for used goods. The products are sorted by condition, with "Like New" being the best. There's a good chance these items haven't been used at all and were returns from customers who have changed their minds, or simply products with blemished packaging. All Amazon Resale items also get the same 30-day return policy as any other Amazon purchase.

Get Bowflex SelectTech Adjustable Dumbbells (Pair) for \$228

Used "Like New" Condition

Update: If the link lands on a "Very Good" condition item, click on the "Other sellers on Amazon" link to see more buying options, including the "Like New" condition item.

Each Bowflex dumbbell is adjustable from 5 pounds all the way up to 52.5 pounds with 15 total increments, so it's essentially like 15 total weights in one package. The SelectTech dumbbells have been around for a long time. Many people, including myself, have used them for well over a decade with no problems. The structural components and moving parts are made of metal wrapped in plastic and rubber to protect against dings, deaden vibrations, and provide better grip. The "Results Series" is just an aesthetic redesign of the original Bowflex SelectTech 552 dumbbells that happened after Johnson Health Tech took over all of Bowflex's assets.

Free weights are a great choice for anyone who wants to work out at home but is limited on space. A pair of dumbbells is all the equipment you need to get access to a wide range of exercises that can target just about every muscle in your upper and lower body. Supplement your free weight exercises with cardio and other workouts that use your own body weight as resistance (like pushups, burpees, or squats) and you may never need to purchase more equipment.

If these are still out of your budget, there are some other options and more fitness deals to check out this week.

Shop More Fitness Deals from Amazon's Spring Sale

How to Follow IGN Deals Recommendations

The IGN Deals team has over 30 years of combined experience finding the best discounts and preorders available

online. If you want the latest updates from our trusted team, here's how to follow our coverage:

Set IGN as a preferred source in Google

IGN Finds on X

IGN Deals on Instagram

IGN Deals on Facebook

IGN Deals on Tiktok

Eric Song is the IGN commerce manager in charge of finding the best gaming and tech deals every day. When Eric isn't hunting for deals for other people at work, he's hunting for deals for himself during his free time.

The SelectTech dumbbells have been around for a long time.

Eric Song

}

{

Crimson Desert fans praise the devs for quickly and directly addressing player complaints and requests: ‘This is what game studios should aspire to be’

Andrej Barovic · Destructoid

As Crimson Desert continues to see a ton of success 10 days after launch, players are now praising developer Pearl Abyss for how quickly it's been addressing issues since the moment the game released. With this latest patch, a lot of concerns were directly resolved, improved, or patched out, while some quirks were even embraced as official mechanics.

Soaring through the Steam charts and peaking at over 270,000 players, Crimson Desert is definitely faring a lot better than it did on launch day. It even managed to rebound its Steam reviews, now being labeled as “very positive”—a significant upgrade over its initial “mixed” reception. And all of this has to do with Pearl Abyss’ patches, which have improved the game to such a degree that I find myself feeling as if I’m playing a different game, especially coming off the review build.

This latest patch added five entirely new mounts to the game, made sprinting a lot less of a hassle by reducing how much button mashing you need to do (a major complaint since release), made cooking more intuitive, made fast travel much faster, and even added certain currencies that make the upgrade process a lot cheaper and easier early on. This is just scratching the surface, though, and fans have been reacting more than positively since it came out.

Visual representation of players hugging the dev team.
Image via Pearl Abyss

In one thread, a player said that “This is what game studios should aspire to be,” praising Pearl Abyss for addressing the sprinting mechanic mentioned above. “Shout out to the devs for listening to people,” said another.

Some showed that you can never please everyone despite your best efforts, as those replying to the thread said, in no small number, that they preferred how things were before the patch. I’m certain this will eventually lead to Pearl Abyss adding in optional control schemes once everything is fleshed out, allowing everyone to tailor their own fun to themselves, much as the game lets you do things at your own pace and leisure.

In another thread, a player said about Pearl Abyss that “they listen, and they deliver,” as well as that it’s the “best dev team ever.”

“They really do. Love the changes,” said another.

It’s certainly amazing to see a developer team actually invest the time and effort necessary to make its game as good as it can possibly be. It’s gotten to the point where so many player complaints have been fixed or just made better that the game now has an 80 percent positive rating on Steam

}

and is slowly but surely climbing in player numbers despite being a singleplayer game that came out 10 days ago.

I love to see it, even if I myself don’t agree with all the changes that have been made. Other developers need to take note.

The post Crimson Desert fans praise the devs for quickly and directly addressing player complaints and requests: ‘This is what game studios should aspire to be’ appeared first on Destructoid .

{

Top 7 Best Wordpress Plugin Of All Time

Unknown · TechCrunch

If you are looking for the best wordpress plugins, then you are at the right place. Here is the list of best wordpress plugins that you should use in your blog to boost SEO, strong your security and know every aspects of your blog . Although creating a good content is one factor but there are many wordpress plugins that perform different actions and add on to your success. So let's start

1. Yoast SEO

Those users who are serious about SEO, Yoast SEO will do the work for them to reach their goals. All they need to do is select a keyword, and the plugin will then optimize your page according to the specified keyword

Yoast offers many popular SEO WordPress plugin functions. It gives you real-time page analysis to optimize your content, images, meta descriptions, titles, and keywords. Yoast also checks the length of your sentences and paragraphs, whether you're using enough transition words or subheadings, how often you use passive voice, and so on. Yoast tells Google whether or not to index a page or a set of pages too.

Let me summarize these points in bullets:

Enhance the readability of your article to reduce bounce rate

Optimize your articles with targetted keywords

Let Google know who you are and what your site is about

Improve your on-page SEO with advanced, real-time guidance and advice on keyword usage, linking, and external linking.

Keep your focus keywords consistent to help rank better on Google.

Preview how your page would appear in the search engine results page (SERP)

Crawl your site daily to ensure Google indexes it as quickly as possible.

Rate your article informing you of any mistakes you might have made so that you can fix them before publishing.

Stay up-to-date with Google's latest algorithm changes and adapt your on-page SEO as needed with smart suggestions from the Yoast SEO plugin. This plugin is always up-to-date.

Free Version is available

Pricing

Premium version=\$89/year that comes with extra functions, allowing you to optimize your content up to five keywords, among other benefits.

2. WP Rocket

A website running WordPress can put a lot of strain on a server, which increases the chances that the website will crash and harm your business. To avoid such an unfortunate situation and ensure that all your pages load quickly, you need a caching plugin like WP Rocket.

WP Rocket plugin designed to increase your website speed. Instead of waiting for pages to be saved to cache, WP Rocket turns on desired caching settings, like page cache and gzip compression. The plugin also activates other features, such as CDN support and lazy image loading, to enhance your site speed.

Features in bullets:

Preloading the cache of pages

Reducing the number of HTTP requests allows websites to load more quickly.

Decreasing bandwidth usage with GZIP compression

Apply optimal browser caching headers (expires)

Minifying and combining JavaScript and CSS files

Remove Unused CSS

Deferred loading of images (LazyLoad)

WebP compatibility

Deferred loading of JavaScript files

Delay JavaScript Execution

Critical Path CSS generation and deferred loading of CSS files

Database optimization

WordPress Heartbeat API control

DNS prefetch

CDN integration

Cloudflare integration

Sucuri integration

Easy import/export of settings

Easy roll back to a previous version

Pricing

}

standard-article(title: [Samsung’s Premium 65” S90 4K OLED Smart TV Drops to Just \$1,099 for the Amazon Spring Sale], author: [Eric Song], source-name: [IGN], images: (), paragraphs: ([As part of its mega Spring Sale , Amazon recently dropped the price on one of Samsung’s best OLED TVs. The highly rated 65” Samsung S90D 4K OLED Smart TV has just dropped in price to \$1,099 with free delivery. This is a new TV with a 1 year Samsung warranty. We recently named the Samsung S90 TV as our top pick for best gaming TV . The Amazon marketplace seller “Walt’s TV” is an authorized Samsung reseller.], [65” Samsung S90D 4K OLED TV for \$1,099], [The Samsung S90D is the prior generation 2024 model. A new “S90F” model is available for about \$250 more, but the differences are minor. The S90D is equipped with Samsung’s proprietary quantum dot (QD) OLED panel. QD OLED panels are brighter than traditional OLED panels without losing the color accuracy, range, and wide-viewing angles that OLEDs are known for. Compared to a traditional LED LCD TV, an OLED TV offers superior image quality, near-infinite blacks, near-infinite contrast ratio, and near-instantaneous response times.], [OLED technology aside, the Samsung S90D also has all the features you’d want in a gaming TV . It has a native 120Hz panel that can be pushed to as high as 144Hz as well as four total HDMI 2.1 ports for running 4K at 120Hz on a PS5 or Xbox Series X gaming console. The S90D also supports variable refresh rate and auto low latency mode. This is an especially good complement to the recently released Playstation 5 Pro console since it can drive more games in quality mode (native 4K @ 120fps) than the standard PS5, which often needs to be set to performance mode (upscaled 4K).], [Is this a good TV for the Nintendo Switch 2?], [Despite being a current generation console, the Nintendo Switch 2 has lax TV requirements compared to the PS5 or Xbox Series X. When the Switch is connected to its dock and a television, it is only capable of outputting a 4K resolution signal at up to 60 frames per second. It can go up to 120fps if the resolution drops to 1080p, but the data rate required for both these options is pretty much the same. Oddly, it doesn’t even get VRR when docked – that’s reserved for handheld mode only. The Switch 2 console, however, will still benefit greatly from the Samsung’s gorgeous OLED panel.], [How to Follow IGN Deals Recommendations], [The IGN Deals team has over 30 years of combined experience finding the best discounts and preorders available online. If you want the latest updates from our trusted team, here’s how to follow our coverage:], [Set IGN as a preferred source in Google], [IGN Deals on X], [IGN Deals on Instagram], [IGN Deals on Facebook], [IGN Deals on Tiktok], [Eric Song is the IGN commerce manager in charge of finding the best gaming and tech deals every day. When Eric isn’t hunting for deals for other people at work, he’s hunting for deals for himself during his free time.],), insert-map: (:), word-count: 499, edited-for-length: false, debug-mode: false,)

{

Fire In The Hole

Penny Arcade

standard-article(title: [Weekly Update 483], author: [Troy Hunt], source-name: [Troy Hunt], images: (), paragraphs: ([Presently sponsored by: Report URI: Guarding you from rogue JavaScript! Don’t get pwned; get real-time alerts & prevent breaches #SecureYourSite], [Building out an IoT environment is a little like the old Maslow’s Hierarchy of Needs. All the stuff on the top is only any good if all the stuff on the bottom is good, starting with power. This week, I couldn’t even get that right, but thankfully, sparky to rescue and ensuite underfloor heating disconnected, and we now have reliable power again. On top of that is the layer that has increasingly been my nemesis - the network. Two days after recording, I’ve just spent the better part of the entire day making a much more concerted effort to adjust channel and power settings on APs, lock clients that don’t move to the APs that make the most sense, and generally just screw around with it until stuff worked. And then I turned off a circuit, turned it back on again, and all hell broke loose ☹️], [style=“width: 170px; display: inline-block; margin-right: 3px;”>], [style=“width: 175px; display: inline-block; margin-right: 3px;”>], [style=“width: 118px; display: inline-block; margin-right: 3px;”>], [style=“width: 120px; display: inline-block;”>], [id=“references”>References], [Sponsored by: 1Password Extended Access Management: Secure every sign-in for every app on every device.],), insert-map: (:), word-count: 205, edited-for-length: false, debug-mode: false,)

I've been trying to have a beer with Kim Swift for about five years now, and it finally happened. She suggested The Bine up in Bothell, and you should draft on that suggestion. I'm afraid I made a pun there without meaning to. But it's also real close to Zulu's , a place you might want to be anyway. Anywhoozle, the last half of the meetup was just her breaking down how Pokopia is a masterclass of game design. A little while later, a sheepish server came over and said, "I don't want to interrupt... but were you talking about Pokopia?"

She was.

}

{

ANALYSIS

Remaking the app store

Benedict Evans

There's very little we can say about Apple's App Store policies that we didn't say when it launched in 2008, in 2011 when it first tightened the rules, or in 2020, when Epic tried to stunt its way out of them. Meanwhile, the world has moved on: generative AI is the centre of tech innovation, excitement and company creation now, not smartphone apps. The Vision Pro might swing the pendulum back to AR/VR and Apple, but not for years. Seen from Silicon Valley, this is an argument from a decade ago.

That time-lag itself is a broader structural problem: the decision-cycle of regulation and due process tends to be longer than the life-cycle of technology platform. Everyone in the regulation industry is very conscious of that, but the alternative is to speculate about the future development of markets that don't exist yet, and that has its own problems, as the FTC found when it tried to stop Meta from buying Within. Pick your trade-off.

And yet, this platform is the primary computing device for over a billion people, and with the DMA we have a big new law from the EU that tells Apple it has to change things. There's also a court ruling in the USA (rather narrower and focused on Apple's anti-steering rule), and perhaps more coming from a rumoured DoJ/FTC lawsuit. Apple has responded with a very broad set of proposals and changes to how things work, or might work. These are not at all what people who wanted change hoped to see. However it's much less clear whether they are what the EU wanted, or indeed whether it's actually possible to build what the EU wanted... presuming that anyone involved in the DMA actually knows what that is.

The challenge for the EU is that the iPhone proposed a new kind of tradeoff for the software industry. IBM's mainframes gave giant companies a computer and software as one sealed system that IBM controlled. In contrast, the PC was an open system that no-one really controlled (not even Microsoft), where anyone could do anything with their computer, and so could any developer. This was freedom, but it also meant that it was easy for you to break your computer and easy for a developer to break it (accidentally or on purpose), and easy for them to steal your data. When there were millions of PCs, mostly owned by pretty technical people, and they were not connected to an open global network, that was one thing, but once billions of people connected to the internet, that was a big problem.

Hence, the iPhone software model proposed a new level of abstraction and a new trade-off. Random third-party developers cannot do whatever they want with your data or break your computer, and indeed you cannot break your computer. The iPhone puts security, privacy and reliability over the freedom to change anything and do anything. This has been a huge success: iOS has the largest mobile app ecosystem, and indeed the largest software ecosystem there

has ever been, with millions of apps and billions of downloads, which makes it rather incomplete to call it 'closed'.

This is the right trade-off for almost everyone. A small number of highly-technical people would like to be able to change anything on their computer even at the cost of opening security holes everywhere: they are free to buy a different device.

Generally, this is a spectrum. Some restrictions are entirely about privacy, security and performance (no, you can't have root access to the file system), but others are also about making life harder for competitors, or extracting rent, or both. Why can't WhatsApp be your default SMS client, say, and why can't Square make a wallet app that uses NFC? There are serious security and ease-of-use answers to this, but also, serious business answers. Banning streaming games services has no security reasoning: it's purely a desire to protect Apple's own ecosystem.

The payment rules are a subset of this, and they have some of the clearest examples. Apple believes that developers should pay it for using the platform (a well-established model on games consoles, but not on PCs or Macs), and it chose to charge them by letting most developers use it for free (Uber and TikTok don't pay anything) but charging a 30% commission for digital content bought inside apps. That sounded simple in 2008, but quickly led to an inverted pyramid of increasingly complex problems, exemptions, carve-outs and edge-cases.

But while you can at least argue whether newspapers should be paying that 30%, Apple again uses this to hurt or block competitors in ways that have no user benefit. The obvious problem right from the beginning was music and ebooks. Neither of these have the gross margin available to pay Apple a 30% commission (meanwhile, Apple's store model has a cap of 10,000 products per app, so you couldn't use it for an ebook store even if you were willing to pay), so they have to use their own subscription, but Apple banned them from telling the user how to do that.

In case there was any doubt, Steve Jobs was very clear about this internally in 2011, in the email below: the new rules introduced in 2011 would preclude some competitors. These rules would be good for users in some ways (see Eddy Cue's comment further down about the app experience) but bad in others.

On one hand, the principle of Apple's sandboxed, managed, curated system (across both the OS sandbox and the store) is great for almost all developers, and users, and Apple. (Apple may or may not see things in that order.) Something over a billion people have an iPhone, and though we sometimes call it 'closed', it has millions of apps and billions of downloads: iOS is the largest open software ecosystem in history. The App Store model is not some kind of aberration incidental to the success of the iPhone: it is a core part of how Apple delivers the promise: a phone that just works and apps that just work.

}

standard-article(title: [Folk Etymologies], author: [Penny Arcade], source-name: [Penny Arcade], images: (), paragraphs: ([As something of a connected phenomenon - my instinct is to say “corollary,” but I think that’s not correct - as Penny Arcade was unearthed by the wily Zillenial, the many fruits of our enterprise became known. The true facts of the name PAX, for example - a name I came up with in five seconds at Costco because we got a call saying the show thing we were trying to do needed a name. I’m always delighted when someone finds the information I tucked in there. Never you mind that the next question is usually “what’s a Penny Arcade” because that’s not what this story is about!!!],), insert-map: (:), word-count: 108, edited-for-length: false, debug-mode: false,)

standard-article(title: [All gods and bonuses in Nova Roma], author: [Arka Sarkar], source-name: [Destructoid], images: (), paragraphs: ([Religion plays a key role in Nova Roma in deciding if your civilization flies to a new high or meets doom. Dedicating temples to the heavenly powers will help your civilization to grow. Failure to do so can result in its perishing.], [Let’s find out all the gods that are present in Nova Roma and the boons each of them offers.], [Table of contents], [All god boons in Nova Roma], [How to gain Favor from the gods in Nova Roma], [The table below has a list of all the gods and the boons they offer.], [God Boons Jupiter +3 happiness in a 5 tile radius. Ceres +1 output to the nearest three farms, vineyards, and orchards in its radius. Neptune +1 fish yields to the nearest five fishing buildings in its radius. Vulcan +1 output to the nearest five fire-based industry Ares +25% damage to the nearest three defense towers in its radius.], [Image via Hooded Horse], [The first step is to build a temple and then dedicate it to a god of your choice. Your decision will be driven by the potential buffs that you want to apply. Once you dedicate a temple to a god, you’ll need to complete Divine Tasks. Once you complete them, you’ll be rewarded with Favor.], [You can then offer Favor to a god in the Small Temple to unlock the benefits (listed above). If you have a Large Temple, you can also throw a festival to gain Favor. Favor is used to unlock new technology spread across several branches. The unlocks will allow you to improve your civilization through time and bring it close to prosperity.], [The post All gods and bonuses in Nova Roma appeared first on Destructoid .],), insert-map: (:), word-count: 307, edited-for-length: false, debug-mode: false,)

Your decision will be driven by the potential buffs that you want to apply.

Arka Sarkar

{

BRIEFS

The Hacker News, The Hacker News (cybersecurity): Most teams have security tools in place. Alerts are firing, dashboards look clean, threat intel is flowing in. On the surface, everything feels under control. But one question usually stays unanswered: Would your defenses actually stop a real attack? That's where things get shaky. A control exists, so it's assumed to work. A detection rule is active, so it's expected to catch something. But very

Daniel Curtis, InfoQ: Nuxt Test Utils has released version 4.0.0, which primarily integrates Vitest v4. This update changes the test environment setup to beforeAll, resolving issues with module-level mocks. It also improves mockNuxtImport for cleaner partial mocking and enhances state management for registered endpoints. The library remains vital for testing in the Nuxt framework, bridging unit and end-to-end testing.

By Daniel Curtis

Connor Makar, Eurogamer: Bungie has implemented a substantial matchmaking change which will aid those playing from isolated regions - think Australia or South America - find games faster.

[Read more](#)

Zack Garvey, PlayStation Blog: class="">> PlayStation Store's Spring Sale begins tomorrow, Wednesday, March 25. For a limited time* you can enjoy deep discounts of up to 75% on a vast range of games that include iconic blockbusters, acclaimed classics, and more.

class="">

class="">

*PlayStation Store Spring Sale promotion runs from Wednesday, March 25 at 00:00 AM PDT/BST/JST and finishes Wednesday April 22 at 11:59 PM PDT/BST/JST**. Each title may have different sale periods. Please refer to the information in the title details page.

class="has-small-font-size"> **The promotion's start and finish time may differ in certain regions. Game selection may differ by region.

class="">

Cerner Engineering, Cerner Engineering: The FHIR standard started as a better way to exchange healthcare data, but it also provides a solid basis for deep analytics and Machine Learning at scale. This post looks at an example from the recent FHIR DevDays conference that does just that. You can also run the interactive FHIR data engineering tutorial used in the conference yourself. Our first step is to bring FHIR data into a data lake – a computational environment where our analysis can easily and efficiently work through petabytes of data.

Victoria Phillips Kennedy, Eurogamer: “I remember being in a really strange time in my life, oddly enough, and a big transition where I felt really uncertain of myself and my place in the world, and this audition came around that was somehow celebrating all the things about myself that I thought were problematic.”

[Read more](#)

The Hacker News, The Hacker News (cybersecurity): Cybersecurity has changed fast. Roles are more specialized, and tooling is more advanced. On paper, this should make organizations more secure. But in practice, many teams struggle with the same basic problems they faced years ago: unclear risk priorities, misaligned tooling decisions, and difficulty explaining security issues in terms the business understands. These challenges do not

Cerner Engineering, Cerner Engineering: At Cerner, we use the latest technologies to transform healthcare. We are the engineers and designers that focus on making you, your family, and friends healthy and happy. We love solving the hard problems that no one else wants to tackle. We embrace polyglot programming, distributed systems, and open source. We also know that it's important to allow engineers to take a step back to learn and try new things. That's why we host periodic hackathons, meetups, tech talks, and an annual internal two-day developers conference.

Bruno Couriol, InfoQ: Max Inden recently explored in a talk at FOSDEM 2026 how the upcoming WebTransport protocol and Web API enhance WebSocket capabilities. WebTransport seeks to provide, among other things, lower latency and transparent network switching for key use cases such as high-frequency financial data streaming, cloud gaming, live streaming, and collaborative editing.

By Bruno Couriol

Robert Purchase, Eurogamer: Hello and welcome back to our regular feature where we write a little about the games we've been playing. This week, Bertie discovers he's not as good at Slay the Spire 2 as he would like to be; Connor continues to fall to the Dark Side and embrace his inner troll; and Marie realises that older games really were more challenging.

[Read more](#)

Cerner Engineering, Cerner Engineering: In the dark ages of development, great software meant packing in the functionality. People began doing more and more with their software. Updates meant newer and more exciting functionality. Sounds great, right? Of course it does, but something went horribly wrong. Slowly we became inundated with cluttered screens as software developers struggled to find a place to put their latest innovative functionality. Buttons began adding up and before we knew it, we were inventing user interface controls like ribbons to hold all the buttons.

The Hacker News, The Hacker News (cybersecurity): Google on Thursday announced a new “advanced flow” for Android sideloading that requires a mandatory 24-hour wait period to install apps from unverified developers in an attempt to balance openness with safety. The new changes come against the backdrop of a developer verification mandate the tech giant announced last year that requires all Android apps to be signed by Google Play Protect before they can be installed on devices. The move is part of a broader effort to improve the security of the Android ecosystem, which has long been a target for malware and other threats. The new “advanced flow” will require users to wait 24 hours before they can install an app from an unverified developer. This is intended to give Google Play Protect time to scan the app for malware. The move is seen as a compromise between the need for openness and the need for security. Some critics argue that the 24-hour wait period is too long and could be a burden for developers. However, Google insists that the move is necessary to protect users from malware. The new changes will be rolled out to all Android devices in the coming weeks.

}

The Civic Digest

Vol. 1, No. 056 · 2026-03-30

Curated and composed automatically.