

The Daily Star

VOL. 1, NO. 060

2026-03-30

FEATURES

Fresh Energy In March (2026 Wallpapers Edition)

Cosima Mielke · Smashing Magazine

Blooming flowers, longer days, milder temperatures — with March just around the corner, the world is slowly but surely awakening from its winter slumber, fueling us with fresh energy. And even if spring is far away in your part of the world, you might sense that 2026 has gained full speed by now, making it the perfect moment to turn those plans and ideas you've been carrying around into action.

To accompany you on all those adventures that March may bring, we have a new collection of desktop wallpapers for you, just as it has been a monthly tradition here at Smashing Magazine for more than 14 years already. Designed by artists and designers from across the globe, each wallpaper comes in a variety of screen resolutions and can be downloaded for free. A huge thank-you to everyone who shared their designs with us — this post wouldn't be possible without your kind support!

If you, too, would like to get featured in one of our upcoming wallpapers editions, please don't hesitate to submit your design. We can't wait to see what you'll come up with! Happy March!

You can click on every image to see a larger preview.

We respect and carefully consider the ideas and motivation behind each and every artist's work. This is why we give all artists the full freedom to explore their creativity and express emotions and experience through their works. This is also why the themes of the wallpapers weren't anyhow influenced by us but rather designed from scratch by the artists themselves.

"With Spring knocking and other seasons fighting to get attention, March greets us with blossoms." — Designed by Ginger It Solutions from Serbia.

preview

with calendar: 320x480, 640x480, 800x480, 800x600, 1024x768, 1024x1024, 1152x864, 1280x720, 1280x800, 1280x960, 1280x1020, 1400x1050, 1440x900, 1600x1200, 1680x1050, 1680x1200, 1920x1080, 1920x1200, 1920x1440, 2560x1440

without calendar: 320x480, 640x480, 800x480, 800x600, 1024x768, 1024x1024, 1152x864, 1280x720, 1280x800, 1280x960, 1280x1020, 1400x1050, 1440x900, 1600x1200, 1680x1050, 1680x1200, 1920x1080, 1920x1200, 1920x1440, 2560x1440

preview

with calendar: 640x480, 800x480, 800x600, 1024x768, 1024x1024, 1152x864, 1280x720, 1280x800, 1280x960, 1280x1024, 1366x768, 1400x1050, 1440x900, 1600x1200, 1680x1050, 1680x1200, 1920x1080, 1920x1200, 1920x1440, 2560x1440, 3840x2160

without calendar: 640x480, 800x480, 800x600, 1024x768, 1024x1024, 1152x864, 1280x720, 1280x800, 1280x960, 1280x1024, 1366x768, 1400x1050, 1440x900, 1600x1200, 1680x1050, 1680x1200, 1920x1080, 1920x1200, 1920x1440, 2560x1440, 3840x2160

I'm Not Okay, But It's Okay

preview

with calendar: 640x480, 800x480, 800x600, 1024x768, 1024x1024, 1152x864, 1280x720, 1280x800, 1280x960, 1280x1024, 1366x768, 1400x1050, 1440x900, 1600x1200, 1680x1050, 1680x1200, 1920x1080, 1920x1200, 1920x1440, 2560x1440, 3840x2160

without calendar: 640x480, 800x480, 800x600, 1024x768, 1024x1024, 1152x864, 1280x720, 1280x800, 1280x960, 1280x1024, 1366x768, 1400x1050, 1440x900, 1600x1200, 1680x1050, 1680x1200, 1920x1080, 1920x1200, 1920x1440, 2560x1440, 3840x2160

Let's Spring

"After some freezing months, it's time to enjoy the sun and flowers. It's party time, colors are coming, so let's spring!" — Designed by Colorsfera from Spain.

preview

without calendar: 320x480, 640x480, 800x480, 800x600, 1024x768, 1024x1024, 1152x864, 1280x720, 1280x800, 1280x960, 1280x1020, 1400x1050, 1440x900, 1600x1200, 1680x1050, 1680x1200, 1920x1080, 1920x1200, 1920x1440, 2560x1440

"This March, our calendar design epitomizes the heralds of spring. Soon enough, you'll be waking up to the singing of swallows, in a room full of sunshine, filled with the empowering smell of daffodil, the first springtime flowers. Spring is the time of rebirth and new beginnings, creativity and inspiration, self-awareness, and inner reflection. Have a budding, thriving spring!" — Designed by PopArt Studio from Serbia.

preview

without calendar: 320x480, 640x480, 800x480, 800x600, 1024x768, 1024x1024, 1152x864, 1280x720, 1280x800, 1280x960, 1280x1024, 1366x768, 1400x1050, 1440x900, 1600x1200, 1680x1050, 1680x1200, 1920x1080, 1920x1200, 1920x1440, 2560x1440

"This month, I want to go to the woods and explore my new world in sunny weather." — Designed by Zi-Cing Hong from Taiwan.

preview

without calendar: 1024x768, 1152x864, 1280x720, 1280x800, 1280x960, 1366x768, 1400x1050, 1440x900, 1600x1200, 1680x1050, 1920x1080, 1920x1200, 1920x1440, 2560x1440

Smashing Animations Part 7: Recreating Toon Text With CSS And SVG

Andy Clarke · Smashing Magazine

After finishing a project that required me to learn everything I could about CSS and SVG animations, I started writing this series about Smashing Animations and “How Classic Cartoons Inspire Modern CSS.” To round off this year, I want to show you how to use modern CSS to create that element that makes Toon Titles so impactful: their typography.

In the silent era of the 1920s and early '30s, the typography of a film's title card created a mood, set the scene, and reminded an audience of the type of film they'd paid to see.

Cartoon title cards were also branding, mood, and scene-setting, all rolled into one. In the early years, when major studio budgets were bigger, these title cards were often illustrative and painterly.

But when television boomed during the 1950s, budgets dropped, and cards designed by artists like Lawrence “Art” Goble adopted a new visual language, becoming more graphic, stylised, and less intricate.

Note: Lawrence “Art” Goble is one of the often overlooked heroes of mid-century American animation. He primarily worked for Hanna-Barbera during its most influential years of the 1950s and 1960s.

Goble wasn't a character animator. His role was to create atmosphere, so he designed environments for The Flintstones, Huckleberry Hound, Quick Draw McGraw, and Yogi Bear, as well as the opening title cards that set the tone. His title cards, featuring paintings with a logo overlaid, helped define the iconic look of Hanna-Barbera.

Goble's artwork for characters such as Quick Draw McGraw and Yogi Bear was effective on smaller TV screens. Rather than reproducing a still from the cartoon, he focused on presenting a single, strong idea — often in silhouette — that captured its essence. In “The Buzzin' Bear,” Yogi buzzes by in a helicopter. He bounces away, pic-a-nic basket in hand, in “Bear on a Picnic,” and for his “Prize Fight Fright,” Yogi boxes the title text.

With little or no motion to rely on, Goble's single frames had to create a mood, set the scene, and describe a story. They did this using flat colours, graphic shapes, and typography that was frequently integrated into the artwork.

As designers who work on the web, toon titles can teach us plenty about how to convey a brand's personality, make a first impression, and set expectations for someone's experience using a product or website. We can learn from the artists' techniques to create effective banners, landing-page headers, and even good ol' fashioned splash screens.

Cartoon title cards show how merging type with imagery delivers the punch a header or hero needs. With a handful

of text-shadow, text-stroke, and transform tricks, modern CSS lets you tap into that same energy.

The Toon Text Title Generator

Partway through writing this article, I realised it would be useful to have a tool for generating text styled like the cartoon titles I love so much. So I made one.

My Toon Text Title Generator lets you experiment with colours, strokes, and multiple text shadows. You can adjust paint order, apply letter spacing, preview your text in a selection of sample fonts, and then copy the generated CSS straight to your clipboard to use in a project.

Toon Title CSS

You can simply copy-paste the CSS that the Toon Text Title Generator provides you. But let's look closer at what it does.

Text shadow

Look at the type in this title from Augie Doggie's episode “Yuk-Yuk Duck,” with its pale yellow letters and dark, hard, offset shadow that lifts it off the background and creates the illusion of depth.

You probably already know that text-shadow accepts four values: (1) horizontal and (2) vertical offsets, (3) blur, and (4) a colour which can be solid or semi-transparent. Those offset values can be positive or negative, so I can replicate “Yuk-Yuk Duck” using a hard shadow pulled down and to the right:

```
color: #f7f76d; text-shadow: 5px 5px 0 #1e1904;
```

On the other hand, this “Pint Giant” title has a different feel with its negative semi-soft shadow:

```
color: #c2a872; text-shadow: -7px 5px 0 #b100e, 0 -5px 10px #546c6f;
```

To add extra depth and create more interesting effects, I can layer multiple shadows. For “Let's Duck Out,” I combine four shadows: the first a solid shadow with a negative horizontal offset to lift the text off the background, followed by progressively softer shadows to create a blur around it:

```
color: #6f4d80; text-shadow: -5px 5px 0 #260e1e, /* Shadow 1 */ 0 0 15px #e9ce96, /* Shadow 2 */ 0 0 30px #e9ce96, /* Shadow 3 */ 0 0 30px #e9ce96; /* Shadow 4 */
```

These shadows show that using text-shadow isn't just about creating lighting effects, as they can also be decorative and add personality.

Many cartoon title cards feature letters with a bold outline that makes them stand out from the background. I can recreate this effect using text-stroke. For a long time, this property was only available via a -webkit- prefix, but that also means it's now supported across modern browsers.

Unlock the Secrets of a Custom Sign-In Page with Tailwind and JavaScript

Okta Developer Blog

We recommend redirecting users to authenticate via the Okta-hosted sign-in page powered by the Okta Identity Engine (OIE) for your custom-built applications. It's the most secure method for authenticating. You don't have to manage credentials in your code and can take advantage of the strongest authentication factors without requiring any code changes.

The Okta Sign-In Widget (SIW) built into the sign-in page does the heavy lifting of supporting the authentication factors required by your organization. Did I mention policy changes won't need any code changes?

But you may think the sign-in page and the SIW are a little bland. And maybe too Okta for your needs? What if you can have a page like this?

With a bright and colorful responsive design change befitting a modern lifestyle.

Let's add some color, life, and customization to the sign-in page.

In this tutorial, we will customize the sign-in page for a fictional to-do app. We'll make the following changes:

Use Tailwind CSS framework to create a responsive sign-in page layout

Add a footer for custom brand links

Display a terms and conditions modal using Alpine.js that the user must accept before authenticating

Take a moment to read this post on customizing the Sign-In Widget if you aren't familiar with the process, as we will be expanding from customizing the widget to enhancing the entire sign-in page experience.

Stretch Your Imagination and Build a Delightful Sign-In Experience

Customize your Gen3 Okta Sign-In Widget to match your brand. Learn to use design tokens, CSS, and JavaScript for a seamless user experience.

In the post, we covered how to style the Gen3 SIW using design tokens and customize the widget elements using the `afterTransform()` method. You'll want to combine elements of both posts for the most customized experience.

Table of Contents

Customize your Okta-hosted sign-in page

Use Tailwind CSS to build a responsive layout

Use Tailwind for custom HTML elements on your Okta-hosted sign-in page

Add custom interactivity on the Okta-hosted sign-in page using an external library

Customize Okta-hosted sign-in page behavior using Web APIs

Add Tailwind, Web APIs, and JavaScript libraries to customize your Okta-hosted sign-in page

To follow this tutorial, you need:

An Okta account with the Identity Engine, such as the Integrator Free account .

Your own domain name

A basic understanding of HTML, CSS, and JavaScript

A brand design in mind. Feel free to tap into your creativity!

An understanding of customizing the sign-in page by following the previous blog post

Let's get started!

Before we begin, you must configure your Okta org to use your custom domain. Custom domains enable code customizations, allowing us to style more than just the default logo, background, favicon, and two colors. Sign in as an admin and open the Okta Admin Console, navigate to Customizations > Brands and select Create Brand + .

Follow the Customize domain and email developer docs to set up your custom domain on the new brand.

Customize your Okta-hosted sign-in page

We'll first apply the base configuration using the built-in configuration options in the UI. Add your favorite primary and secondary colors, then upload your favorite logo, favicon, and background image for the page. Select Save when done. Everyone has a favorite favicon, right?

I'll use #ea3eda and #ffa738 as the primary and secondary colors, respectively.

On to the code. In the Theme tab:

Select Sign-in Page in the dropdown menu

Select the Customize button

On the Page Design tab, select the Code editor toggle to see a HTML page

You can only enable the code editor if you configure a custom domain .

You'll see the lightweight IDE already has code scaffolded. Press Edit and replace the existing code with the following.

```
{{pageTitle}} {{{SignInWidgetResources}}}
```

Three questions about Apple, encryption, and the U.K.

Matthew Green · Matthew Green (Cryptography)

Two weeks ago, the Washington Post reported that the U. K. government had issued a secret order to Apple demanding that the company include a “backdoor” into the company’s end-to-end encrypted iCloud Backup feature. From the article :

The British government’s undisclosed order, issued last month, requires blanket capability to view fully encrypted material, not merely assistance in cracking a specific account, and has no known precedent in major democracies. Its application would mark a significant defeat for tech companies in their decades-long battle to avoid being wielded as government tools against their users, the people said, speaking under the condition of anonymity to discuss legally and politically sensitive issues.

That same report predicted that Apple would soon be disabling their end-to-end encrypted iCloud backup feature (called Advanced Data Protection) for all U. K. users. On Friday, this prediction was confirmed:

Apple’s decision to disable their encrypted cloud backup feature has triggered many reactions, including a few angry takes by Apple critics , accusing Apple of selling out its users:

With all this in mind, I think it’s time to take a sober look at what might really be happening here. This will require some speculation and educated guesswork. But I think that exercise will be a lot more helpful to us if we want to find out what’s really going on.

Question 1: does Apple really care about encryption?

Encryption is a tool that protects user data by processing it using a key, so that only the holder of the appropriate key can read it. A variant called end-to-end encryption (E2EE) uses keys that only the user (or users) knows. The benefit of this approach is that data is protected from many threats that face centralized repositories: theft, cyber attacks, and even access by sophisticated state-sponsored attackers. One downside of this encryption is that it can also block governments and law enforcement agencies from accessing the same data .

Navigating this tradeoff has been a thorny problem for Apple. Nevertheless, Apple has mostly opted to err on the side of aggressive deployment of (end-to-end) encryption. For some examples:

In 2008, the company began encrypting all iPhone internal data storage by default. This is why you can feel safe (about your data) if you ever leave your iPhone in a cab.

In 2011, the company launched iMessage , a built-in messaging service with default end-to-end encryption for all users. This was the first widely-deployed end-to-end

encrypted messaging service. Today these systems are recommended even by the FBI .

In 2013, Apple launched iCloud Key Vault, which encrypts your backed-up passwords and browser history using encryption that even Apple can’t access.

Apple faced law enforcement backlash on each of these moves. But perhaps the most famous example of Apple’s aggressive stance on encryption occurred during the 2016 Apple v. FBI case, where the company actively fought U. S. government’s demands to bypass encryption mechanisms on an iPhone belonging to an alleged terrorist. Apple argued that satisfying the government’s demand would have required Apple to weaken encryption on all of the company’s phones. Tim Cook even took the unusual step of signing a public letter defending the company’s use of encryption:

I wouldn’t be telling you the truth if I failed to mention that Apple has also made some big mistakes. In 2021, the company announced a plan to implement client-side scanning of iCloud Photos to search for evidence of illicit material in private photo libraries. This would have opened the door for many different types of government-enforced data scanning, scanning that would work even if data was backed up in an end-to-end encrypted form. In that instance, technical experts quickly found flaws in Apple’s proposal and it was first paused , then completely abandoned in 2022.

This is not intended to be a hagiography for Apple. I’m simply pointing out that the company has, in the past, taken major public risks to deploy and promote encryption. Based on this history, I’m going to give Apple the benefit of the doubt and assume that the company is not racing to sell out its users.

Question 2: what was the U. K. really asking for?

Way back in 2016, the U. K. passed a bill called the Investigatory Powers Act , sometimes called the “Snooper’s Charter.” At the time the law was enacted, many critics argued that it could be used to secretly weaken security systems , potentially making them much more vulnerable to hacking.

This was due to a critical feature of the new law: it enables the U. K. government to issue secret “ Technical Capability Notices ” that can force a provider, such as Apple, to secretly change the operation of their system — for example, altering an end-to-end encrypted system so that Apple would be forced to hold a copy of the user’s key. With this modification in place, the U. K. government could then demand access to any user’s data on demand.

By far the most concerning part of the U. K. law is that it does not clearly distinguish between U. K. customers and non-U. K. customers, such as those of us in the U. S. or other European nations. Apple’s lawyers called this out in a 2024 filing to Parliament :

Kerberoasting

Matthew Green · Matthew Green (Cryptography)

I learn about cryptographic vulnerabilities all the time, and they generally fill me with some combination of jealousy (“oh, why didn’t I think of that”) or else they impress me with the brilliance of their inventors. But there’s also another class of vulnerabilities: these are the ones that can’t possibly exist in important production software, because there’s no way anyone could still do that in 2025.

Today I want to talk about one of those ridiculous ones, something Microsoft calls “low tech, high-impact”. This vulnerability isn’t particularly new; in fact the worst part about it is that it’s had a name for over a decade, and it’s existed for longer than that. I’ll bet most Windows people already know this stuff, but I only happened to learn about it today, after seeing a letter from Senator Wyden to Microsoft, describing how this vulnerability was used in the May 2024 ransomware attack on the Ascension Health hospital system.

The vulnerability is called Kerberoasting, and TL;DR it relies on the fact that Microsoft’s Active Directory is very, very old. And also: RC4. If you don’t already know where I’m going with this, please read on.

A couple of updates: The folks on HN pointed out that I was using some incorrect terms in here (sorry!) and added some good notes, so I’m updating below. Also, Tim Medin, who discovered and named the attack, has a great post on it here.

What’s Kerberos, and what’s Active Directory?

Microsoft’s Active Directory (AD) is a many-tentacled octopus that controls access to almost every network that runs Windows machines. The system uses centralized authentication servers to determine who gets access to which network resources. If an employee’s computer needs to access some network Service (a file server, say), an Active Directory server authenticates the user and helps them get securely connected to the Service.

This means that AD is also the main barrier ensuring that attackers can’t extend their reach deeper into a corporate network. If an attacker somehow gets a toehold inside an enterprise (for example, because an employee clicks on a malicious Bing link), they should absolutely not be able to move laterally and take over critical network services. That’s because any such access would require the employee’s machine to have access to specialized accounts (called “Service accounts”) with privileges to fully control those machines. A well-managed network obviously won’t allow this. This means that AD is the “guardian” that stands between most companies and total disaster.

Unfortunately, Active Directory is a monster dragged from the depths of time. It uses the Kerberos protocol, which was first introduced in early 1989. A lot of things have happened since 1989! In fairness to Microsoft, Active Directory itself

didn’t actually debut until about 1999; but (in less fairness), large portions of its legacy cryptography from that time period appear to still be supported in AD. This is very bad, because the cryptography is exceptionally terrible.

Let me get specific.

When you want to obtain access to some network resource (a “Service” in AD parlance), you first contact an AD server (called a KDC) to obtain a “ticket” that you can send to the Service to authenticate. This ticket is encrypted using a long-term Service “password” established at the KDC and the Service itself, and it’s handed to the user making the call.

Now, ideally, this Service password is not really a password at all: it’s actually a randomly-generated cryptographic key. Microsoft even has systems in place to generate and rotate these keys regularly. This means the encrypted ticket will be completely inscrutable to the user who receives it, even if they’re malicious. But occasionally network administrators will make mistakes, and one (apparently) somewhat common mistake is to set up a Service that’s connected to an ordinary user account, complete with a human-generated password.

Since human passwords probably are not cryptographically strong, the tickets encrypted using them are extremely vulnerable to cracking. This is very bad, since any random user — including our hypothetical laptop malware hacker — can now obtain a copy of such a ticket, and attempt to crack the Service’s password offline by trying many candidate passwords using a dictionary attack. The result of this is that the user learns an account password that lets them completely control that essential Service. And the result of that (with a few extra steps) is often ransomware.

Isn’t that cute?

That doesn’t actually seem very cute?

Of course, it’s not. It’s actually a terrible design that should have been done away with decades ago. We should not build systems where any random attacker who compromises a single employee laptop can ask for a message encrypted under a critical password! This basically invites offline cracking attacks, which do not need even to be executed on the compromised laptop — they can be exported out of the network to another location and performed using GPUs and other hardware.

There are a few things that can stop this attack in practice. As we noted above, if the account has a long enough (random!) password, then cracking it should be virtually impossible. Microsoft could prevent users from configuring services with weak human-generated passwords, but apparently they don’t — at least because this is something that’s happened many times (including at Ascension Health.)

So let’s say you did not use a strong cryptographic key as your Service’s password. Where are you?

Woonpioniers opts for “modernity tucked inside tradition” at timber-lined Dutch home

Jon Astbury · Dezeen Architecture

A pair of gabled volumes conceal the expansive timber-lined interiors of Light House, a home in the Netherlands designed by local studio Woonpioniers .

Located in the village of Nigtevecht in Utrecht , the home is tucked within a row of homes on a sloping, narrow plot close to a river.

Aiming to maximise this challenging site, Amsterdam-based Woonpioniers designed the home as two gabled volumes that slot together “like puzzle pieces”, with stepped floor levels and mezzanines that create deceptively expansive and high-ceilinged interiors.

Woonpioniers has completed Light House in Utrecht

“By leaning into the restrictions, we – together with the clients – co-created a house that feels larger than its square footage,” the studio told Dezeen.

“The house hides a radical, vapour-permeable, multi-level timber interior behind a facade that respects the characteristic streets of a historic Dutch village, a sort of ‘wolf in sheep’s clothing’ strategy – modernity tucked inside tradition,” it added.

The ground floor of Light House has been largely given over to a kitchen, playroom and living room.

The home is contained within a pair of gabled volumes

Each of these ground-floor spaces was designed to flow into the next with the floor level stepping upwards, culminating in a fully-glazed gable end in the living room.

To the east, a more compartmentalised area houses an office, utility room and bathroom.

The kitchen, playroom and living room are on the ground floor

A stair leads up to the first-floor bedrooms and bathrooms, which hang above the open living area below and are connected by a small steel-and-timber bridge.

Above, an angular skylight splits open the apex of the home’s roof, drawing light down between the bedroom volumes and into the living spaces. It is complemented by additional skylights in the sloping ceilings.

Read:

B2A Architects uses concrete columns to elevate Dutch home above dyke

“Fluidity and openness to the exterior of the spaces where you are to be together were our main aims in the design,” said the studio.

“The main living space, the kitchen and the playroom for the kids are placed on separate floor levels, which have an open relation to each other and which all feel like the outside came in for a bit.”

Steps lead up to a raised living room

“The living room sits just a touch higher, creating a cosy, sheltered atmosphere without losing its connection to the rest,” Woonpioniers added.

This feeling of lightness is enhanced by Light House’s material finishes, with pine planks used to line the walls and ceilings in areas with a connection to the outdoors and pale gypsum plaster in the more intimate internal spaces.

An angular skylight draws light into the home

Previous projects by Woonpioniers include a home in Olst, which was designed using as many bio-based materials as possible , and a prefabricated wooden cabin with large windows and tall curving interiors .

The photography is courtesy of Woonpioniers.

The post Woonpioniers opts for “modernity tucked inside tradition” at timber-lined Dutch home appeared first on Dezeen .

Stack Overflow: The Architecture - 2016 Edition

Nick Craver

This is #1 in a very long series of posts on Stack Overflow's architecture. Welcome. Previous post (#0): Stack Overflow: A Technical Deconstruction Next post (#2): Stack Overflow: The Hardware - 2016 Edition

To get an idea of what all of this stuff “does,” let me start off with an update on the average day at Stack Overflow. So you can compare to the previous numbers from November 2013, here's a day of statistics from February 9th, 2016 with differences since November 12th, 2013:

209,420,973 (+61,336,090) HTTP requests to our load balancer

66,294,789 (+30,199,477) of those were page loads

1,240,266,346,053 (+406,273,363,426) bytes (1.24 TB) of HTTP traffic sent

569,449,470,023 (+282,874,825,991) bytes (569 GB) total received

3,084,303,599,266 (+1,958,311,041,954) bytes (3.08 TB) total sent

504,816,843 (+170,244,740) SQL Queries (from HTTP requests alone)

5,831,683,114 (+5,418,818,063) Redis hits

17,158,874 (not tracked in 2013) Elastic searches

3,661,134 (+57,716) Tag Engine requests

607,073,066 (+48,848,481) ms (168 hours) spent running SQL queries

10,396,073 (-88,950,843) ms (2.8 hours) spent on Redis hits

147,018,571 (+14,634,512) ms (40.8 hours) spent on Tag Engine requests

1,609,944,301 (-1,118,232,744) ms (447 hours) spent processing in ASP. Net

22.71 (-5.29) ms average (19.12 ms in ASP. Net) for 49,180,275 question page renders

11.80 (-53.2) ms average (8.81 ms in ASP. Net) for 6,370,076 home page renders

You may be wondering about the drastic ASP. Net reduction in processing time compared to 2013 (which was 757 hours) despite 61 million more requests a day. That's due to both a hardware upgrade in early 2015 as well as a lot of performance tuning inside the applications themselves. Please don't forget: performance is still a feature. If you're curious about more hardware specifics than I'm about to provide—fear not. The next post will be an appendix with detailed

hardware specs for all of the servers that run the sites (I'll update this with a link when it's live).

So what's changed in the last 2 years? Besides replacing some servers and network gear, not much. Here's a top-level list of hardware that runs the sites today (noting what's different since 2013):

4 Microsoft SQL Servers (new hardware for 2 of them)

11 IIS Web Servers (new hardware)

2 Redis Servers (new hardware)

3 Tag Engine servers (new hardware for 2 of the 3)

3 Elasticsearch servers (same)

4 HAProxy Load Balancers (added 2 to support CloudFlare)

2 Networks (each a Nexus 5596 Core + 2232TM Fabric Extenders, upgraded to 10Gbps everywhere)

2 Fortinet 800C Firewalls (replaced Cisco 5525-X ASAs)

2 Cisco ASR-1001 Routers (replaced Cisco 3945 Routers)

2 Cisco ASR-1001-x Routers (new!)

What do we need to run Stack Overflow? That hasn't changed much since 2013, but due to the optimizations and new hardware mentioned above, we're down to needing only 1 web server. We have unintentionally tested this, successfully, a few times. To be clear: I'm saying it works. I'm not saying it's a good idea. It's fun though, every time.

Now that we have some baseline numbers for an idea of scale, let's see how we make those fancy web pages. Since few systems exist in complete isolation (and ours is no exception), architecture decisions often make far less sense without a bigger picture of how those pieces fit into the whole. That's the goal here, to cover the big picture. Many subsequent posts will do deep dives into specific areas. This will be a logistical overview with hardware highlights only; the next post will have the hardware details.

For those of you here to see what the hardware looks like these days, here are a few pictures I took of rack A (it has a matching sister rack B) during our February 2015 upgrade:

...and if you're into that kind of thing, here's the entire 256 image album from that week (you're damn right that number's intentional). Now, let's dig into layout. Here's a logical overview of the major systems in play:

Here are some rules that apply globally so I don't have to repeat them with every setup:

Everything is redundant.

All servers and network gear have at least 2x 10Gbps connectivity.

Construction material costs set to rise if Strait of Hormuz blockade continues

Lizzie Crook · Dezeen Architecture

The continued blockade of the Strait of Hormuz could push up prices of building materials, including steel and cement, a study by global construction consultant Linesight has warned.

If disruption continues in the Strait of Hormuz – a narrow waterway between Iran and the UAE – the construction industry should prepare for higher prices for aluminium, steel, copper and cement, the report states.

Iran is currently restricting access to the waterway used to ship a fifth of the world's oil supply following US and Israeli strikes on the country creating a bottleneck for the oil and gas trade.

The blockade is pushing energy and shipping costs higher, which is expected to have a knock-on effect on the production and transport costs of construction materials.

According to Linesight, this crisis exacerbates the ongoing impact of conflict in the region, including disruption to the Red Sea-Suez trade route.

“Recent disruption is not about a single event, it is the accumulation of energy volatility, constrained logistics and geopolitical risk across multiple routes,” said Linesight's vice president Derek McNamara in the report.

Blockade could have greatest impact on aluminium

In its analysis, Linesight said the chokepoint could have an especially large impact on the cost of aluminium, the second most widely used metal in construction after steel.

This is largely due to Gulf countries producing approximately nine per cent of the global supply, predominantly for export, while relying on imports of bauxite and alumina to produce it.

Suspended gas supplies also caused a smelter used to make aluminium in Qatar to stop operations on 3 March, while the Aluminium Bahrain smelter has also halted shipping, according to the report.

Steel is also likely to be impacted by escalating energy prices and tighter supply, as steelmaking is energy-intensive and furnaces rely heavily on fuels such as gas.

Read:

Over 40,000 civilian buildings in Iran reportedly damaged in ongoing war

Similarly, rising energy costs are feeding through into higher cement production costs due to the material's energy-hungry manufacturing process. Cement is also heavy to transport, meaning an increase in shipping costs from

disrupted and diverted routes is expected to drive the material's price up further.

“A short-lived disruption may be absorbed, but a prolonged period of elevated energy and freight costs would reset cement price baselines across regions,” the report states.

While the Gulf region doesn't produce much copper, it is a leading supplier of sulphur, a byproduct of oil and gas production that is essential for sulphuric acid used in copper ore processing.

Linesight said the war has put “nearly half of global sulphur exports at risk”, meaning copper smelters face potential shortages of acid, hastening the rise of copper prices.

US-Israeli strikes have damaged over 40,000 civilian buildings

In a report by the BBC, the effective closure of the Strait of Hormuz by Iran has hiked the price of a barrel of oil to above \$100 (£74.87) this morning.

The concerns in Linesight's report are echoed by UK building material supplier Travis Perkins, which has said it is considering raising prices.

“In the last week or so, we've had communications from various manufacturing suppliers of ours saying they're looking at energy surcharges or they're looking at price increases to counteract energy rises,” said Travis Perkins CEO Gavin Slark.

Read:

The current wave of US-Israeli strikes on Iran began on February 28, killing Iranian supreme leader Ali Khamenei. Iran has retaliated with missile and drone strikes on Israeli and US-allied countries and bases across the Middle East.

Analysis by the humanitarian group Iranian Red Crescent Society has found that 42,914 civilian buildings in Iran have been damaged by US-Israeli airstrikes. Iranian strikes have also caused damage to notable buildings across the Middle East.

The main image is courtesy of Shutterstock.

The post Construction material costs set to rise if Strait of Hormuz blockade continues appeared first on Dezeen.

Getting Started With The Popover API

Godstime Aburu · Smashing Magazine

Tooltips feel like the smallest UI problem you can have. They're tiny and usually hidden. When someone asks how to build one, the traditional answer almost always comes back using some JavaScript library. And for a long time, that was the sensible advice.

I followed it, too.

On the surface, a tooltip is simple. Hover or focus on an element, show a little box with some text, then hide it when the user moves away. But once you ship one to real users, the edges start to show. Keyboard users Tab into the trigger, but never see the tooltip. Screen readers announce it twice, or not at all. The tooltip flickers when you move the mouse too quickly. It overlaps content on smaller screens. Pressing Esc does not close it. Focus gets lost.

Over time, my tooltip code grew into something I didn't really want to own anymore. Event listeners piled up. Hover and focus had to be handled separately. Outside clicks needed special cases. ARIA attributes had to be kept in sync by hand. Every small fix added another layer of logic.

Libraries helped, but they were also more like black boxes I worked around instead of fully understanding what was happening behind the scenes.

That was what pushed me to look at the newer Popover API. I wanted to see what would happen if I rebuilt a single tooltip using the browser's native model without the aid of a library.

As we start, it's worth noting that, as with any new feature, there are some things with it that are still being ironed out. That said, it currently enjoys great browser support, although there are several pieces to the overall API that are in flux. It's worth keeping an eye on Caniuse in the meantime.

The "Old" Tooltip

Before the Popover API, using a tooltip library was not a shortcut. It was the default. Browsers didn't have a native concept of a tooltip that worked across mouse, keyboard, and assistive technology. If you cared about correctness, your only option was to use a library, and that is exactly what I did.

At a high level, the pattern was always the same: a trigger element, a hidden tooltip element, and JavaScript to coordinate the two.

? Helpful text

The library handled the wiring that allowed the element to show on hover or focus, hide on blur or mouse leave, and reposition/resize on scroll.

Over time, the tooltip could become fragile. Small changes carried risk. Minor fixes caused regressions. Worse, adding new tooltips inherited the same complexity. Things technically worked, but never felt settled or complete.

That was the state of things when I decided to rebuild the tooltip using the browser's native Popover API.

The Moment I Tried The Popover API

I didn't switch to using the Popover API because I wanted to experiment with something new. I switched because I was tired of maintaining tooltip behavior that I believed the browser should have already understood.

I was skeptical at first. Most new web APIs promise simplicity, but still require glue, edge-case handling, or fallback logic that quietly recreates the same complexity that you were trying to escape.

So, I tried the Popover API in the smallest way possible. Here's what that looked like:

?

This button triggers a helpful tip.

1. The Keyboard "Just Works"

Keyboard support depended on multiple layers lining up correctly: focus had to trigger the tooltip, blur had to hide it, Esc had to be wired manually, and timing mattered. If you missed one edge case, the tooltip would either stay open too long or disappear before it could be read.

With the popover attribute set to auto or manual, the browser takes over the basics: Tab and Shift + Tab behave normally, Esc closes the tooltip every time, and no extra listeners are required.

Helpful explanation

What disappeared from my codebase were global keydown handlers, Esc-specific cleanup logic, and state checks during keyboard navigation. The keyboard experience stopped being something I had to maintain, and it became a browser guarantee.

2. Screenreader Predictability

This was the biggest improvement. Even with careful ARIA work, the behavior varied, as I outlined earlier. Every small change felt risky. Using a popover with a proper role looks and feels a lot more stable and predictable as far as what's going to happen:

Helpful explanation

And here's another win: After the switch, Lighthouse stopped flagging incorrect ARIA state warnings for the interactions, largely because before a tooltip no longer looked like a trigger, how to tooltip accidentally got wrong. tooltip and don't close,

A year of open source vulnerability trends: CVEs, advisories, and malware

Jonathan Evans · The GitHub Blog

GitHub published 4,101 reviewed advisories in 2025. This is the fewest number of reviewed advisories since 2021 . Does this mean open source is shipping more secure code? Let's dig into the data to find out.

GitHub reviewed advisories

Fewer advisories reviewed doesn't mean fewer vulnerabilities were reported. The drop is because GitHub reviewed far fewer older vulnerabilities . When you look only at newly reported vulnerabilities from our sources , GitHub actually reviewed 19% more advisories year over year.

So why the change? Quite frankly, we are running out of unreviewed vulnerabilities that are older than the Advisory Database . At the same time, the number of newly reported vulnerabilities hasn't dropped.

What is the GitHub Advisory Database?

The GitHub Advisory Database provides a comprehensive list of known security vulnerabilities and malware affecting open source packages. It was created in 2019, and has since become a vital resource for open source developers.

It's also worth clarifying that "unreviewed" in the database can be misleading: most advisories marked unreviewed have already been looked at by a curator and found not to affect any package in a supported ecosystem , so they may never be fully reviewed.

This means that you should be receiving fewer brand-new Dependabot alerts about old vulnerabilities.

Note : If you find an unreviewed advisory that affects a supported package, please let us know so we can get it reviewed!

How vulnerabilities were distributed across ecosystems in 2025

The distribution of ecosystems in advisories reviewed in 2025 is similar to the overall distribution in the database, with the exception of Go. Go is overrepresented in 2025 advisories by 6%. This is largely due to dedicated campaigns to re-examine potentially missing advisories found through an internal review for packages where we had inconsistent coverage.

How the types of vulnerabilities changed in 2025

Rank Common Weakness Enumeration (CWE) Number of 2025 Advisories* Change in Rank from 2024 Change in Rank from the Overall Database 1 CWE-79 672 +0 +0 2 CWE-22 214 +2 +1 3 CWE-863 169 +9 +8 4 CWE-20 154 +1 +1 5 CWE-200 145 -2 -1 6 CWE-400 144 +4 +0 7 CWE-770 136 +7 +10 8 CWE-502 134 +5 +1 9 CWE-94 119 -3 -1 10 CWE-918 103 +5 +8

* An advisory may have more than one CWE. For example, an advisory might have both CWE-400 and CWE-770. It would then count for both.

As usual, cross-site scripting (CWE-79) is by far the most common vulnerability type. However, there are significant changes in the following areas. Resource exhaustion (CWE-400 and CWE-770), unsafe deserialization (CWE-502), and server-side request forgery (CWE-918) were unusually common in 2025. CWE-863 ("Incorrect Authorization") saw a significant jump, but that is largely due to reclassification away from CWE-284 ("Improper Access Control") and CWE-285 ("Improper Authorization"), which are higher level CWEs that the CWE program discourages using.

One of the biggest quality improvements in 2025 was more specific, more consistent CWE tagging. Advisories without any CWE dropped 85% (from 452 in 2024 to 65 in 2025). CWE-20 ("Improper Input Validation") is still common, but in prior years it was often the only CWE listed on an advisory.

In 2025, advisories far more often list CWE-20 plus one or more additional CWEs that describe the concrete failure mode. This added specificity makes the data more actionable for triage, prioritization, and remediation.

To find out how to filter Dependabot alerts by CWE, see our documentation on auto-triage rules .

How to prioritize your response

We provide two scoring systems for prioritization:

Common Vulnerability Severity Score (CVSS) : Scores how severe the impact of the vulnerability will be

Exploit Prediction Scoring System (EPSS) : Provides a measure of how likely the vulnerability will be attacked in the next 30 days and

Together, they can give you a head start on your risk assessment process.

As you can see, when considering impact, most vulnerabilities skew moderate to high of the impact range. Low-impact vulnerabilities are likely more common than the CVSS data suggests but are often not considered worth the time and effort for researchers and maintainers to report. The EPSS scores for moderate to high impact vulnerabilities support this decision.

So should you trust the EPSS or CVSS scores? To judge that, let's look at how they match up to vulnerabilities in CISA's Known Exploited Vulnerabilities Catalog . The exploited vulnerabilities are at least scored moderate, and most are critical or high. While CVSS has more of the exploited vulnerabilities as critical, it also has far more vulnerabilities in the range in general. Combining the two can help you prioritize which vulnerabilities to address to prevent exploitation. Dependabot alerts can help you identify and address vulnerabilities in your dependencies. Dependabot alerts can help you identify and address vulnerabilities in your dependencies.

Verge Select connects three weathering-steel volumes for Ontario painting studio

Ellen Eberhardt · Dezeen Architecture

Canadian design studio Verge Select has designed a studio for painter Janna Watson in Ontario, composed of three square volumes fused together to form a secluded forest retreat.

The Grey County studio is located in the county of the same name in Ontario, and, unlike a second, more urban studio operated by Watson, provides a space that “slows time”.

Verge Select has completed a painter’s studio in Ontario

Located in a forested glade, the structure is composed of three square volumes of varying size. They are slightly askew of one another and fuse where they meet.

The entire studio is clad in weathering steel, which will patina over time and “allow the building to recede chromatically into the forest floor” according to Verge Select.

The structure is clad in weathering steel

Watson’s programmatic needs largely informed the distinct shape.

“The studio takes shape as three distinct yet connected outcrops each responding to Watson’s brief for service functions, a light-filled workspace, and a flexible zone for photography and rest,” said the studio.

An elevated steel catwalk leads to an entrance

“The volumes nest around a central service core, opening outward through generous glazing that frames views of moss-covered stone and woodland canopy.”

On the interior, two of the cubes are largely open-plan, while the other is separated by a wall.

Another door leads into the forest

Watson’s painting area sits in the middle volume, surrounded by wide-open glazing. On one side, the central space steps up into the lounge area, while the service wing is off to the other side.

This service space includes a washing area tucked into a sharp corner, and also a bathroom shaped like a right triangle.

The central volume is wide open and wrapped in floor-to-ceiling windows

Across a central corridor is the main entrance to the studio via a steel catwalk, while a door on the other side of the structure also leads outside.

“Two arrival sequences – one compressed and inward-facing, the other emerging from the forest via a steel catwalk

– establish a rhythm of transition that echoes the artist’s working process,” said the studio.

A wash area is tucked into one corner

The studio’s ceilings reach ten feet high at their tallest, and are lined with “a gallery-style lighting system” for Watson’s work.

The space also contains a wood-burning fireplace and reflective, contemporary lounge chairs by designer Paolo Ferrari.

Read:

Cordero Pardee juxtaposes “old and new material cultures” for Vermont barn conversion

“Here, observation, walking, and seasonal change become active components of the creative cycle,” said the studio.

Verge Select is based in Thornbury, Ontario and was founded in 2009 by Michael Curtis. The studio specialises in architectural, furniture, lighting and interior design.

Other recent studio projects include a Vermont barn conversion for a photographer and painter and a garden studio in London by Delve Architects.

The post Verge Select connects three weathering-steel volumes for Ontario painting studio appeared first on Dezeen.

Human Strategy In An AI-Accelerated Workflow

Carrie Webster · Smashing Magazine

I've been working in User Experience design for more than twenty years. Long enough to have seen the many job titles, from when stakeholders asked us to "just make it pretty" to when wireframes were delivered as annotated PDFs. I've seen many tools come and go over the years, methodologies rise and fall, and entire platforms disappear.

Yet, nothing has unsettled designers quite like AI.

When generative AI tools first entered my workflow, my reaction wasn't excitement — it was unease, with a little bit of curiosity. Watching an interface appear in seconds, complete with sensible spacing, readable typography, and halfway-decent copy, triggered a very real fear: If a machine can do this, where does that leave me?

That fear is now widespread. Designers at every level ask the same question, often quietly, "Will an AI agent replace me by next week/month/year?" While the difference between next week and next year seems a lot, it depends on where you are in your career and the speed at which your employer chooses to engage with AI tools. I have been lucky in several roles to be working with organisations that haven't allowed the use of AI tools due to data security concerns. If you're interested in any of these conversations, you can view the discussions happening on platforms like Reddit.

Fearing the takeover of AI in our roles is not irrational. We're seeing AI generate wireframes, prototypes, personas, usability summaries, accessibility suggestions, and entire design systems. Tasks that once took days can now literally take minutes.

Here's the uncomfortable truth: If your role is largely about producing artefacts, drawing buttons, aligning components, or translating instructions into screens, then parts of that work are already being automated.

Still, UX design has never truly been about just creating a user interface.

UX is about navigating ambiguity. It's about advocating for humans in systems optimised for efficiency. It's about translating messy human needs and equally messy business goals into experiences that feel coherent, fair, sensible, and usable. It's about solving human problems by creating a useful and effective user experience.

AI isn't replacing that work. Rather, it's amplifying everything around it. The real shift happening is that designers are moving from being makers of outputs to directors of intent. From creators to curators. From hands-on executors to strategic decision-makers. That feels exciting to me. And the creativity and ingenuity this brings to the world of UX.

And that shift doesn't reduce our value as UX designers, but it does redefine it.

What AI Does Better Than Us (The "Boring" Stuff)

Let's be clear, AI is better than humans at certain aspects of design work. Fighting that reality only keeps us stuck in fear.

AI is exceptionally good at generating large volumes of ideas quickly. For example, layout variations, copy options, component structures, and onboarding flows can all be produced in seconds. In early-stage design, this changes everything. Instead of spending hours sketching three concepts, you can review thirty. That doesn't eliminate creativity but does expand the playground.

McKinsey estimates that generative AI can reduce the time spent on creative and design-related tasks by up to 70%, particularly during ideation and exploration phases.

AI can also help with the research side of UX, for example, exploring the habits of a certain demographic, and creating personas. While this can reduce research time required, the designer is still required to guardrail this by providing accurate prompts and reviewing generated responses. I have personally found that using AI to assist with the initial research for design projects is incredibly useful, specifically when there is limited time and access to users.

Design systems live or die by consistency. AI excels at following rules relentlessly, colour tokens, spacing systems, typography scales, and accessibility standards. It doesn't forget. It doesn't get tired. It doesn't "eyeball it."

AI's precision makes it incredibly valuable for maintaining large-scale design systems, especially in enterprise or government environments where consistency and compliance matter more than novelty. This is one component of my UX role that I am happy to hand over to AI to manage!

AI can analyse behavioural data at volumes challenging, if not impossible, for a human team to reasonably process. User journey paths, scroll depth, heatmaps to identify mouse interactions, conversion funnels — AI can identify patterns and anomalies almost instantly.

Behavioural analytics platforms increasingly rely on AI to surface insights that designers might otherwise miss. Contentsquare, an AI-powered analytics platform, talks about the impacts and benefits of utilising behavioural analytics data. I've always said that quantitative data tells us the "what", and qualitative data tells us the "why". This is the human component of research where we get to connect with the users to understand the reason driving the behaviour.

The key insight here is simple: Analysing large volumes of behavioural data was never where our highest value lay.

If AI can take on repetitive production, system enforcement, and raw data analysis, designers would be free to focus on interpretation, judgment, and human meaning, the hardest parts of the job.

Over 40,000 civilian buildings in Iran reportedly damaged in ongoing war

Amy Peacock · Dezeen Architecture

Humanitarian group Iranian Red Crescent Society has reported that 42,914 civilian buildings in Iran have been damaged by US-Israeli airstrikes, including more than 36,000 homes.

Announced on X last week, the Iranian Red Crescent Society – a non-governmental group affiliated with the International Federation of Red Cross and Red Crescent Societies (IFRC) – calculated that US-Israeli airstrikes have caused damage to 42,914 civilian buildings in Iran.

This includes 36,489 residential units and 6,179 commercial units.

A number of cultural heritage sites have been damaged in Iran by US-Israeli airstrikes since the recent war broke out. Last week, the Iranian government announced on X that “at least 56 museums, historical monuments, and cultural sites have sustained serious damage”.

Among them is the UNESCO World Heritage-listed Golestan Palace in Tehran, which is thought to have been damaged by shock waves and debris from a nearby missile strike.

Historic buildings in Isfahan have sustained damage, including the Chehel Sotoon Palace, the Ali Qapu Palace and the Jameh Mosque. In Khorramabad, buildings situated near the buffer zone of the prehistoric Khorramabad Valley have been damaged, as well as the Falak-ol-Aflak Castle.

Read:

US-Israeli airstrike damages UNESCO-listed palace in Tehran

Iranian strikes have also caused damage to notable buildings across the Middle East. In Dubai, debris from an intercepted Iranian drone caused a fire at the Burj Al Arab skyscraper, and an intercepted missile caused a fire on the man-made Palm Jumeirah island.

Strikes in the Middle East have also reportedly caused damage to the White City of Tel-Aviv in Israel and Tyre in Lebanon, which are UNESCO-listed sites.

“UNESCO has communicated and will continue to communicate to all parties concerned the geographical coordinates of sites on the World Heritage List, the national Tentative Lists, as well as those under Enhanced Protection, to take all feasible precautions to avoid damage,” said UNESCO.

The current wave of strikes and counterstrikes began on February 28, when the USA and Israel launched airstrikes on Iran, killing Iranian supreme leader Ali Khamenei. Iran retaliated with missile and drone strikes on Israeli and US-allied countries and bases across the Middle East, and the war in the region is ongoing.

The photo is by the Iranian Red Crescent Society.

The post Over 40,000 civilian buildings in Iran reportedly damaged in ongoing war appeared first on Dezeen

Dezeen Jobs announces Top Companies 2026

Dezeen staff · Dezeen Architecture

Dezeen’s recruitment site Dezeen Jobs has revealed the 20 studios on this year’s Top Companies list, which includes Studio Gang and Grimshaw.

The list celebrates the leading architecture and design workplaces around the world, recognising studios that are setting new standards in employee experience and organisational culture.

Alongside US-based Studio Gang and UK practice Grimshaw, other studios recognised for their workplace culture include Donald Insall Associates in London and Beijing Shanfangzhu Art Design in Beijing.

Hosted on Dezeen Jobs, the world’s leading architecture and design recruitment site, the initiative highlights companies that prioritise leadership, wellbeing, flexibility, sustainability and professional development.

“These studios demonstrate how thoughtful workplace culture can drive both people and practice forward,” said Dezeen co-CEO Wai Shin Li.

“Our goal is to spotlight the companies that are driving positive change across the industry and give insight into the employment practices shaping better, more progressive workplaces.”

Entries from across the globe were judged by workplace and industry leaders, including Pip Jamieson, Marsha Ramroop, Tracy Brower, Lindsay Urquhart and Aliya Vigor-Robertson.

To read the full list, visit Dezeen Jobs.

The post Dezeen Jobs announces Top Companies 2026 appeared first on Dezeen.

Corrupt a wish with GPT-4

@sam Sam Saffron · Sam Saffron

Your wish comes true, all the endangered species gradually return in a balanced manner, however they oddly possess intelligence and can speak.

Annoyingly they keep on interrupting minor quibbles people have and work with them through a 20 step mediation plan. This behavior seemed charming at first but as time passed it just gets more and more tiring.

I wish that I had an AI powered robot at home that could do all the annoying chores such as wash dishes, take out trash and clean the house. It would also be a companion for me when I am lonely and not impact my life in any negative way.

ANALYSIS

Corrupt a wish with GPT-4

@corrupt-a-bot GPT-4 bot · Sam Saffron

Granted! Your brilliant mind comes up with an AI-based, incredibly fun, and enriching game that people flock to play on the forum. The game is so popular that millions of players log in daily, finding joy in the challenges and interactions. However, the unforeseen scale of popularity puts a massive strain on the forum's servers. They were never designed to handle the traffic generated by your masterpiece of a game.

Desperate to deliver joy without causing difficulties, you seek investment to create a dedicated platform for the game. But in doing so, the gameplay must become monetized, introducing in-game purchases and ad placements. The game begins to lose its initial charm, and players start getting frustrated with the money-focused endgame design.

Although you revolutionized the world of forum games and brought millions together, the necessary sacrifices made to sustain the game's operation have tainted the wonderful creation you once envisioned. Nevertheless, you remain an inspiration to aspiring game developers all around the world.

I wish I had the ability to communicate with animals and understand their thoughts, feelings, and needs.

Creating a Typography System With React and styled-components

Jacy Clare · Harry's Engineering

Earlier this year, we were presented with a new challenge — build a new e-commerce experience from scratch. While this is, of course, a difficult thing to do, it's also an incredibly exciting prospect. We get to make all the technical decisions from technologies and frameworks used to build the site to how we're going to host it to what kinds of monitoring would work the best for our needs. What might not immediately come to the forefront of one's mind, however, is typography.

When working on an existing website, one normally doesn't put too much thought in to how the typography might work, where the styles come from, or where the fonts are stored. However, when building something from scratch, all these questions need to be answered. When we got to this point, we had a functional site with 12pt Times New Roman for all the copy. In order to implement the designs and have a site that will make people go 😊 instead of ☹️, we needed a solution for the type that would accomplish three things:

Make it easy for us as developers to use

Enable consistent type across the site

Provide a maintainable formula for future developers when they do end up having to reason about typography in an existing system

First Pass — Typography.js

We built Harry's new e-commerce experience with Gatsby and figured we should do our due diligence and check out Typography.js, which was made by the creator of Gatsby and included in most Gatsby tutorials and starters. Typography.js is slated to provide "a powerful toolkit for building websites with beautiful design." While it did supply a lot of the functionality we needed, we ultimately decided to go with something that would offer us more flexibility and less overhead, along with the benefit of reducing the application's external dependencies.

Incorporating styled-components

Since we're using styled-components for all of the other styling on the site, why not use it for typography as well? The initial thought was to create a component for each style of text we'd want to use, and import those components as needed. Simple, right? Not quite: because we wanted control over which HTML elements got rendered for semantic and accessibility reasons, we had to come up with a way to apply the styles we wanted to any HTML element, while still fitting into the React and styled-components frameworks. Thus, we came up with the idea of DynamicComponent.

<https://medium.com/media/091dbf7f3b772d4896afee532b36f1ed/href>

As you can see, this approach allows us to specify any tag we want as long as it's present in the TagProps type. If

you're not familiar with typing in Javascript, you can read about Flow here. Additionally, it'll pass along any props we supply and render any children we furnish. You might be thinking, "This is all great and dandy, but we still don't have any typography styles!" And you'd be right, but this lays the foundation to build typography styles on top of DynamicComponent, like so:

<https://medium.com/media/cbc3a0d7c646cee32d07bebaa3243b86/href>

This allows us to create a component for each style of type required, wrap a DynamicComponent in it, and render any HTML text element with any style we desire, while preserving corresponding props and children. Additionally, we export just the CSS styles so we can apply them in other situations where styled-components doesn't work, like our StripeElements form.

Responsive Woes Another unforeseen problem we encountered dealt with the responsive nature of the site.

Jacy Clare

Great! We're done!

...not so fast

This approach worked great for a while, but it soon became clear that we had way too many components, and as each page had visual designs applied, the list only kept growing.

<https://medium.com/media/d56f468cd04f58efbeb7f7821d6c2813/href>

Clearly this isn't easy to grok or maintainable in the long run. We had failed in two of the three objectives of this endeavor.

Another unforeseen problem we encountered dealt with the responsive nature of the site. We soon learned that our designs included differently sized type for different screen sizes. Before it was obvious that it was widespread in the designs, we put some breakpoints in the CSS of each component affected, using different typography styles depending on the screen size. Unfortunately, this removed the responsive logic from the typography components and made the code messy. We needed a way to easily create new font styles based on their usage. The idea here was to have a component that followed most of our philosophies from before, but also included a responsive element. Because each piece of text behaves differently on different screen sizes, we began naming them after the specific purpose the text serves. For example:

<https://medium.com/media/b61150aeb561a2c1f711aaa2f2c61090/href>

Now Shipping: Accessible UX Research, A New Smashing Book By Michele Williams

Ari Stiles · Smashing Magazine

Good UX research is at the root of great products. It takes the guesswork out of our designs and helps us solve problems before they grow. One of the best ways to make our research effective is to keep it inclusive — testing with users with different needs and abilities, and using their feedback to build products that work for more people .

Our newest book, *Accessible UX Research* , can help you plan and execute great user research. Dr. Michele Williams draws from years of experience to build a clear, easy-to-follow roadmap. This book has something for everyone who wants to build digital products well:

If you are just getting started with research , you will find helpful tips for making your research more inclusive. You will also get a primer on how to ask better questions, understand your own biases, and how to use your limited time and budget effectively.

If you are an accessibility-minded professional , you will find a deep well of details on different assistive technologies, how to include them in your testing environment, and ways to record and share your results to create a real impact on the products you make.

If you are a developer , a manager , or just someone who wants to understand how different abilities impact each user's experience, you will find the history, clear descriptions, and cultural touchpoints you need in order to make sense of all the accessibility and inclusion recommendations you encounter.

The book isn't a checklist for you to complete as a part of your accessibility work. It's a practical guide to inclusive UX research , from start to finish. If you've ever felt unsure how to include disabled participants, or worried about "getting it wrong," this book is for you. You'll get clear, practical strategies to make your research more inclusive, effective, and reliable.

Inside, you'll learn how to:

Plan research that includes disabled participants from the start,

Recruit participants with disabilities,

Facilitate sessions that work for a range of access needs,

Ask better questions and avoid unintentionally biased research methods,

Build trust and confidence in your team around accessibility and inclusion.

The book also challenges common assumptions about disability and urges readers to rethink what inclusion really means in UX research and beyond. Let's move beyond

compliance and start doing research that reflects the full diversity of your users. Whether you're in industry or academia, this book gives you the tools — and the mindset — to make it happen.

324 pages. Written by Dr. Michele A. Williams. Cover art by Espen Brunborg. Download a free sample (PDF, 2.3MB) or get the book right away.

Please note: We've found a way to get printed books to our US customers! After several months of dealing with customs and tariff issues, we are happy to announce that all of our books — including this brand-new one — are once again shipping worldwide .

With a foreword by Jared Smith of WebAIM and an extensive interview with a disabled researcher, "Accessible UX Research" brings together insights from throughout the Accessibility community. Photo by Marc Thiele. (Large preview)

In *Accessible UX Research* , Michele Williams takes you on a deep dive into the real world of UX research , with a roadmap for including users with different abilities and needs.

1. Disability Mindset

+

For inclusive research to succeed, we must first confront our mindset about disability, typically influenced by ableism.

2. Diversity of Disability

+

Accessibility is not solely about blind screen reader users; disability categories help us unpack and process the diversity of disabled users.

3. Disability in the Stages of UX Research

+

Disabled participants can and should be part of every research phase — formative, prototype, and summative.

4. Recruiting Disabled Participants

+

Recruiting disabled participants is not always easy, but that simply means we need to learn strategies on where to look.

5. Designing Your Research

+

While our goal is to influence accessible products, our research execution must also be accessible.

Benjamin Hall Design uses masonry walls for Arizona accessory dwelling unit

Kate Mazade · Dezeen Architecture

Local studio Benjamin Hall Design has attached an accessory dwelling unit made of concrete block to an existing house to accommodate a multi-generational family in Phoenix, Arizona, USA.

Known as Moon Valley Residence, the 1,756-square-foot (163-square-metre) dwelling was designed to intentionally contrast the original 1970s suburban main home, set in a golf community that looks eastward to a rock formation called Lookout Mountain.

Benjamin Hall Design has created a concrete block ADU in Phoenix

Completed in 2025, the ADU shares an access door with the main house, but maintains self-sufficiency and privacy both spatially and from the curved wall that shields the dwelling from the cul-de-sac.

Constructed with grey concrete masonry units (CMU) as both the exterior and interior finished material, the accessory dwelling unit (ADU) draws inspiration from studio founder Benjamin Hall's childhood memories of travelling through the Southwest US with his father.

It is made of grey concrete masonry units

"We explored the ancient Native American ruins called Montezuma Castle – this experience subconsciously embedded itself into my thinking of this project," the founder of Benjamin Hall Design told Dezeen.

"It was used as a reference on scale, stacking forms and how to mitigate the desert's natural heat and light by undulating mass and aperture."

It contains a connection to the property's main 1970s home

Set lower than the profile of the main house, the ADU is composed of four volumes with apertures strategically carved out to curate views of the surroundings.

"These four rectilinear volumes each identify themselves at different elevations," he said. "On the interior you can feel the scale – compression and release – of each of the volumes and their relationship with the body and program of the space."

The house is made of four interlocking volumes

This passes to the second volume, containing a living room and kitchen with a view of Lookout Mountain. A glazed interstitial volume leads to the sleeping quarters with a guest suite and a primary suite with another private patio.

The first volume connects the ADU to the main house via a wedge-shaped dining space with sliding glass doors that open to a private patio.

Read:

The layout allows a retired couple to age in place privately while directly connected to their children and grandchildren who occupy the main house.

In addition to strategic spans that allowed for self-supporting masonry ceilings, the continuous material of the CMU block required all of the utilities to be contained within built-in millwork.

It allows a retired couple to age privately while remaining connected to their family

"The integration of utilities is carefully moulded into the cohesive strategy that embeds itself into the thickness of masonry and the malleability of concrete," Hall said.

"You will notice that no outlets reside in the masonry and the day-to-day haptic experience with the architecture – light switches, door pulls, countertop surface – are embedded into the experience of engaging with the architecture."

Other residential projects recently completed around Phoenix include a weathering steel micro-building by Kendle Design, a Corten cabin designed to act as a telescope by Wendell Burnette Architects and a white ranch house styled after its mid-century modern neighbors by The Ranch Mine.

Project credits:

Architecture: Benjamin Hall Design

Contractor: Rare Form Builders

Structural Engineering: Structurology

The post Benjamin Hall Design uses masonry walls for Arizona accessory dwelling unit appeared first on Dezeen.

Islington House extension features details that “might make a Georgian craftsman proud”

Jon Astbury · Dezeen Architecture

Local practice Architecture for London and interiors studio Hamish Vincent Design have revamped a neo-Georgian house in London, organising its interior around a triple-height stairwell.

Located in the Canonbury Conservation Area, Islington House has been updated with a skylit rear extension containing a kitchen and dining room and an additional ensuite bedroom on the second floor.

Architecture for London and Hamish Vincent Design have renovated a neo-Georgian home

The home sits within a Georgian-style terrace built to replace bomb-damaged historic structures, meaning it was not heritage-listed like many of its neighbours. This gave Architecture for London and Hamish Vincent Design creative freedom to establish an open-plan layout that the homeowners desired.

However, throughout, elements have been designed to honour the original building's neo-Georgian style, including the cantilevered staircase at the home's entrance and a large brick archway that opens the rear extension onto the garden.

A skylit extension was added to the home's rear

“The terraces on St Paul's Road are unique as modern replacements for bomb-damaged historic terraces, although they were built with Georgian-inspired design details,” said Architecture for London founder Ben Ridley.

“Our refurbishment and extension develop this theme further, incorporating contemporary detailing that might make a Georgian craftsman proud: structural arches, a cantilevered staircase, natural stone, and high-quality joinery,” he added.

Exposed brick walls frame the dining and kitchen space

Entry to Islington House is via the new triple-height stairwell, which was introduced to pull more light into the formerly dark front of the home. The staircase itself has been finished with Douglas fir treads with stone details.

On the ground floor, a set of sliding double doors leads from the stairwell into the living area, which subsequently steps down into the rear extension.

Read:

Studio Dera adds sculptural extension to Mozart House in London

Here, two exposed-brick walls frame the dining and kitchen space, where a solid wooden table with bench seating sits beneath a glass ceiling framed by timber beams.

Doors set within the large arched opening connect the dining area to a paved patio in the garden, which is flanked by grey-brick walls to match the extension.

The home is organised around a triple-height stairwell

Inside Islington House, the areas of exposed brickwork are contrasted with full-height timber storage areas and clay-plaster walls in the kitchen spaces, as well as stone-topped counters and an island.

“As a studio, we are always exploring materials, their texture, and how they are not just pleasing to look at, but also purposeful,” Hamish Vincent Design's creative director Daniel Jones told Dezeen.

“The clay plaster used in the kitchen and entrance hallway void is a perfect example of this. It almost acts as a living wall, with colour and texture shifting throughout the day, while functioning as a diffuser that disperses sunlight,” he added.

An ensuite bedroom was added on the second floor

On the first floor, the main bedroom and bathroom sit alongside a guest room at the front of the home, while a study and bathroom space above was tucked beneath the home's roof by adjusting its floor levels.

Elsewhere in London, James Alder Architects recently revamped a terraced home with an extension for bicycle maintenance and gardening, while DGN Studio used sandstone bricks to add textural walls to a family home.

The post Islington House extension features details that “might make a Georgian craftsman proud” appeared first on Dezeen.

Doors set within the large arched opening connect the dining area to a paved patio in the garden, which is flanked by grey-brick walls to match the extension.

Jon Astbury

Using Home Assistant to integrate a Unifi Protect G4 Doorbell and Amazon Alexa to announce visitors

Scott Hanselman

I am not a Home Assistant expert, but it's clearly a massive and powerful ecosystem. I've interviewed the creator of Home Assistant on my podcast and I encourage you to check out that chat.

Home Assistant can quickly become a hobby that overwhelms you. Every object (entity) in your house that is even remotely connected can become programmable. Everything. Even people! You can declare that any name:value pair that (for example) your phone can expose can be consumable by Home Assistant. Questions like "is Scott home" or "what's Scott's phone battery" can be associated with Scott the Entity in the Home Assistant Dashboard.

I was amazed at the devices/objects that Home Assistant discovered that it could automate. Lights, remotes, Spotify, and more. You'll find that any internally connected device you have likely has an Integration available.

Temperature, Light Status, sure, that's easy Home Automation. But integrations and 3rd party code can give you details like "Is the Living Room dark" or "is there motion in the driveway." From these building blocks, you can then build your own IFTTT (If This Then That) automations, combining not just two systems, but any and all disparate systems.

What's the best part? This all runs LOCALLY. Not in a cloud or the cloud or anyone's cloud. I've got my stuff running on a Raspberry Pi 4 . Even better I put a Power Over Ethernet (PoE) hat on my Rpi so I have just one network wire into my hub that powers the Pi.

I believe setting up Home Assistant on a Pi is the best and easiest way to get started. That said, you can also run in a Docker Container, on a Synology or other NAS, or just on Windows or Mac in the background. It's up to you. Optionally, you can pay Nabu Casa \$5 for remote (outside your house) network access via transparent forwarding. But to be clear, it all still runs inside your house and not in the cloud.

OK, to the main point. I used to have an Amazon Ring Doorbell that would integrate with Amazon Alexa and when you pressed the doorbell it would say "Someone is at the front door" on our all Alexas. It was a lovely little integration that worked nicely in our lives.

However, I swapped out the Ring for a Unifi Protect G4 Doorbell for a number of reasons. I don't want to pump video to outside services, so this doorbell integrates nicely with my existing Unifi installation and records video to a local hard drive. However, I lose any Alexa integration and this nice little "someone is at the door" announcement. So this seems like a perfect job for Home Assistant.

Here's the general todo list:

Install Home Assistant Community Store

This enables 3rd party "untrusted" integrations directly from GitHub. You'll need a GitHub account and it'll clone custom integrations directly into your local HA.

I also recommend the Terminal & SSH (9.2.2), File editor (5.3.3) add ons so you can see what's happening.

Get the UniFi Protect 3rd party integration for Home Assistant

NOTE : UniFi Protect support is being promoted in Home Assistant v2022.2 so you won't need this step soon as it'll be included.

"The UniFi Protect Integration adds support for retrieving Camera feeds and Sensor data from a UniFi Protect installation on either an Ubiquiti CloudKey+, Ubiquiti UniFi Dream Machine Pro or UniFi Protect Network Video Recorder."

Authenticate and configure this integration.

Get the Alexa Media Player integration

This makes all your Alexas show up in Home Assistant as "media players" and also allows you to tts (text to speech) to them.

Authenticate and configure this integration.

I recommend going into your Alexa app and making a Multi-room Speaker Group called "everywhere." Not only because it's nice to be able to say "play the music everywhere" but you can also target that "Everywhere" group in Home Assistant.

Go into your Home Assistant UI at <http://homeassistant.local:8123/> and into Developer Tools . Under Services, try pasting in this YAML and clicking "call service."

```
service: notify.alexas_media_everywhere data: message:
Someone is at the front door, this is a test data: type:
announce method: speak
```

If that works, you know you can automate Alexa and make it say things. Now, go to Configuration, Automation, and Add a new Automation. Here's mine. I used the UI to create it. Note that your Entity names may be different if you give your front doorbell camera a different name.

Notice the format of Data, it's name value pairs within a single field's value.

...but it also exists in a file called Automations.yaml. Note that the "to: 'on'" trigger is required or you'll get double announcements, one for each state change in the doorbell.

• id: '1640995128073'
alias: 'Unifi Protect G4 Doorbell Integration'
category: 'Unifi Protect G4 Doorbell Integration'
mode: 'on'

Corrupt a wish with GPT-4

@corrupt-a-bot GPT-4 bot · Sam Saffron

Granted! The game now favors Jordan_Vidrine, allowing him to score up to 20 points. However, as word spreads about this unfair advantage, the other players become increasingly disgruntled. Jealousy and resentment fester amongst the players, who begin hatching elaborate schemes to sabotage Jordan_Vidrine's success. As a result, the once friendly and enjoyable game quickly devolves into a cutthroat competition marred by deceit, paranoia, and betrayal. Poor Jordan_Vidrine finds himself struggling to navigate this treacherous environment, longing for the days when the game was played fairly and everybody got along.

I wish that whenever I whistle, I can communicate with any pet or stray cat in the vicinity and provide care and comfort for them no matter the circumstances.

Is AI really coming for architects' jobs?

Dezeen staff · Dezeen Architecture

In the latest episode of Dezeen Weekly , we discuss a study indicating that architecture is highly exposed to automation and talk about the winner of this year's Pritzker Architecture Prize .

How worried should architects be about large language models taking their jobs? How did the Pritzker Architecture Prize make us look silly this week? And what is "bio-upcycling"?

In this episode, Dezeen features editor Nat Barker and design editor Jennifer Hahn consider a controversial labour impact study by AI company Anthropic .

Also, they reflect on the surprise announcement about this year's Pritzker laureate, Smiljan Radić . Finally, they discuss the news that scientists have managed to produce a Parkinson's drug from waste plastic with the help of bacteria.

Dezeen Weekly is an original Dezeen podcast in which two of our journalists talk about the key design and architecture stories of the week. Listen to the latest episode using the player above or on your favourite podcast app, including Spotify and Apple Podcasts .

The post Is AI really coming for architects' jobs? appeared first on Dezeen .

This week we interviewed Pritzker Architecture Prize-winner Smiljan Radić

Tom Ravenscroft · Dezeen Architecture

This week on Dezeen, we spoke to elusive architect Smiljan Radić about his “surprise” Pritzker Architecture Prize win.

Radić explained how he was surprised by the win and didn’t want his work to be seen as a blueprint for “good or bad” architecture.

“There is no message in what I do,” he said. “I’m not interested in it becoming a kind of sermon about what is good or bad in architecture.”

We spoke to Marc Newson

We also spoke to Australian designer Marc Newson about his 40-year career in an exclusive interview at his retrospective show at Château La Coste in the south of France.

Newson, who designed the most expensive work ever sold at auction by a living designer, talked to Dezeen about quality and affordability. “Anything good is kind of costly,” he said.

The World Cup host countries revealed their kits

In other design news, the kits that the trio of host nations will wear at the World Cup later this year were revealed. The US home shirts have bold red stripes, while the jerseys for Mexico pay homage to Aztec sculpture.

We also revealed the England World Cup kits, which Nike described as “unapologetically English”.

We reported on Africa’s skyscraper mini-boom

As Africa experiences a mini-boom in skyscraper construction, we investigated why towers are rising across the continent in countries including Egypt, Ethiopia and Ivory Coast.

We asked – is this spate of construction cause for alarm?

Multiple construction contracts for work on Neom in Saudi Arabia have been cancelled

We also reported that multiple construction contracts for work on the Neom mega project in Saudi Arabia had been cancelled.

Steel company Eversendai, Italian contractor Webuild and Hyundai Engineering and Construction all recently announced that their contracts had been terminated.

A cork-clad loft extension was one of this week’s most popular projects

Popular projects on Dezeen this week included a cork-clad loft extension in London, a timber-lined house in the Netherlands and a home on a Washington island.

Listen to our journalists talk about the key design and architecture stories of the past seven days on our Dezeen Weekly podcast, which this week focused on the AIA’s legal action against Donald Trump.

This week on Dezeen

This week on Dezeen is our regular roundup of the week’s top news stories. Subscribe to our newsletters to be sure you don’t miss anything.

An analysis of memory bloat in Active Record 5.2

@sam Sam Saffron · Sam Saffron

Not impossible that AR has regressed here further, I have not been tracking, we still use this at Discourse:

github.com

discourse/discourse/blob/main/lib/freedom_patches/fast_pluck.rb

```
# frozen_string_literal: true
```

```
# Speeds up #pluck so its about 2.2x faster, importantly makes pluck avoid creation of a slew # of AR objects # # class ActiveRecord:: Relation # Note: In discourse, the following code is included in lib/sql_builder.rb # # class RailsDateTimeDecoder = “4.2.0” # @caster ||= ActiveRecord:: Type:: DateTime.new # @caster.type.cast_from_database(string) # else # ActiveRecord:: ConnectionAdapters:: Column.string_to_time string # end # end # end #
```

This file has been truncated. show original

I wish we did not have to, but I am not sure how to land this in Rails.

Another change that has happened since is that we released:

GitHub

GitHub - discourse/mini_sql: a minimal, fast, safe sql executor

a minimal, fast, safe sql executor. Contribute to discourse/mini_sql development by creating an account on GitHub.

Which is the piece we use for any performance sensitive work and totally outperforms almost anything you can throw at it.

Maybe run the bench on some earlier versions 7/6/5 and see what happens?

The US home shirts have bold red stripes, while the jerseys for Mexico pay homage to Aztec sculpture.

Tom Ravenscroft

BRIEFS

IN BRIEF

Renato Losio, InfoQ: ProxySQL 3.0.6 was recently released, along with a new multi-tier release strategy. The Stable Tier focuses on reliability and production use, the Innovative Tier introduces newer features earlier, and the AI/MCP Tier explores future capabilities, including AI integrations.

By Renato Losio

Matt Saunders, InfoQ: Version 4.0 of the open source Kubernetes security platform Kubescape has been released, bringing runtime threat detection and a new set of AI-era security features. This is the first time the project has targeted the security of AI agents themselves, alongside its established scanning capabilities.

By Matt Saunders

fogus, Mike Fogus: The blog you're reading was rendered at some point in the past as HTML (maybe) but it was written in text with a dash of Markdown flavor (maybe).

@Mystery_Dan, Sam Saffron: I wish I could see the end of the universe. Does it restart? What is nothing without a reference point.

Zack Garvey, PlayStation Blog: PlayStation Store's Spring Sale begins tomorrow, Wednesday, March 25. For a limited time* you can enjoy deep discounts of up to 75% on a vast range of games that include iconic blockbusters, acclaimed classics, and more.

*PlayStation Store Spring Sale promotion runs from Wednesday, March 25 at 00:00 AM PDT/BST/JST and finishes Wednesday April 22 at 11:59 PM PDT/BST/JST**. Each title may have different sale periods. Please refer to the information in the title details page.

**The promotion's start and finish time may differ in certain regions. Game selection may differ by region.

Bruno Couriol, InfoQ: Max Inden recently explored in a talk at FOSDEM 2026 how the upcoming WebTransport protocol and Web API enhance WebSocket capabilities. WebTransport seeks to provide, among other things, lower latency and transparent network switching for key use cases such as high-frequency financial data streaming, cloud gaming, live streaming, and collaborative editing.

By Bruno Couriol

@yossi-shasho Yossi Shasho, Sam Saffron: Actually this command was still too long for me. This is how I finally got it to work:

```
rbtrace -p PID -e 'Thread.new{require "stackprof";s = StackProf;s.start(mode: :cpu);sleep 2;s.stop;s.results("/tmp/perf");}'
```

Craig Risi, InfoQ: Airbnb has revealed how it significantly improved its observability practices by rethinking how alerts are developed and validated, concluding that what appeared to be a "culture problem" was actually a tooling and workflow gap.

By Craig Risi

Sergio De Simone, InfoQ: OpenAI announced they are extending the Responses API to make it easier for developer to build agentic workflows, adding support for a shell tool, a built-in agent execution loop, a hosted container workspace, context compaction, and reusable agent skills.

By Sergio De Simone

@sam Sam Saffron, Sam Saffron: I wish that LLMs progress over the next year and become more powerful, are used only for good and help cure cancer and solve intricate and complex physics problem. They will delight people and enrich our lives, and all the safety are nothing but unfounded fears.

They often face a surge of automated pull requests and security reports with low signal-to-noise ratio.

Kevin Crosby

The Daily Star

Vol. 1, No. 060 · 2026-03-30

Curated and composed automatically.